

Convolutional Neural Network 2


CSE 849 Deep Learning
Spring 2025

Zijun Cui

Notation

- “Channel”: the number of filters
- “Depth”: for each filter, the depth equals to the number of feature maps from the previous layer

Convolutional Filters



Original: Mandrill



Smoothed with
Gaussian kernel



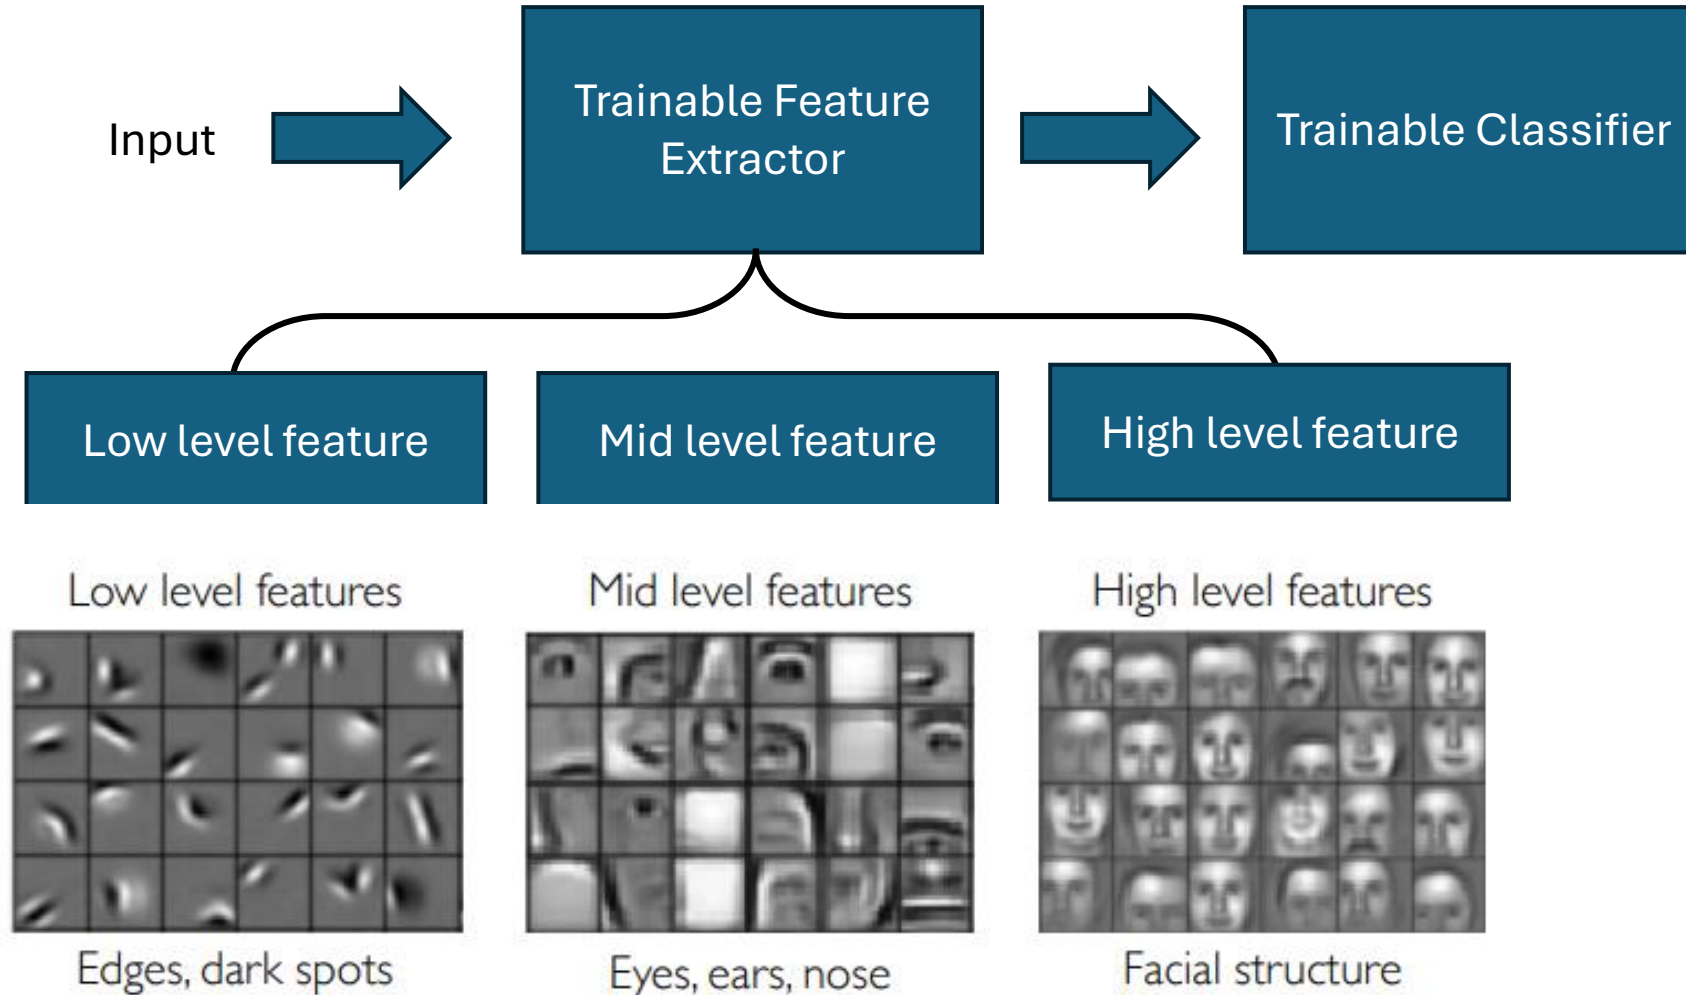
Original Image



Edges

More can be found in computer vision

Recall: Deep Learning = Learning Representations



Outline

- An CNN model for image classification
- More recent CNN architecture designs
 - Group Convolution and ResNeXt
 - Neural Architecture Search and EfficientNets
 - RegNets and NFNets

Setting everything together


- Typical image classification task
 - Assuming maxpooling..

Convolutional Neural Networks



- Input: 1 or 3 images
 - Grey scale or color
 - Will assume color to be generic

Convolutional Neural Networks




- Input: 3 pictures

Preprocessing

- Large images are a problem
 - Too large
 - Compute and memory intensive to process
- Sometimes scaled to smaller sizes, e.g. 128x128 or even 32x32
 - Based on how much will fit on your GPU
 - Typically cropped to *square* images
 - Filters are also typically square


Convolutional Neural Networks



$I \times I$ image

- Input is convolved with a set of K_1 filters
 - Typically K_1 is a power of 2, e.g. 2, 4, 8, 16, 32,...
 - Filters are typically 5x5, 3x3, or even 1x1


Convolutional Neural Networks



K_1 total filters

Filter size: $L \times L \times 3$


Small enough to capture fine features
(particularly important for scaled-down images)



$I \times I$ image

- Input is convolved with a set of K_1 filters
 - Typically K_1 is a power of 2, e.g. 2, 4, 8, 16, 32,...
 - Filters are typically 5x5, 3x3, or even 1x1


Convolutional Neural Networks



K_1 total filters

Filter size: $L \times L \times 3$

Small enough to capture fine features
(particularly important for scaled-down images)




$I \times I$ image

What on earth is this?

- Input is convolved with a set of K_1 filters
 - Typically K_1 is a power of 2, e.g. 2, 4, 8, 16, 32,...
 - Filters are typically 5x5, 3x3, or even 1x1

The 1x1 filter



- A 1x1 filter is simply a perceptron that operates over the *depth* of the stack of maps, but has no spatial extent
 - Takes one pixel from each of the maps (at a given location) as input

Convolutional Neural Networks



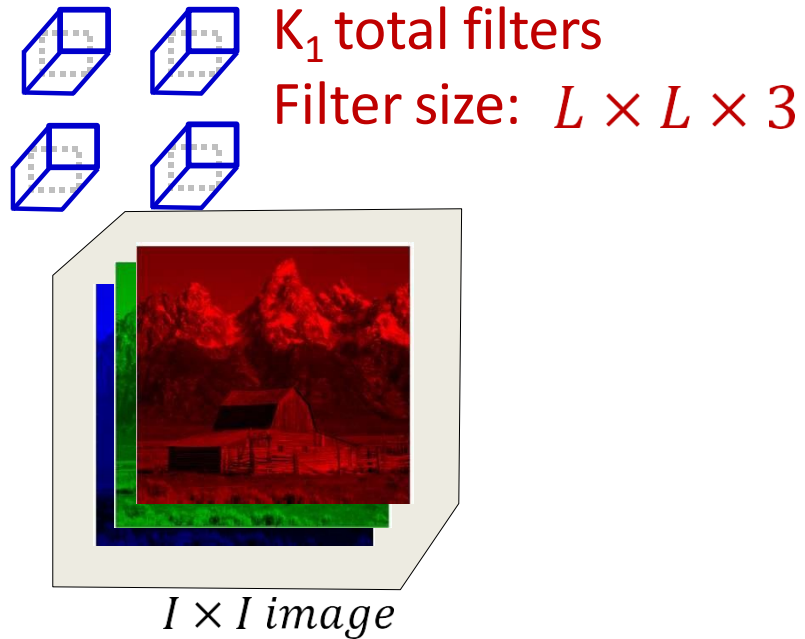
$I \times I$ image

Parameters to choose: K_1 , L and S

1. Number of filters K_1
2. Size of filters $L \times L \times 3 + bias$
3. Stride of convolution S

- Input is convolved with a set of K_1 filters
 - Typically K_1 is a power of 2, e.g. 2, 4, 8, 16, 32,...
 - Filters are typically 5x5(x3), 3x3(x3), or even 1x1(x3)
 - **Typical stride:** 1 or 2

Convolutional Neural Networks





- The input may be zero-padded according to the size of the chosen filters




Convolutional Neural Networks

K_1 filters of size:

$$L \times L \times 3$$



$I \times I$ image



The layer includes a convolution operation followed by an activation (typically RELU)

$$z_m^{(1)}(i, j) = \sum_{c \in \{R, G, B\}} \sum_{k=1}^L \sum_{l=1}^L w_m^{(1)}(c, k, l) I_c(i + k, j + l) + b_m^{(1)}$$

$$Y_m^{(1)}(i, j) = f(z_m^{(1)}(i, j))$$

- **First convolutional layer:** Several convolutional filters
 - Filters are “3-D” (third dimension is color)
 - Convolution followed typically by a RELU activation
- Each filter creates a single 2-D output map




Learnable parameters in the first convolutional layer

- The first convolutional layer comprises K_1 filters, each of size $L \times L \times 3$
 - Spatial span: $L \times L$
 - Depth : 3 (3 colors)
- This represents a total of $K_1(3L^2 + 1)$ parameters
 - “+ 1” because each filter also has a bias
- All of these parameters must be learned

Convolutional Neural Networks

Filter size:

$$L \times L \times 3$$



The layer pools $P \times P$ blocks of $Y_m^{(1)}$ into a single value. It employs a stride D between adjacent blocks.

First pooling/downsampling layer: From each $P \times P$ block of each map, *pool* down to a single value

- For max pooling, during training keep track of which position had the highest value

Convolutional Neural Networks



$$W_m: 3 \times L \times L$$

$$m = 1 \dots K_1$$



$$W_m: K_2 \times L_3 \times L_3$$

$$m = 1 \dots K_3$$




$$z_m^{(n)}(i, j) = \sum_{r=1}^{K_{n-1}} \sum_{k=1}^{L_n} \sum_{l=1}^{L_n} w_m^{(n)}(r, k, l) Y_r^{(n-1)}(i + k, j + l) + b_m^{(n)}$$

$$Y_m^{(n)}(i, j) = f(z_m^{(n)}(i, j))$$

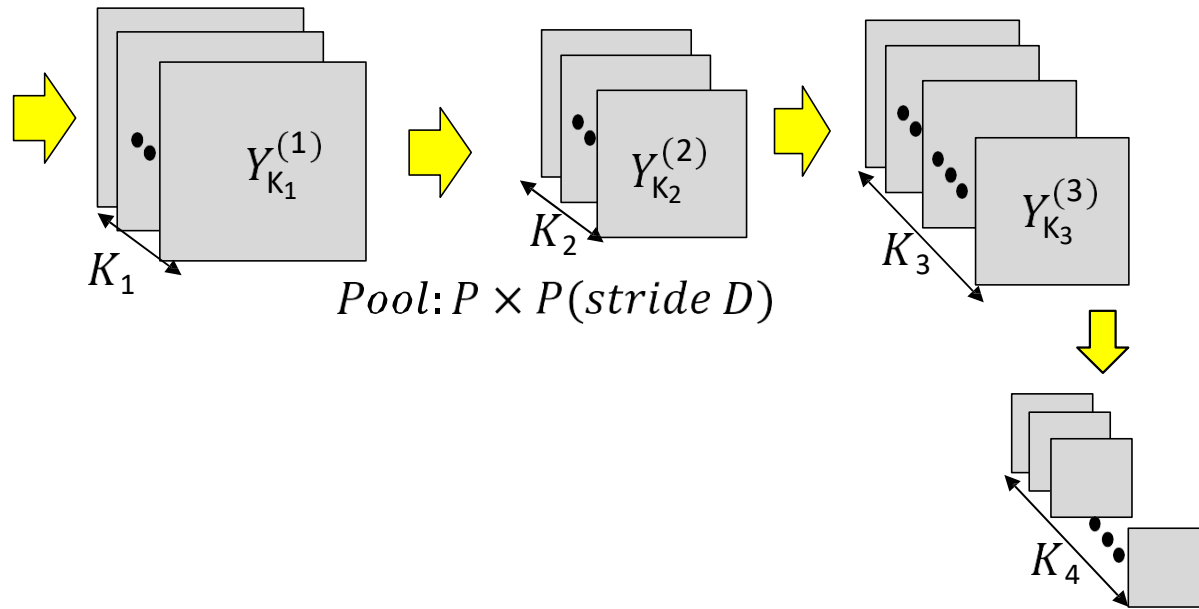
- **Second convolutional layer:** K_3 3-D filters resulting in K_3 2-D maps

Convolutional Neural Networks

$$W_m: 3 \times L \times L$$
$$m = 1 \dots K_1$$




$$W_m: K_2 \times L_3 \times L_3$$
$$m = 1 \dots K_3$$





Convolutional Neural Networks

$$W_m: 3 \times L \times L$$
$$m = 1 \dots K_1$$



$$W_m: K_2 \times L_3 \times L_3$$
$$m = 1 \dots K_3$$



- This continues for several layers until the final convolved output is fed to a softmax
 - Or a full MLP


The Size of the Layers

- Each convolution layer with stride 1 typically maintains the size of the image
 - With appropriate zero padding
 - If performed *without* zero padding it will decrease the size of the input
- Each convolution layer will generally *increase* the *number* of maps from the previous layer
 - In general, the number of convolutional filters increases with layers
- Each pooling layer with stride D *decreases* the *size* of the maps by a factor of D
- Filters within a layer must all be the same size, but sizes may vary with layer
 - Similarly for pooling


Parameters to choose (design choices)

- Number of convolutional and downsampling layers
 - And arrangement (order in which they follow one another)
- For each convolution layer:
 - Number of filters K_i
 - Spatial extent of filter $L_i \times L_i$
 - The “depth” of the filter is fixed by the number of filters in the previous layer K_{i-1}
 - The stride S_i
- For each downsampling/pooling layer:
 - Spatial extent of filter $P_i \times P_i$
 - The stride D_i
- For the final MLP:
 - Number of layers, and number of neurons in each layer


CNN Architectures




AlexNet



VGG




GoogLeNet




ResNet

ImageNet Classification Challenge



ImageNet Classification Challenge



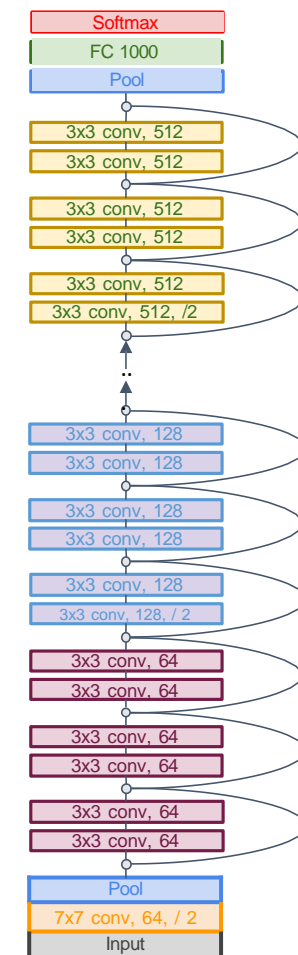

More recent CNN architectures



Post-ResNet Architectures

ResNet made it possible to increase accuracy with larger, deeper models

Many followup architectures emphasize **efficiency**: can we improve accuracy while controlling for model “complexity”?



Measures of Model Complexity

Parameters: How many learnable parameters does the model have?

Floating Point Operations (FLOPs): How many arithmetic operations does it take to compute the forward pass of the model?

Watch out, lots of subtlety here:



- Many papers only count operations in conv layers (ignore ReLU, pooling, BatchNorm)
- Most papers use “1 FLOP” = “1 multiply and 1 addition” so dot product of two N-dim vectors takes N FLOPs
- Other sources (e.g. NVIDIA marketing material) count “1 multiply and one addition” = 2 FLOPs, so dot product of two N-dim vectors takes 2N FLOPs

Network Runtime: How long does a forward pass of the model take on real hardware?

Standard Convolution

- Standard Convolution (groups=1)
- All convolutional kernels touch all C_{in} channels of the input
- Input: $C_{in} \times H \times W$
- Weight: $C_{out} \times C_{in} \times K \times K$
- Output: $C_{out} \times H' \times W'$
- Define: 1 FLOP" = "1 multiply and 1 addition
- For each output element,
$$\text{FLOP} = C_{in} \times K \times K$$
- In total, the number of output elements is
$$C_{out} \times H' \times W'$$
- Hence, the total FLOPS is
$$C_{out} C_{in} K^2 H' W'$$


Comparing Complexity




Key ingredient: Grouped / Separable convolution

Convolution Layer


Each filter has the
same number of
channels as the input



Input: $C_{in} \times H \times W$



Weights: $C_{out} \times C_{in} \times K \times K$




Output: $C_{out} \times H' \times W'$

Recall: Convolution Layer

Each filter has the same number of channels as the input


Each plane of the output depends on the full input and one filter




Recall: Convolution Layer

Each filter has the same number of channels as the input


Each plane of the output depends on the full input and one filter



Input: $C_{in} \times H \times W$



Weights: $C_{out} \times C_{in} \times K \times K$




Output: $C_{out} \times H' \times W'$

Recall: Convolution Layer

Each filter has the same number of channels as the input


Each plane of the output depends on the full input and one filter




Recall: Convolution Layer

Each filter has the same number of channels as the input


Each plane of the output depends on the full input and one filter



Input: $C_{in} \times H \times W$




Weights: $C_{out} \times C_{in} \times K \times K$



Output: $C_{out} \times H' \times W'$

Grouped Convolution

Divide channels of input into G
groups with (C_{in}/G) channels each




Example:
 $G = 2$

Input: $C_{in} \times H \times W$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each



Example:
 $G = 2$

Divide filters into G groups;
each group looks at a **subset** of input channels




Input: $C_{in} \times H \times W$

Weights: $C_{out} \times (C_{in}/G) \times K \times K$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each




Example:
 $G = 2$

Divide filters into G groups;
each group looks at a
subset of input channels



Each plane of the output
depends on one filter and a
subset of the input channels




Input: $C_{in} \times H \times W$

Weights: $C_{out} \times (C_{in}/G) \times K \times K$

Output: $C_{out} \times H' \times W'$

Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each



Example:
 $G = 2$


H

W

C_{in}


Input: $C_{in} \times H \times W$

Divide filters into G groups;
each group looks at a
subset of input channels



Weights: $C_{out} \times (C_{in}/G) \times K \times K$


Each plane of the output
depends on one filter and a
subset of the input channels



Output: $C_{out} \times H' \times W'$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each




Input: $C_{in} \times H \times W$

Divide filters into G groups;
each group looks at a **subset** of input channels



Weights: $C_{out} \times (C_{in}/G) \times K \times K$


Each plane of the output depends on one filter and a **subset** of the input channels



Output: $C_{out} \times H' \times W'$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each




Input: $C_{in} \times H \times W$

Divide filters into G groups;
each group looks at a **subset** of input channels



Weights: $C_{out} \times (C_{in}/G) \times K \times K$


Each plane of the output depends on one filter and a **subset** of the input channels



Output: $C_{out} \times H' \times W'$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each




Input: $C_{in} \times H \times W$

Divide filters into G groups;
each group looks at a **subset** of input channels



Weights: $C_{out} \times (C_{in}/G) \times K \times K$


Each plane of the output depends on one filter and a **subset** of the input channels



Output: $C_{out} \times H' \times W'$


Grouped Convolution

Divide channels of input into G **groups** with (C_{in}/G) channels each



Input: $C_{in} \times H \times W$


Example:
 $G = 4$



Weights: $C_{out} \times (C_{in}/G) \times K \times K$

Divide filters into G groups;
each group looks at a **subset** of input channels


Each plane of the output
depends on one filter and a **subset** of the input channels



Output: $C_{out} \times H' \times W'$


Special Case: Depthwise Convolution

Number of groups equals
number of input channels




Input: $C_{in} \times H \times W$

Common to also set $C_{out} = G$



Weights: $C_{out} \times 1 \times K \times K$


Output only mixes *spatial*
information from input;
channel information not mixed



Output: $C_{out} \times H' \times W'$


Special Case: Depthwise Convolution

Number of groups equals
number of input channels




Input: $C_{in} \times H \times W$

Can still have multiple filters per group
(e.g. $C_{out} = 2C_{in}$)



Weights: $C_{out} \times 1 \times K \times K$

Output only mixes *spatial*
information from input;
channel information not mixed



Output: $C_{out} \times H' \times W'$

Grouped Convolution vs Standard Convolution

Grouped Convolution (G groups):

G parallel conv layers; each “sees”

C_{in}/G input channels and produces

C_{out}/G output channels

Input: $C_{in} \times H \times W$

Split to $G \times [(C_{in} / G) \times H \times W]$

Weight: $G \times (C_{out} / G) \times (C_{in} \times G) \times K \times K$

G parallel convolutions

Output: $G \times [(C_{out} / G) \times H' \times W']$

Concat to $C_{out} \times H' \times W'$

FLOPs: $C_{out}C_{in}K^2H'W'/G$

• Standard Convolution (groups=1)

• All convolutional kernels touch all C_{in} channels of the input

• Input: $C_{in} \times H \times W$

• Weight: $C_{out} \times C_{in} \times K \times K$

• Output: $C_{out} \times H' \times W'$

• **FLOPs: $C_{out}C_{in}K^2H'W'$**

Using G groups reduces FLOPs by a factor of G!

Grouped Convolution in PyTorch


PyTorch convolution gives an option for groups!

Conv2d

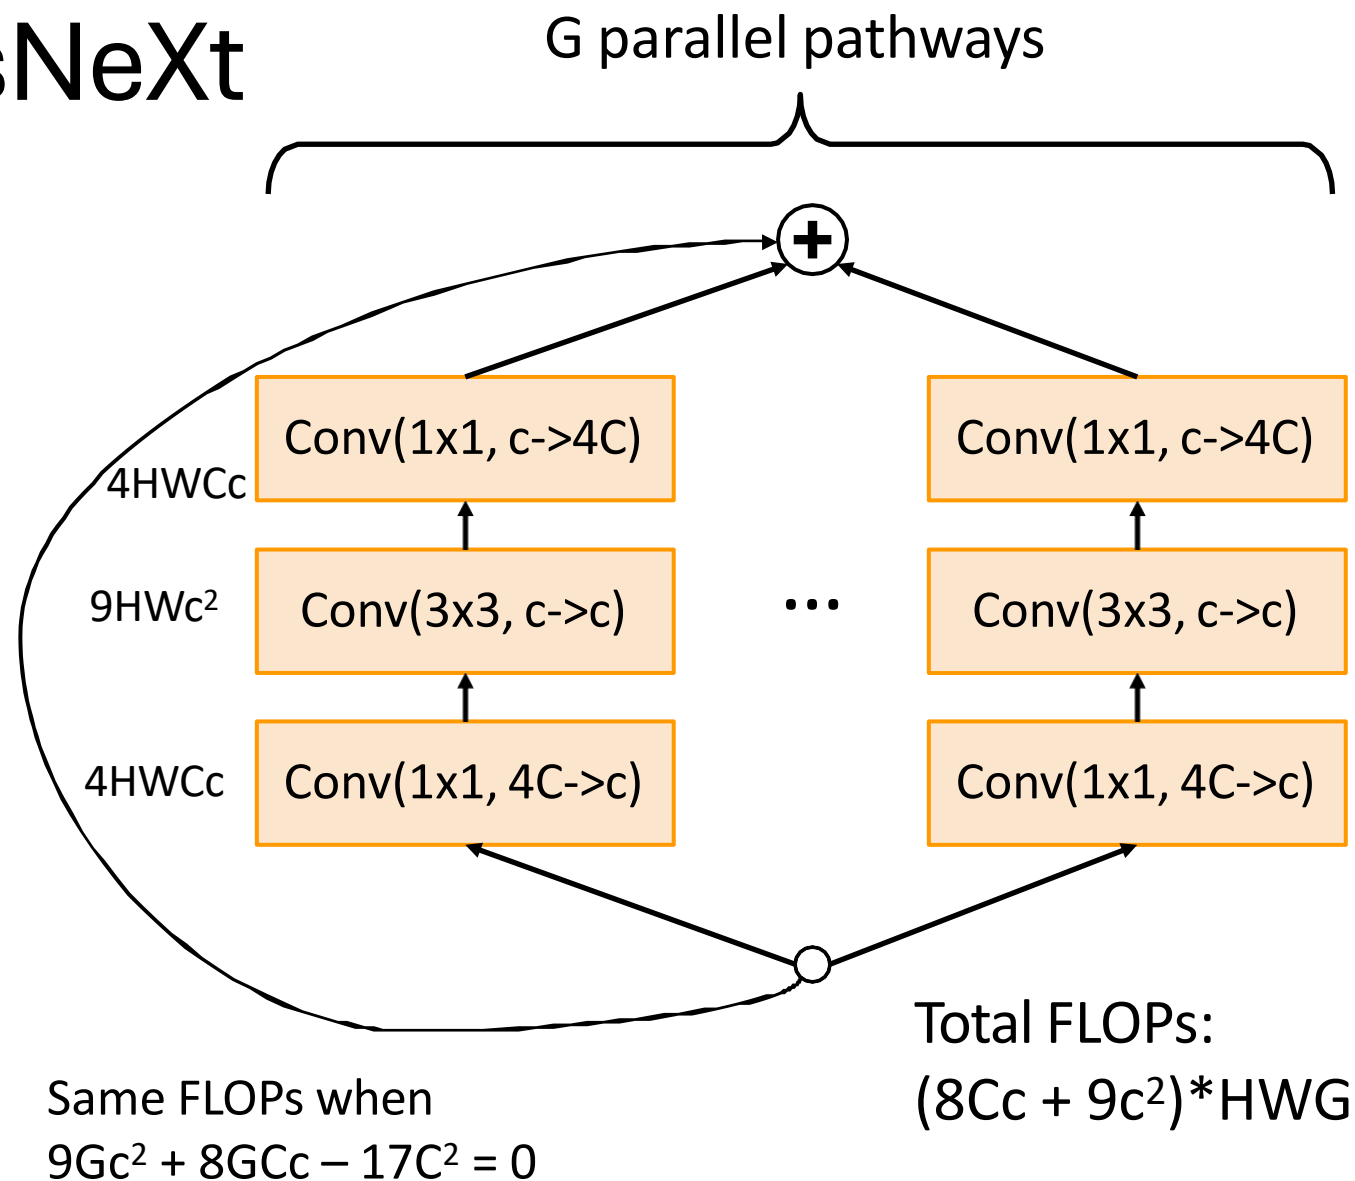
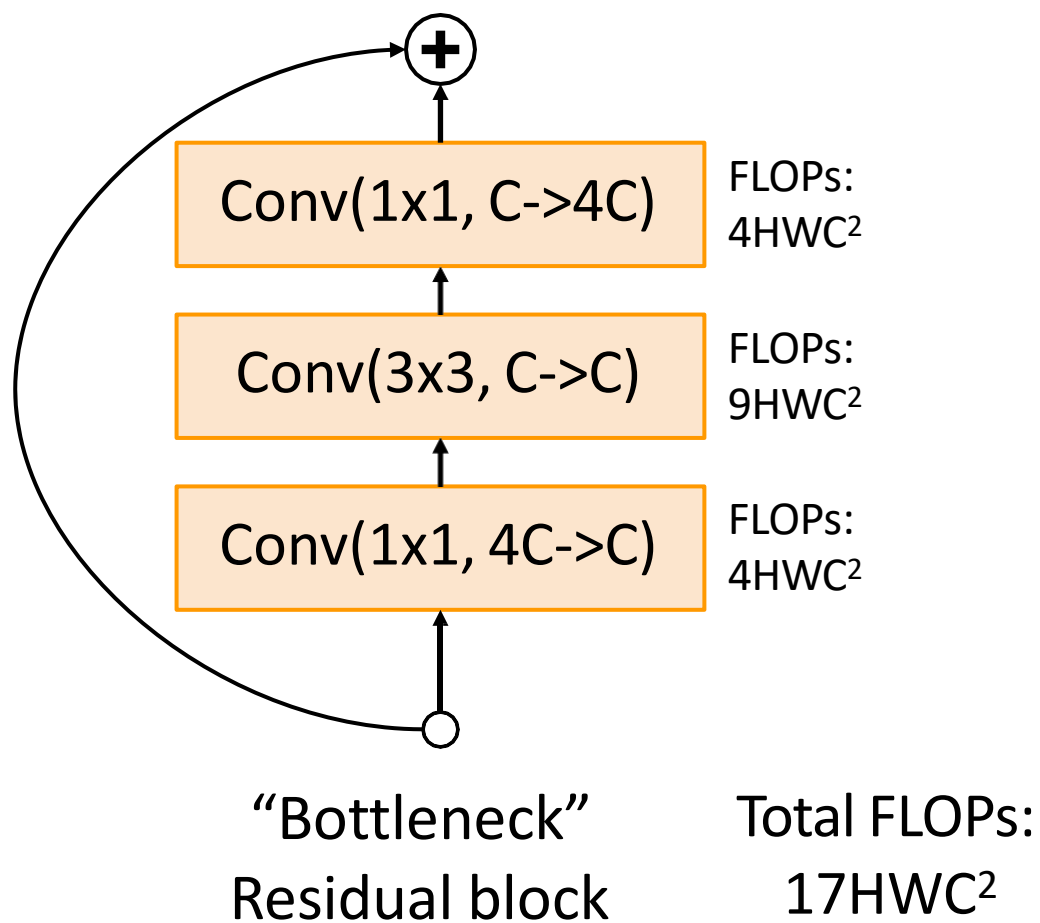
```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size,  
stride=1, padding=0, dilation=1, groups=1, bias=True,  
padding_mode='zeros')
```

[\[SOURCE\]](#)

Improving ResNets

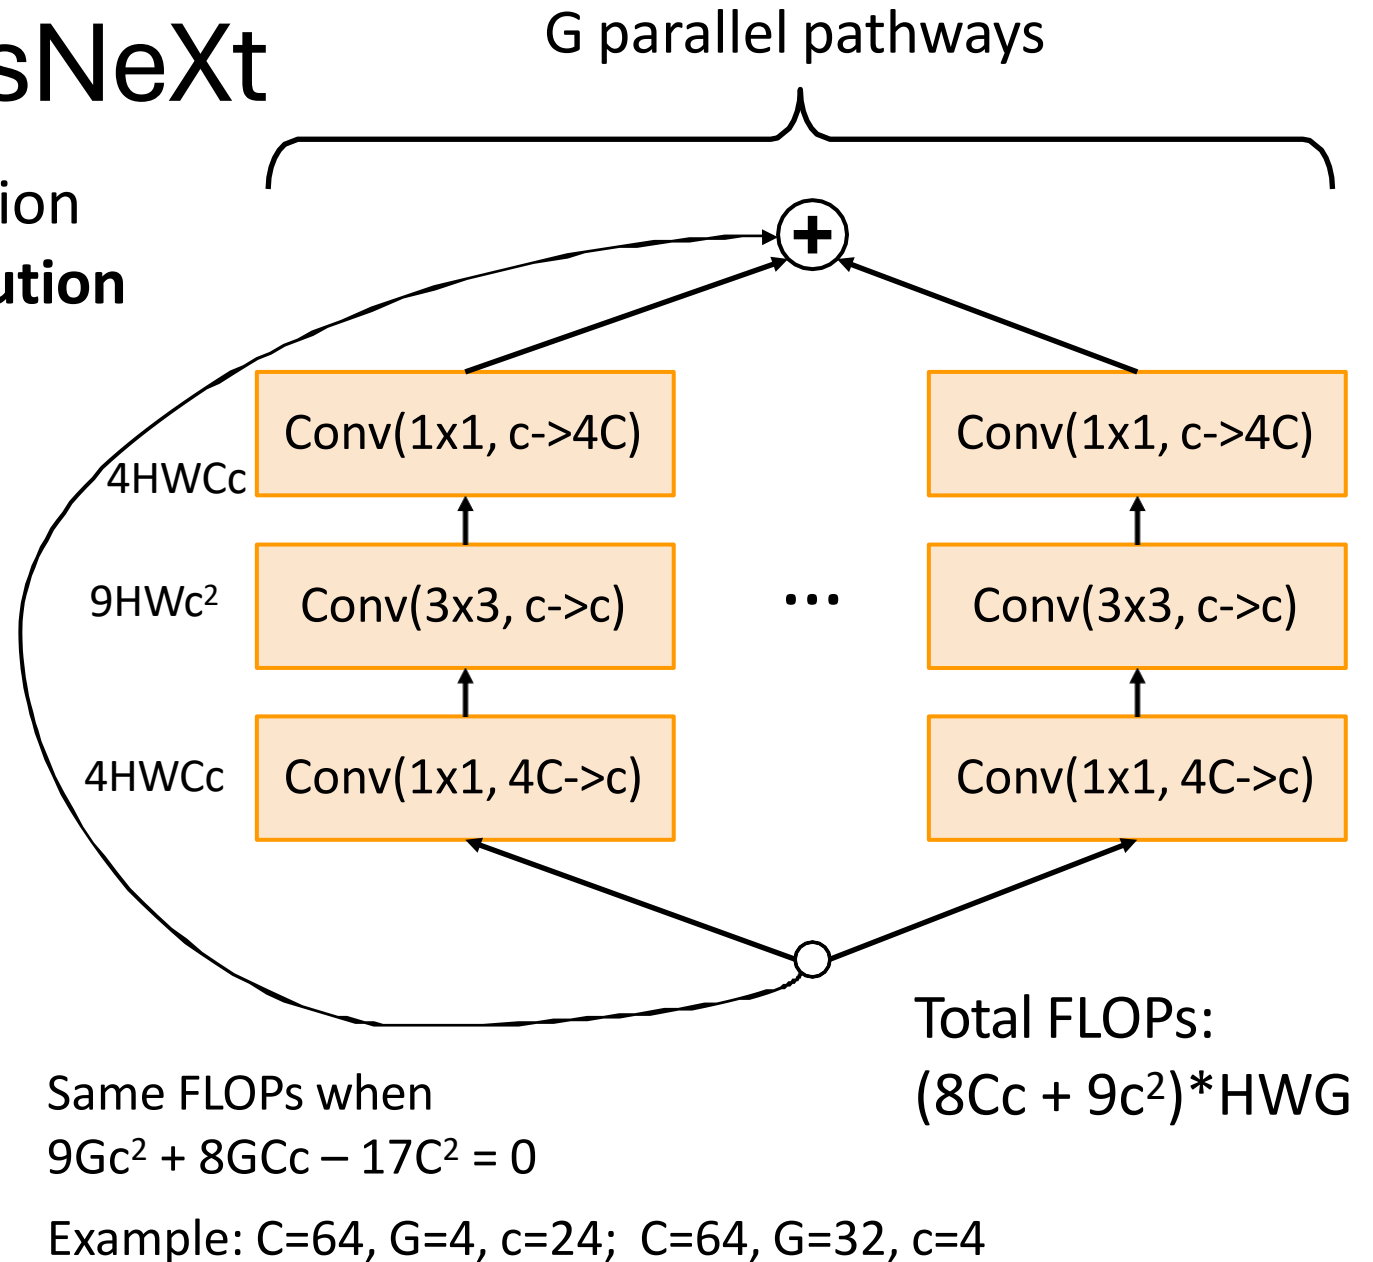
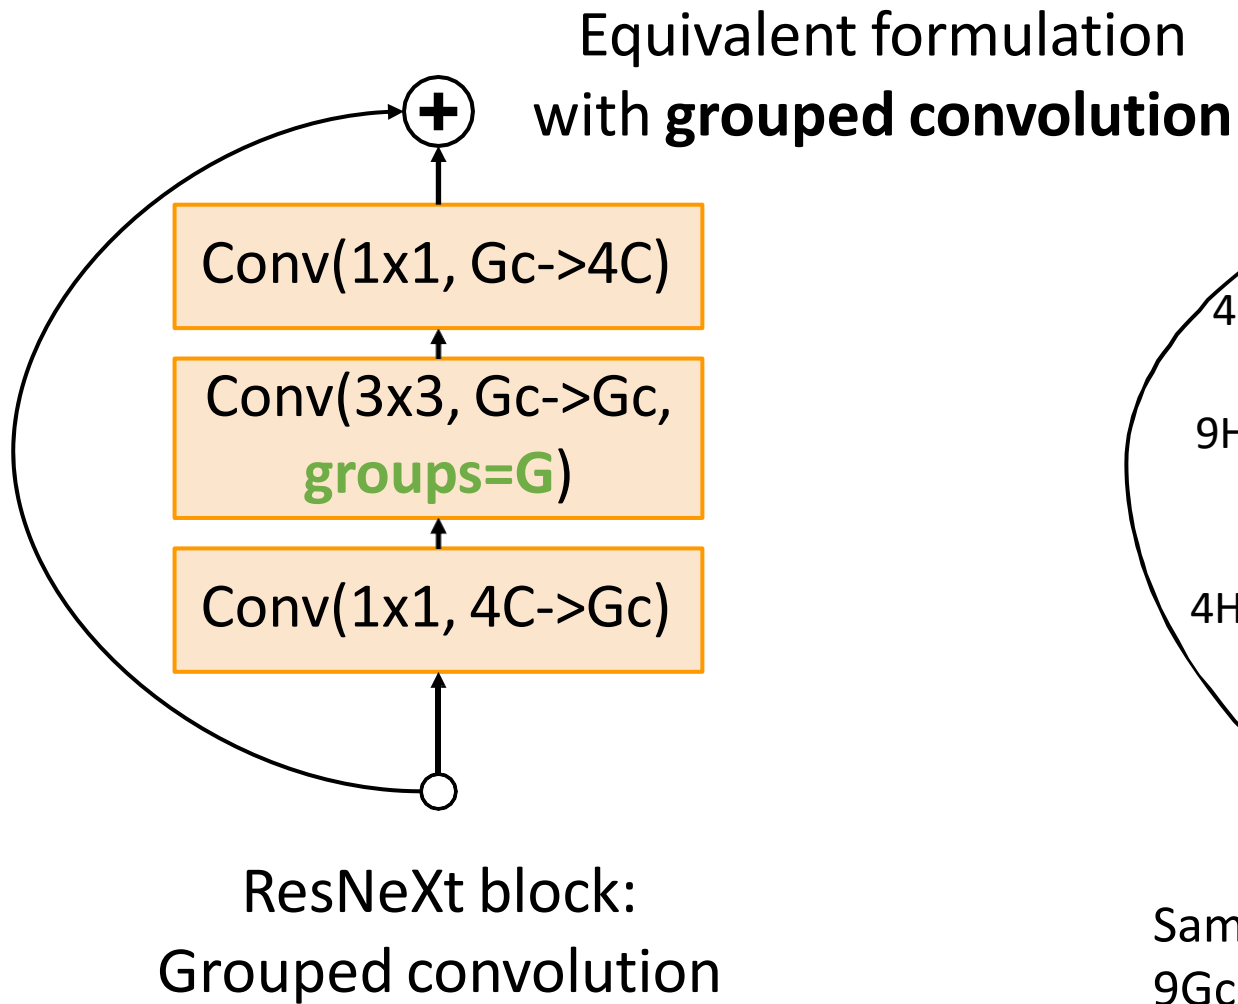


Improving ResNets: ResNeXt



Example: $C=64$, $G=4$, $c=24$; $C=64$, $G=32$, $c=4$

Improving ResNets: ResNeXt



ResNeXt: Maintain computation by adding groups!

Model	Groups	Group width	Top-1 Error
ResNet-50	1	64	23.9
ResNeXt-50	2	40	23
ResNeXt-50	4	24	22.6
ResNeXt-50	8	14	22.3
ResNeXt-50	32	4	22.2

Model	Groups	Group width	Top-1 Error
ResNet-101	1	64	22.0
ResNeXt-101	2	40	21.7
ResNeXt-101	4	24	21.4
ResNeXt-101	8	14	21.3
ResNeXt-101	32	4	21.2

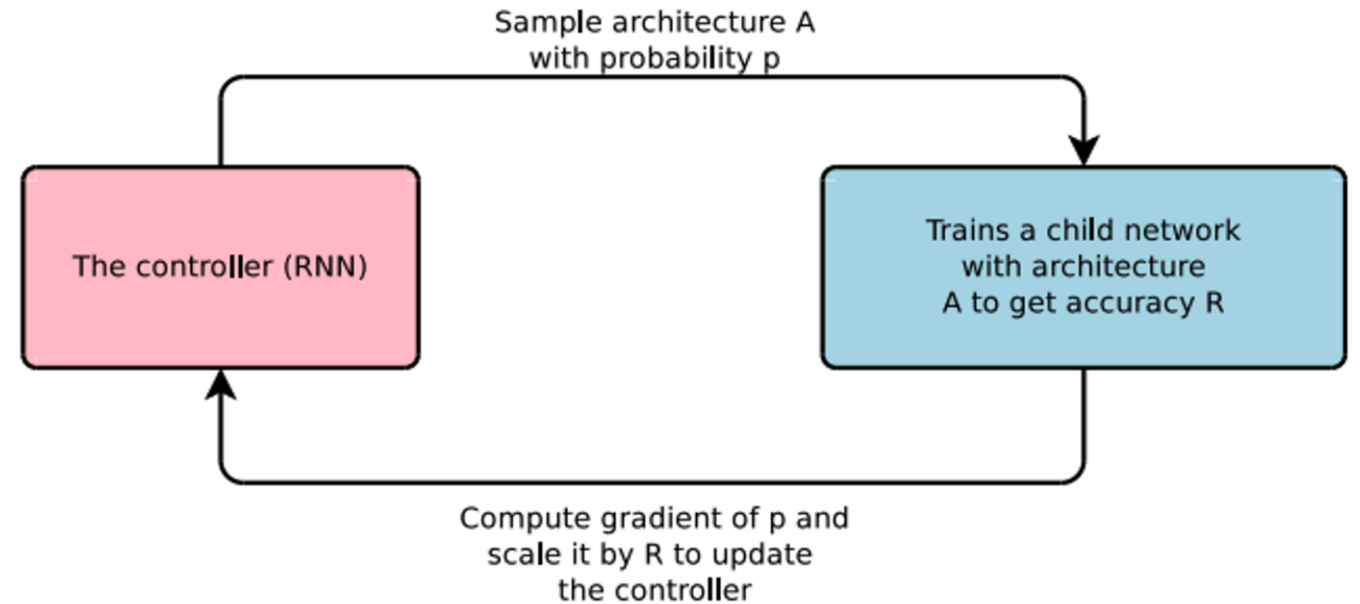
Adding groups improves performance **with same FLOPs!**

Often denoted e.g. ResNeXt-50-32x4d: 32 groups,
Blocks in first stage have 4 channels per group (#channels still doubles at each stage)

Neural Architecture Search (NAS)

Designing neural network architectures is hard – let's automate it!

- One network (**controller**) outputs network architectures
- Sample **child networks** from controller and train them
- After training a batch of child networks, make a gradient step on controller network (Using **policy gradient**)
- Over time, controller learns to output good architectures!




Neural Architecture Search (NAS)

- Search for reusable “block” designs which can use the following operators:
 - Identity
 - 1x1 conv
 - 3x3 conv
 - 3x3 dilated conv
 - 1x7 then 7x1 conv
 - 1x3 then 3x1 conv
 - 3x3, 5x5, or 7x7 depthwise-separable conv
 - 3x3 avg pool
 - 3x3, 5x5, or 7x7 max pool

The “Normal cell” maintains the same image resolution

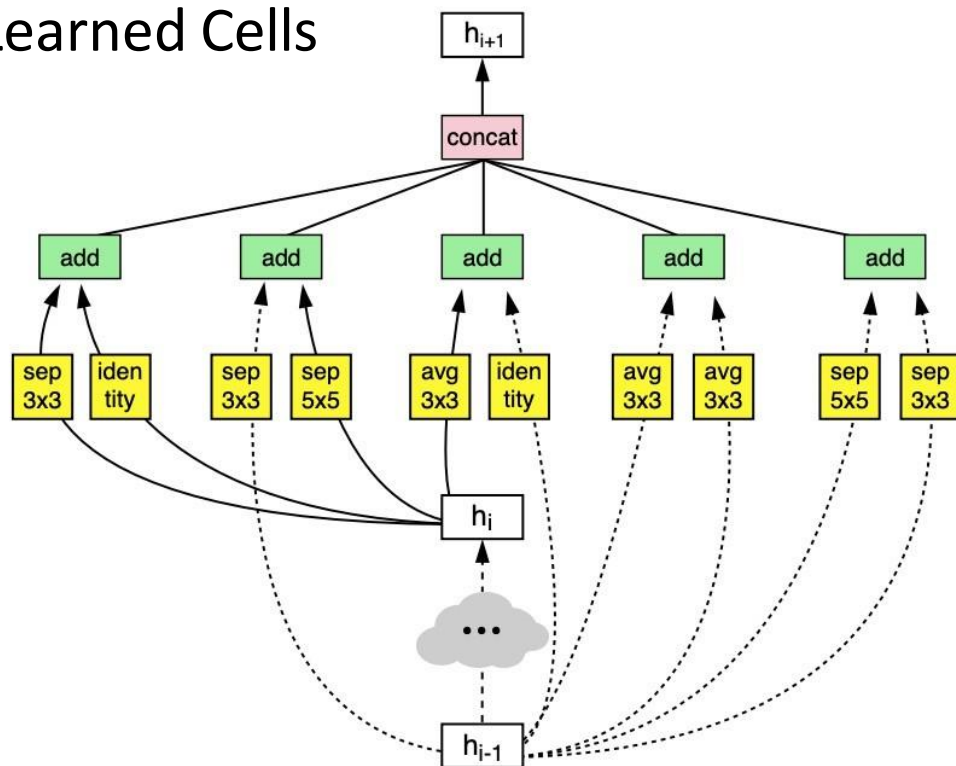
The “Reduction cell” reduces image resolution by 2x

Combine two learned cells in a regular pattern to create overall architecture

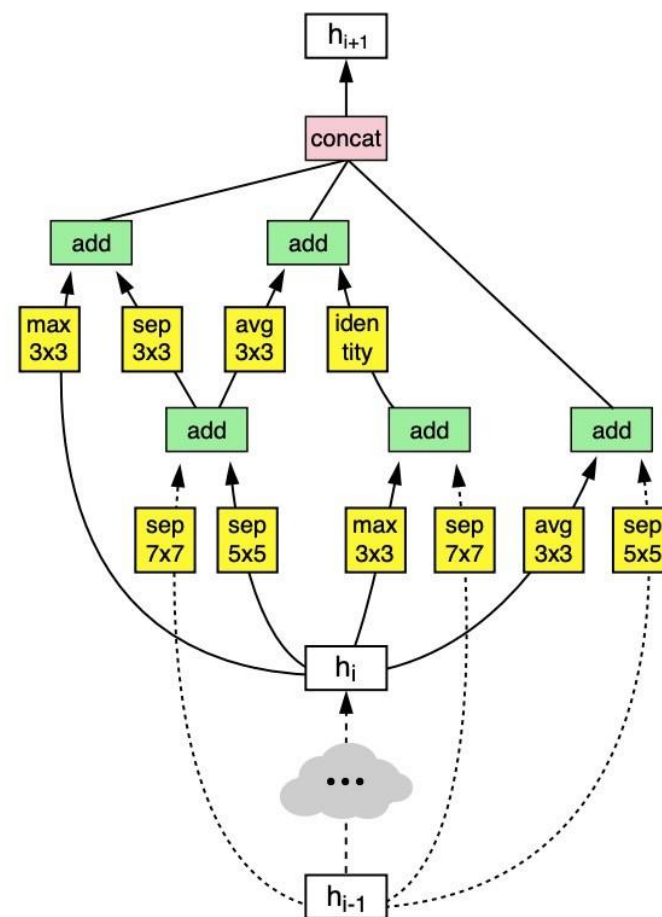


Neural Architecture Search (NAS)


Learned Cells



Normal Cell





Reduction Cell



Zoph and Le, "Neural Architecture Search with Reinforcement Learning", ICLR 2017

Zoph et al, "Learning transferable architectures for scalable image recognition", CVPR 2018

Neural Architecture Search (NAS)



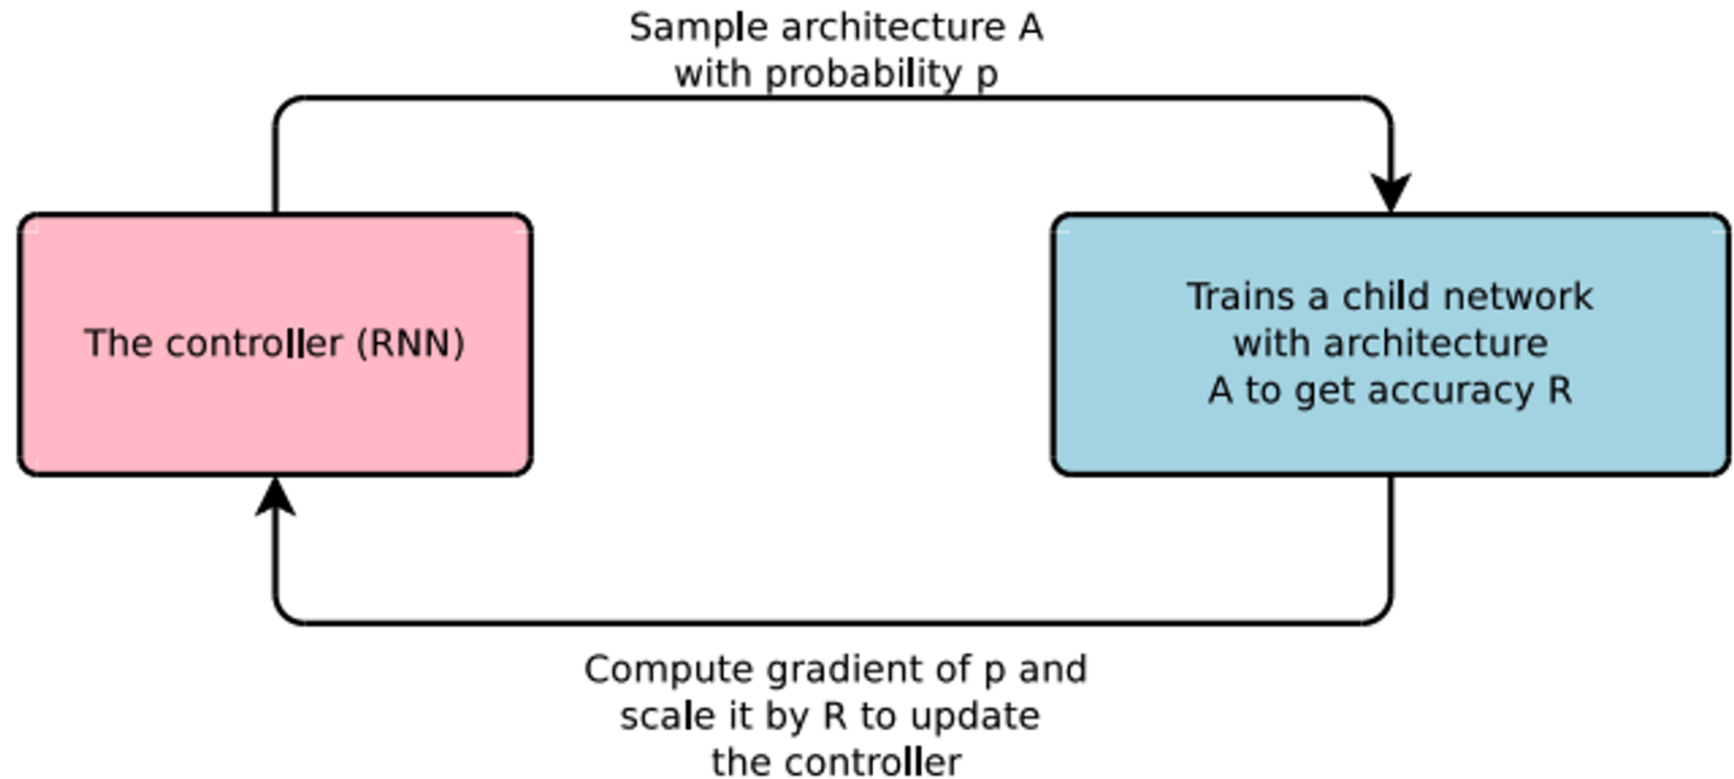
Zoph and Le, "Neural Architecture Search with Reinforcement Learning", ICLR 2017

Zoph et al, "Learning transferable architectures for scalable image recognition", CVPR 2018

Big Problem: NAS is Very Expensive!


Original NAS paper: Each update to the controller requires training **800 child models** for 50 epochs on CIFAR10; Total of **12,800** child models are trained

Later work improved efficiency, but still expensive





Model Scaling




Starting from a given architecture, how should you **scale it up** to improve performance?



Model Scaling: EfficientNets




Model Scaling: EfficientNets



Scaling any of width, depth, or resolution has diminishing returns.
For optimal results, need to scale them all jointly!

Model Scaling: EfficientNets



Big problem: *Real-world runtime does not correlate well with FLOPs!*

- Runtime depends on the device (mobile CPU, server CPU, GPU, TPU); A model which is fast on one device may be slow on another
- Depthwise convolutions are efficient on mobile, but not on GPU / TPU – they become memory-bound
- The “naïve” FLOP counting we have done for convolutions can be incorrect – alternate conv algorithms can reduce FLOPs in some settings (FFT for large kernels, Winograd for 3x3 conv)
- EfficientNet was designed to minimize FLOPs, not actual runtime – so it is surprisingly slow!

Beyond NAS – back to hand-designed models!


Instead of using NAD
smartly tweak ResNet-style models to improve performance, scaling,
runtime on GPU / TPU

RegNets: Simple block design, optimize macro architecture and scaling

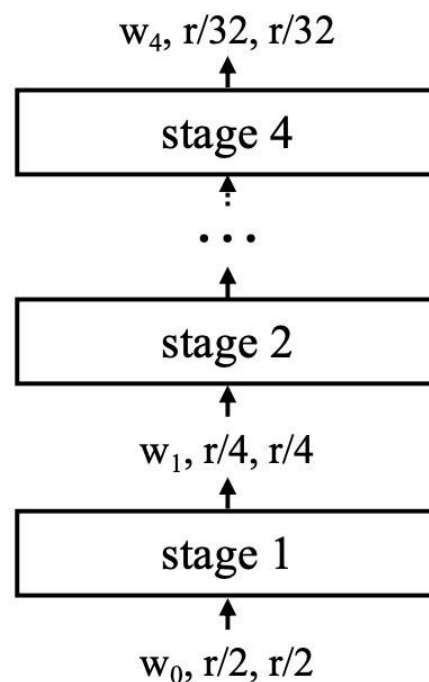
NFNets: Remove Batch Normalization

RegNets: Network Design Spaces

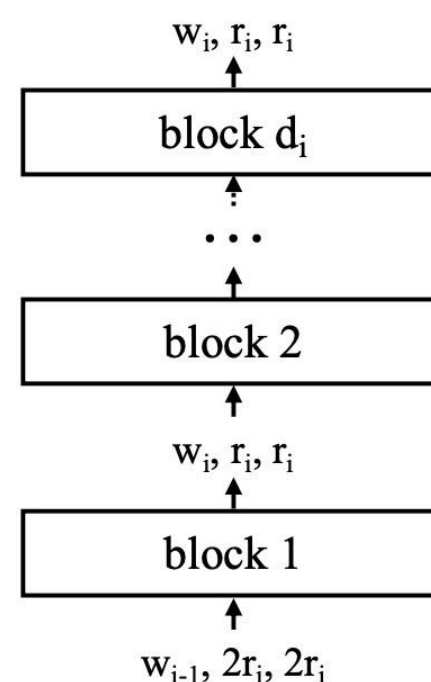
Network design is simple: **Stem** of 3x3 convs, a **body** of 4 *stages*, and a **head**; Each stage has multiple **blocks**: First block downsamples by 2x, others keep resolution the same



(a) network



(b) body



(c) stage i

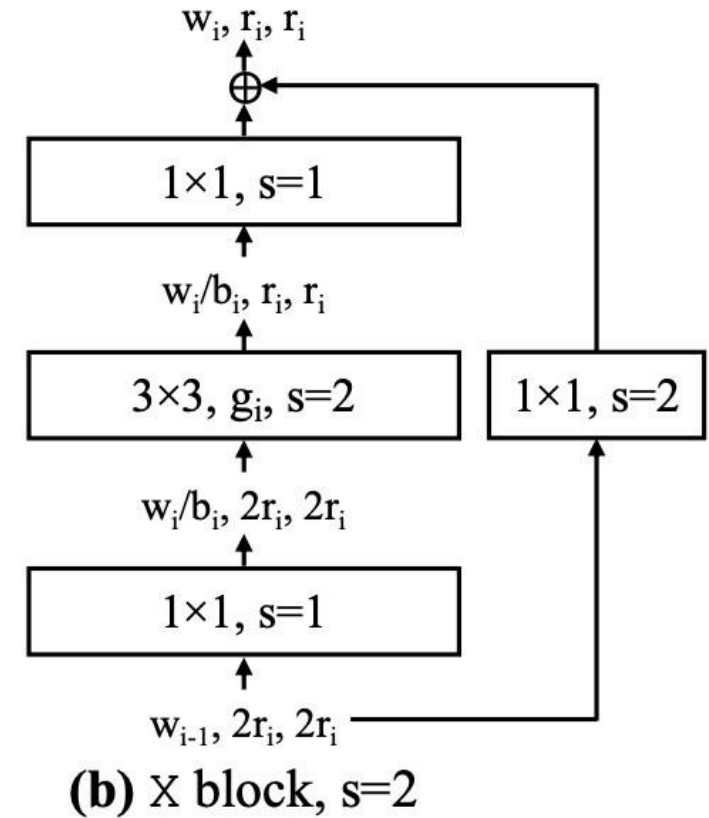

RegNets: Network Design Spaces

Block design is simple,
generalizes ResNext

Each stage has 4 parameters:



- Number of blocks
- Number of input channels w
- Bottleneck ratio b
- Group width g

The *design space* for the network
has just 16 parameters



RegNets: Network Design Spaces

Randomly sample architectures from the design space, examine trends:



RegNets: Network Design Spaces



Use results to *refine* the design space: Reduce degrees of freedom from 16 to bias toward better-performing architectures:

- Share bottleneck ratio across all stages (16 \rightarrow 13 params)
- Share group width across all stages (13 \rightarrow 10 params)
- Force width, blocks per stage to increase *linearly* across stages

Final design space has 6 parameters:

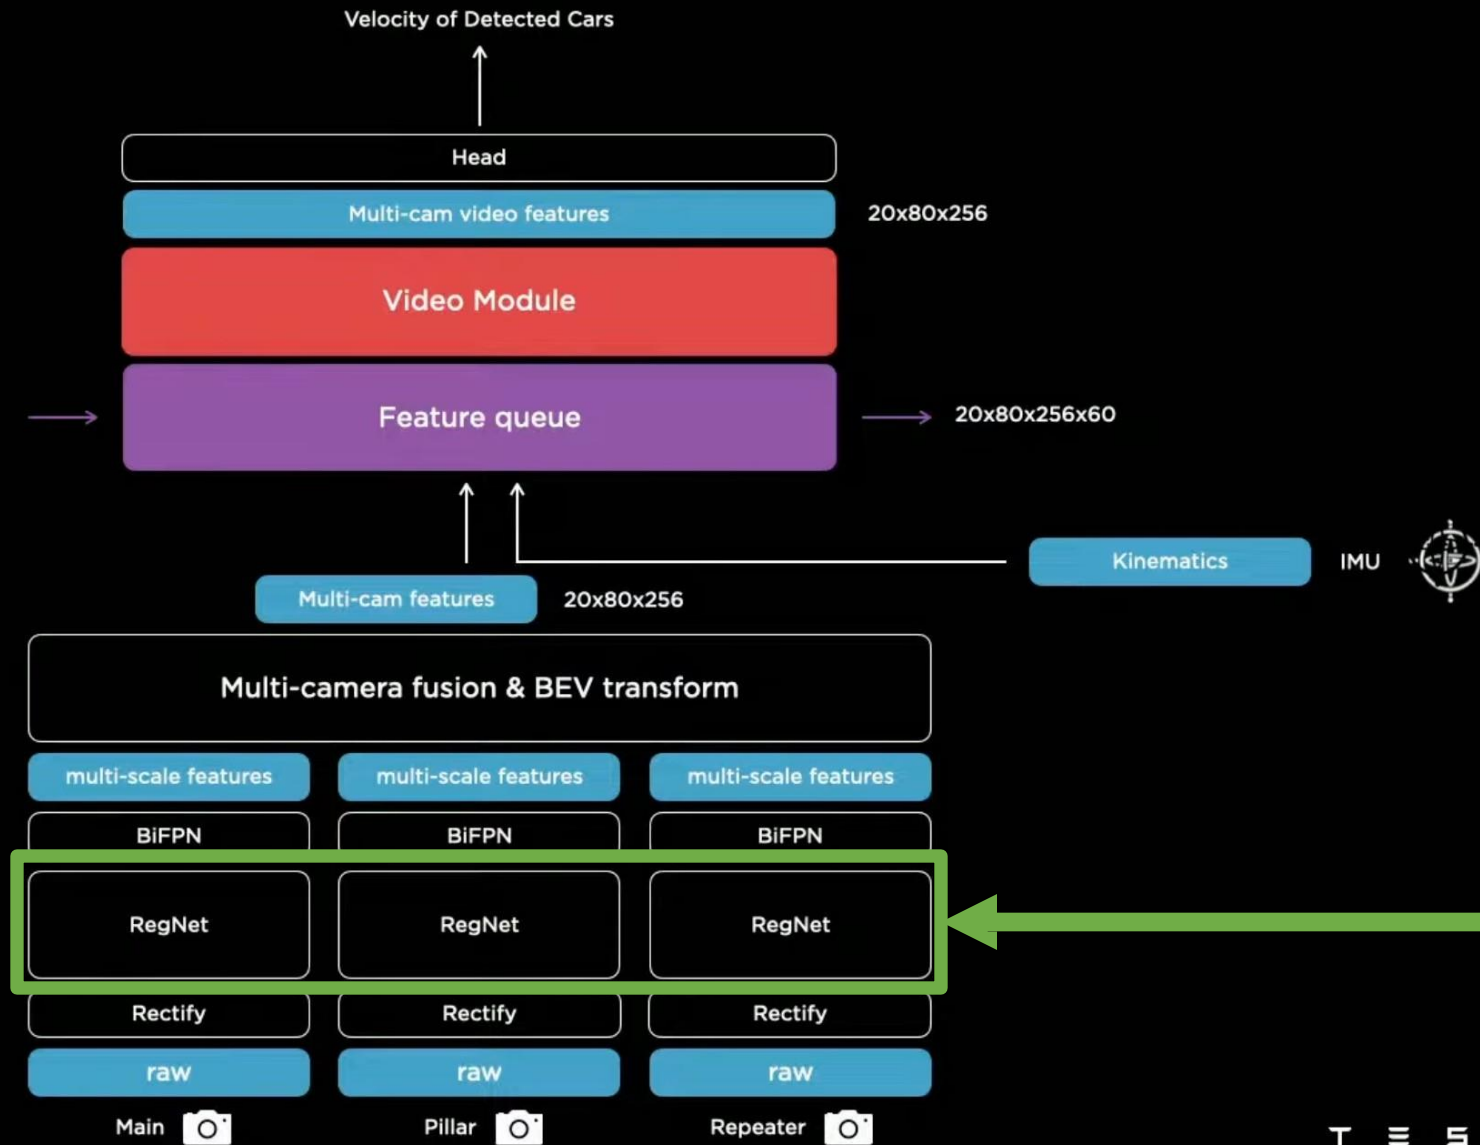
- Overall depth d , bottleneck ratio b , group width g
- Initial width w_0 , width growth rate w_a , blocks per stage w_m

RegNets: Network Design Spaces



At same FLOPs, RegNet models get similar accuracy as EfficientNets
but are up to 5x faster in training (each iteration is faster)

Video Neural Net Architecture



Tesla Vision system
uses RegNets to
process inputs
from each camera

Tesla AI Day 2021,
<https://www.youtube.com/watch?v=j0z4FweCy4M>

Training ResNets without Batch Normalization

- Batch Normalization has good properties:
 - Makes it easy to train deep networks ≥ 10 layers
 - Makes learning rates, initialization less critical
 - Adds regularization
 - "Free" at inference: can be merged into linear layers
- But also has bad properties:
 - Doesn't work with small minibatches
 - Different behavior at train and test
 - Slow at training time

NFNets are ResNets without Batch Normalization!

Brock et al, "Characterizing Signal Propagation to Close the Performance Gap in Unnormalized ResNets", ICLR 2021


Brock et al, "High-Performance Large-Scale Image Recognition without Normalization", ICML 2021

NFNets

Consider a pre-activation ResNet block $x_{\ell+1} = f_{\ell}(x_{\ell}) + x_{\ell}$

Problem: Variance grows with each block:

$$\text{Var}(x_{\ell+1}) = \text{Var}(x_{\ell}) + \text{Var}(f_{\ell}(x_{\ell}))$$



Brock et al, "Characterizing Signal Propagation to Close the Performance Gap in Unnormalized ResNets", ICLR 2021

Brock et al, "High-Performance Large-Scale Image Recognition without Normalization", ICML 2021

He et al, "Identity Mappings in Deep Residual Networks", ECCV 2016

NFNets: Scaled Residual Blocks

Consider a pre-activation ResNet block $x_{\ell+1} = f_{\ell}(x_{\ell}) + x_{\ell}$


Problem: Variance grows with each block:

$$\text{Var}(x_{\ell+1}) = \text{Var}(x_{\ell}) + \text{Var}(f_{\ell}(x_{\ell}))$$

Solution: Re-parameterize block:

$$x_{\ell+1} = x_{\ell} + \alpha f_{\ell}(x_{\ell}/\beta_{\ell})$$

α is a hyperparameter, $\beta_{\ell} = \sqrt{\text{Var}(x_{\ell})}$ at initialization;
both are constants during training



NFNets: Scaled Residual Blocks

Consider a pre-activation ResNet block $x_{\ell+1} = f_{\ell}(x_{\ell}) + x_{\ell}$

Problem: Variance grows with each block:

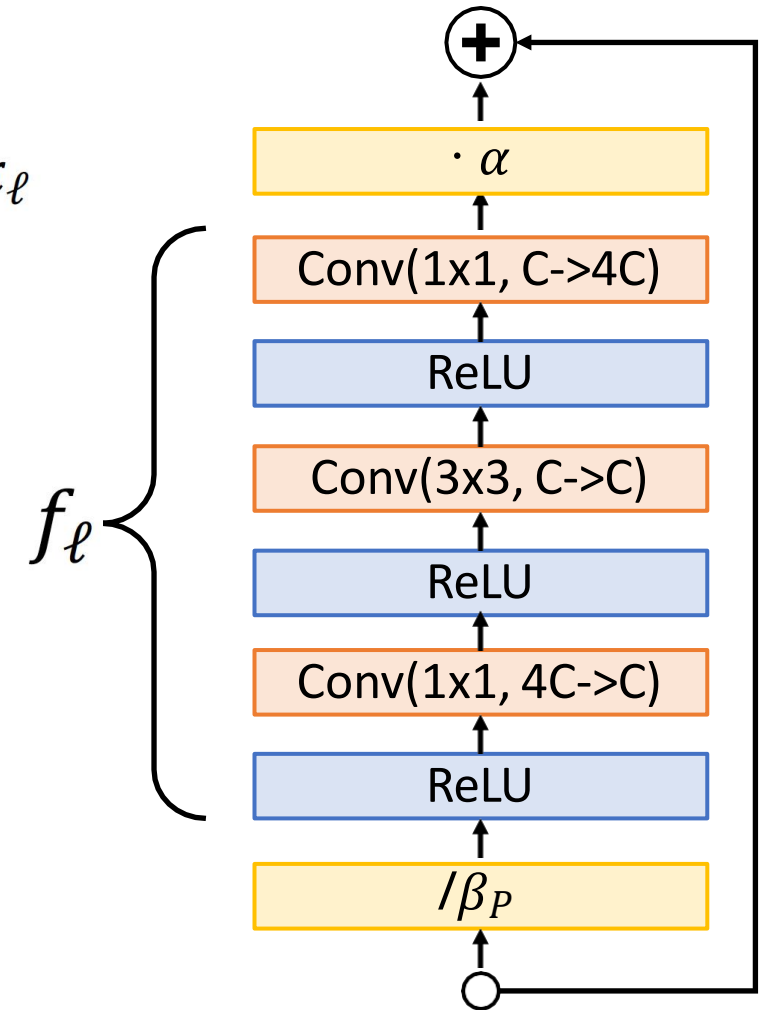
$$\text{Var}(x_{\ell+1}) = \text{Var}(x_{\ell}) + \text{Var}(f_{\ell}(x_{\ell}))$$

Solution: Re-parameterize block:


$$x_{\ell+1} = x_{\ell} + \alpha f_{\ell}(x_{\ell}/\beta_{\ell})$$

α is a hyperparameter, $\beta_{\ell} = \sqrt{\text{Var}(x_{\ell})}$ at initialization;
both are constants during training

Now $\text{Var}(x_{\ell+1}) = \text{Var}(x_{\ell}) + \alpha^2$; resets to $1 + \alpha^2$
after each downsampling block



NFNet



Always be careful with plots like this
– different papers use different metric for x-axis:

- FLOPs
- Params
- Test-time runtime
- Training-time runtime
- Runtime on CPU / GPU / TPU /...