

Recurrent Neural Networks


CSE 849 Deep Learning
Spring 2025

Zijun Cui

Modelling Series

- In many situations one must consider a *series* of inputs to produce an output
 - Outputs too may be a series

What did I say?



- Speech Recognition
 - Analyze a series of spectral vectors, determine what was said
- Note: Inputs are sequences of vectors. Output is a classification result


What is he talking about?

The Steelers, meanwhile, continue to struggle to make stops on defense. They've allowed, on average, 30 points a game, and have shown no signs of improving anytime soon.

“Football” or “basketball”?


- Text analysis
 - E.g. analyze document, identify topic
 - Input series of words, output classification output
 - E.g. read English, output French
 - Input series of words, output series of words

Should I invest..




- Note: Inputs are sequences of vectors. Output may be scalar or vector
 - Should I invest, vs. should I not invest in X?
 - Decision must be taken considering how things have fared over time
- Must consider the series of stock values in the past several days to decide if it is wise to invest today

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network



- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

The stock predictor network




- The sliding predictor
 - Look at the last few days
 - This is just a convolutional neural net applied to series data
 - Also called a *Time-Delay neural network*

Finite-response model

- This is a *finite response* system
 - Something that happens *today* only affects the output of the system for N days into the future
 - N is the *width* of the system


$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-N})$$

Finite-response model




- Something that happens *today* only affects the output of the system for N days into the future
 - **Predictions consider N days of history**
- To consider more of the past to make predictions, you must increase the “history” considered by the system

Finite-response




- Problem: Increasing the “history” makes the network more complex
 - require more computational resources

Systems often have long-term dependencies



- Longer-term trends –
 - Weekly trends in the market
 - Monthly trends in the market
 - Annual trends

We want *infinite* memory




- Required: *Infinite* response systems
 - What happens today can continue to affect the output forever
 - Possibly with weaker and weaker influence

$$Y_t = f(X_t, X_{t-1}, \dots, X_{t-\infty})$$


Examples of infinite response systems

$$Y_t = f(X_t, Y_{t-1})$$


- Required: Define initial state: Y_{-1} for $t = 0$
 - An input X_0 at $t = 0$ produces Y_0
 - Y_0 produces Y_1 which produces Y_2 and so on until Y_∞ *even if* $X_1 \dots X_\infty$ are 0
 - i.e. even if there are no further inputs!
 - **A single input influences the output for the rest of time!**
- This is an instance of a NARX network
 - “nonlinear autoregressive network with exogenous inputs”
 - $Y_t = f(X_{0:t}, Y_{0:t-1})$
 - *Output* contains information about the entire past




- A NARX net with recursion from the output




- A NARX net with recursion from the output




- A NARX net with recursion from the output




- A NARX net with recursion from the output




- A NARX net with recursion from the output



- A NARX net with recursion from the output

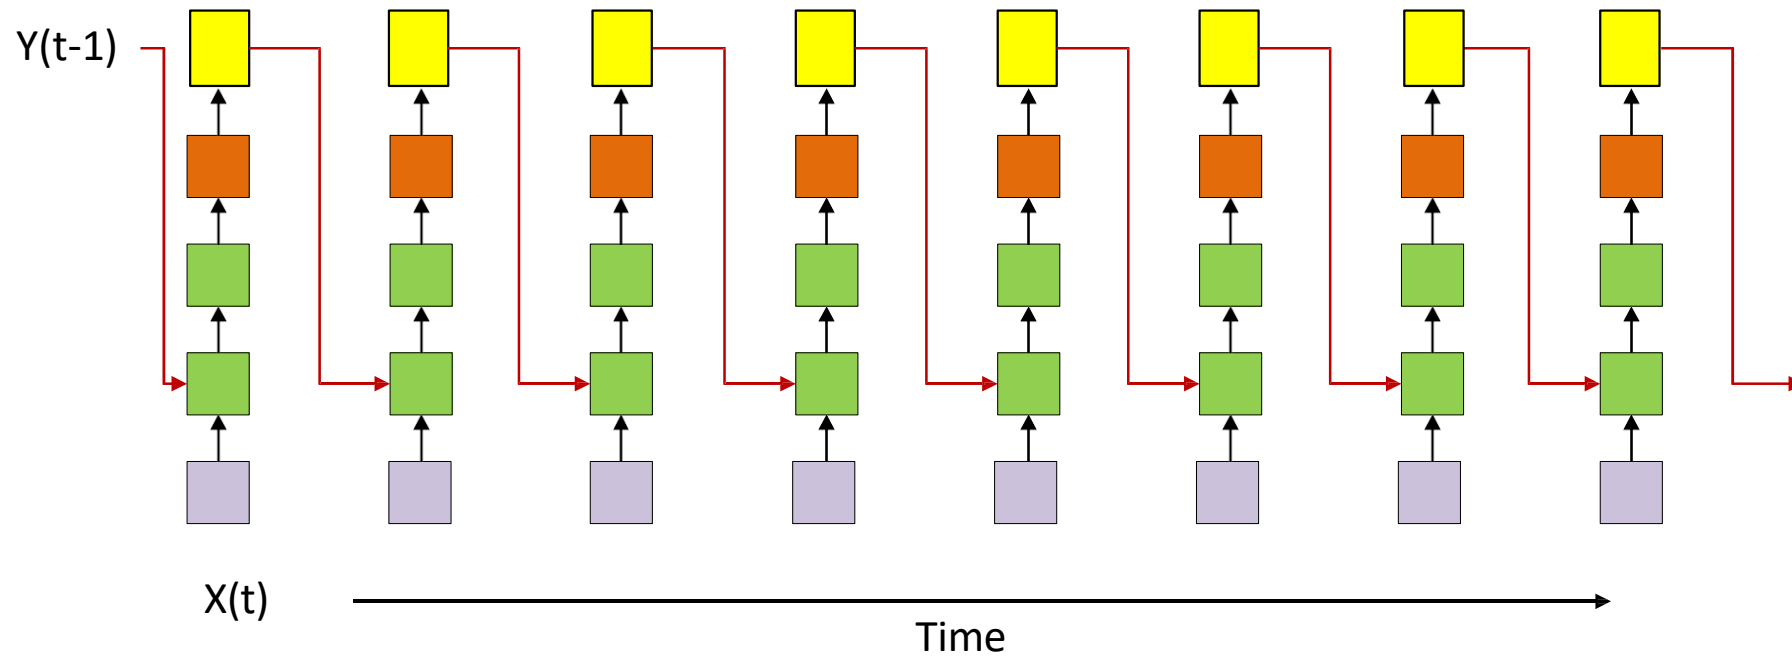


- A NARX net with recursion from the output



- A NARX net with recursion from the output

A more complete representation



- A NARX net with recursion from the output
- Showing all computations
- All columns are identical
- *An input at $t=0$ affects outputs forever*

NARX Networks

- Very popular for time-series prediction
 - Weather
 - Stock markets
 - As alternate system models in tracking systems
 - *Language*
- Note: here the “memory” of the past is in the output itself, and *not in the network*

Let's make memory more explicit

- Task is to “remember” the past
- Introduce an explicit *memory* variable whose *job* it is to remember

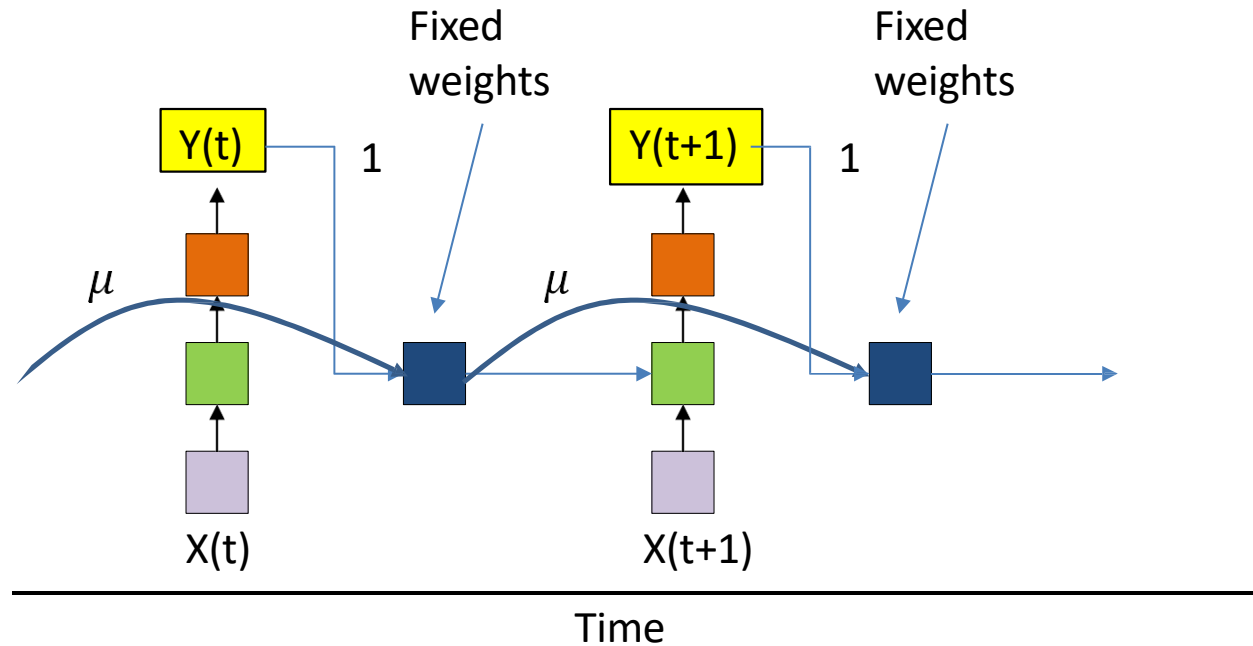
$$m_t = r(y_{t-1}, h_{t-1}, m_{t-1})$$

$$h_t = f(x_t, m_t)$$

$$y_t = g(h_t)$$

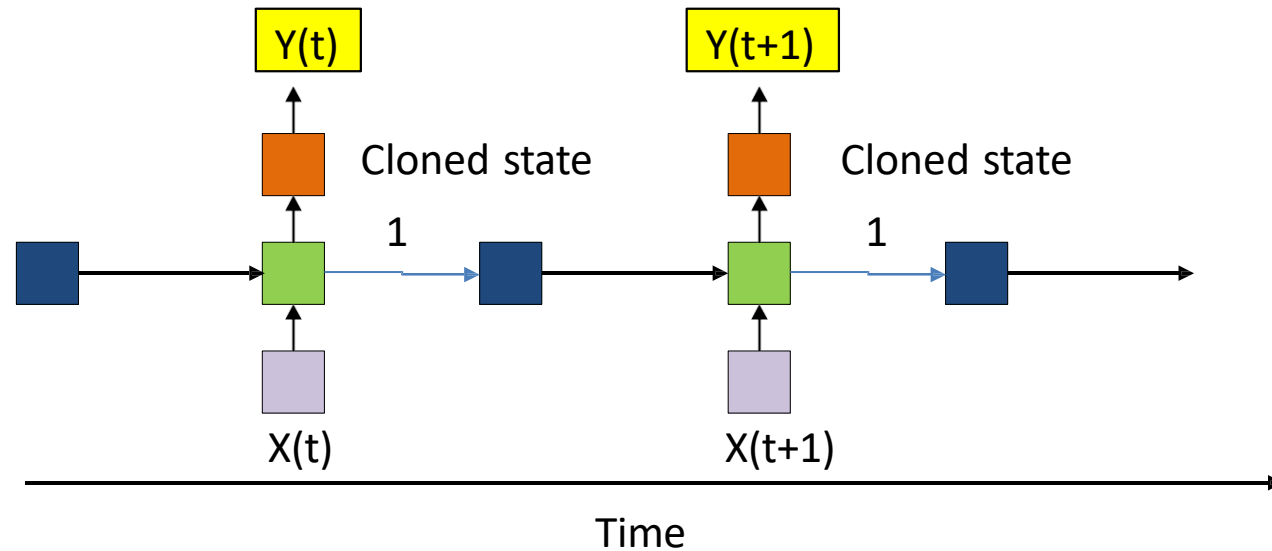
- m_t is a “memory” variable
 - Generally stored in a “memory” unit
 - Used to “remember” the past

Jordan Network



- Memory unit simply retains a running average of past outputs
 - “Serial order: A parallel distributed processing approach”, M.I.Jordan, 1986
 - Memory has fixed structure; does not “learn” to remember
 - The running average of outputs considers entire past, rather than immediate past

Elman Networks



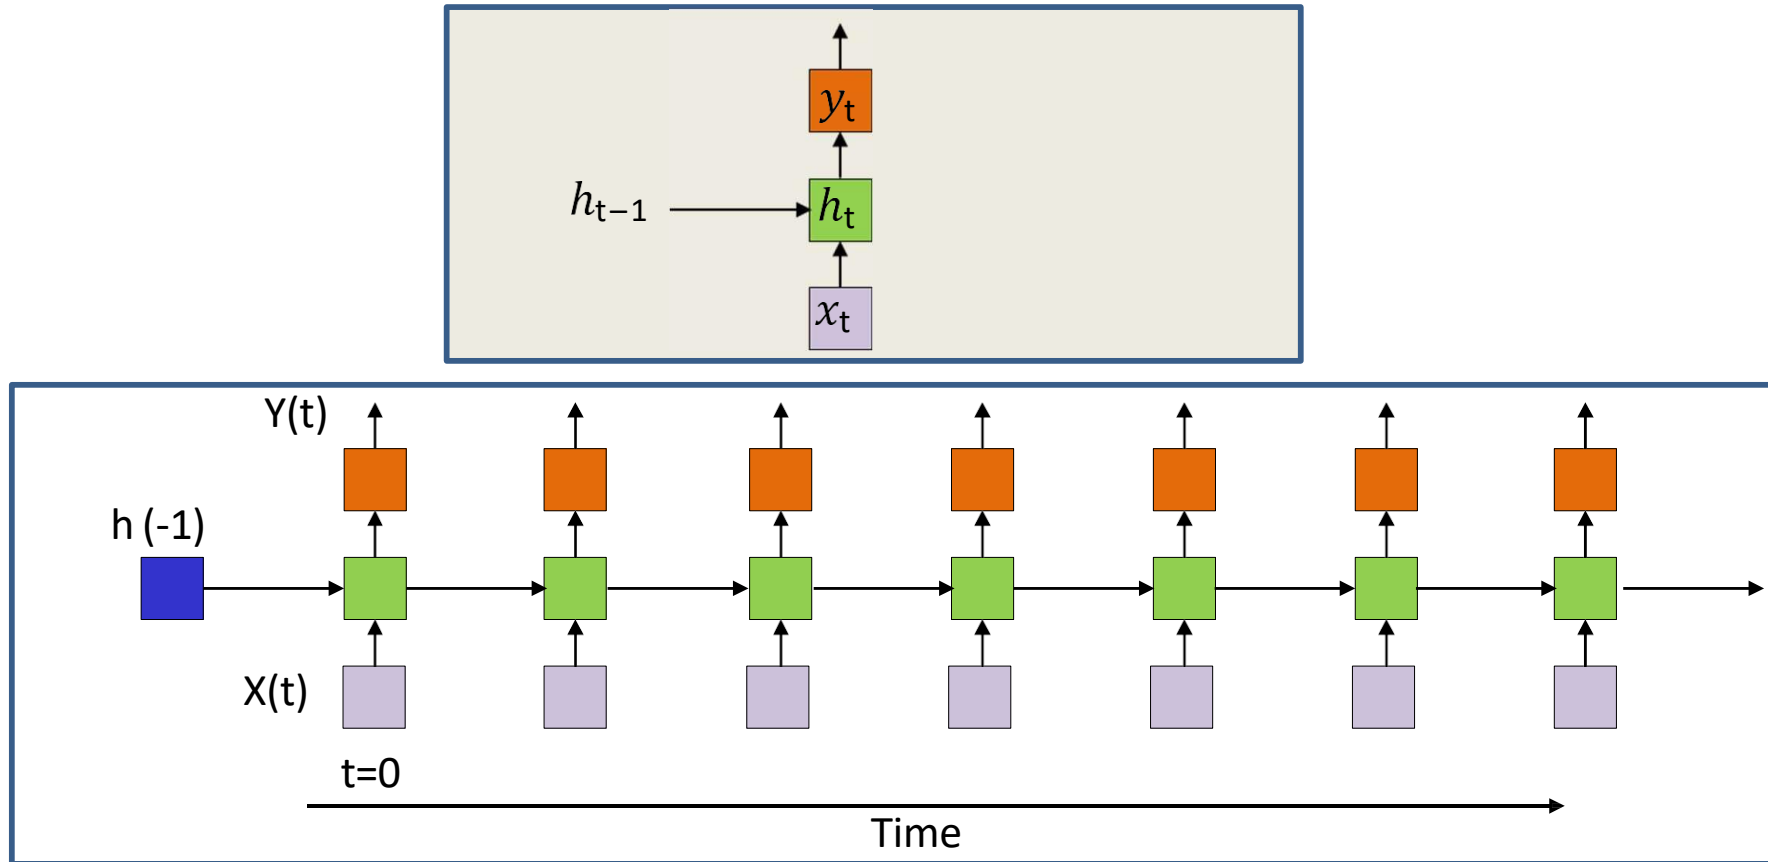
- Separate memory state from output
 - “Context” units that carry historical state
 - “Finding structure in time”, Jeffrey Elman, Cognitive Science, 1990
- Only the weight *from* the memory unit to the hidden unit is learned
 - But during training no gradient is backpropagated over the “1” link

An alternate model for infinite response systems: **the state-space model**

$$\begin{aligned}h_t &= f(x_t, h_{t-1}) \\ y_t &= g(h_t)\end{aligned}$$

- h_t is the *state* of the network
 - *State* summarizes information about the entire past
 - Model directly embeds the memory in the state
- Need to define initial state h_{-1}
- This is a *recurrent* neural network

The simple state-space model




- The state (green) at any time is determined by the input at that time, and the state at the previous time
- *An input at $t=0$ affects outputs forever*
- Also known as a recurrent neural net

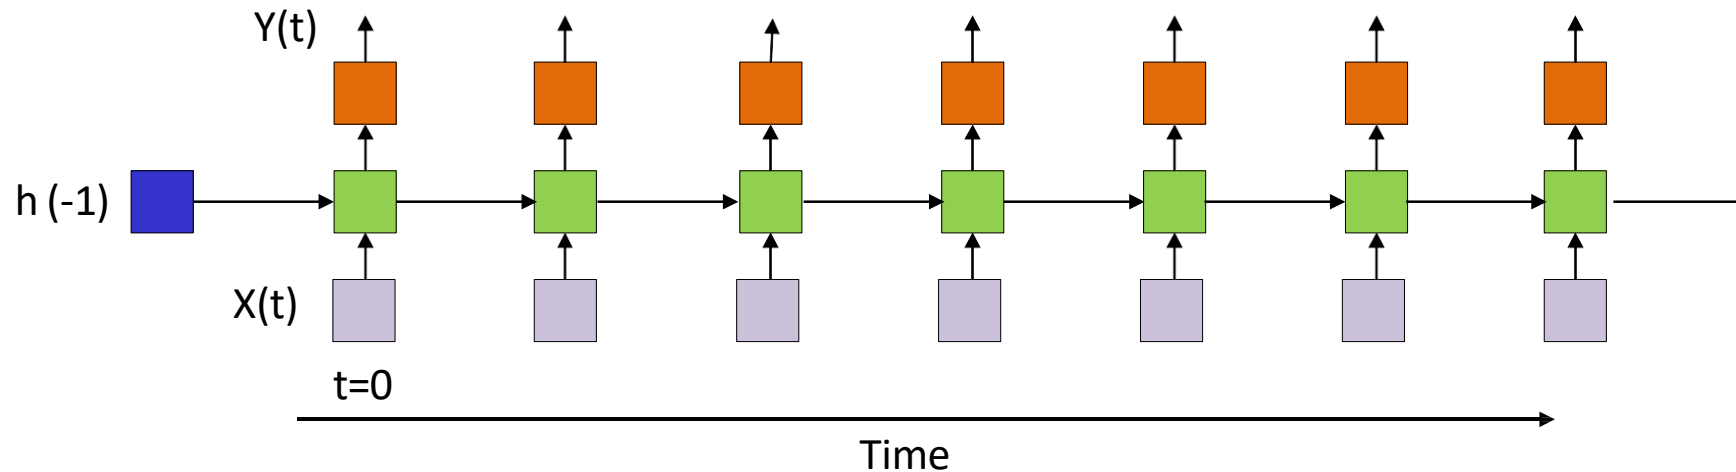
An alternate model for infinite response systems: the state-space model

$$h_t = f(x_t, h_{t-1})$$
$$y_t = g(h_t)$$

- h_t is the *state* of the network
- Need to define initial state h_{-1}
- The state can be *arbitrarily complex*

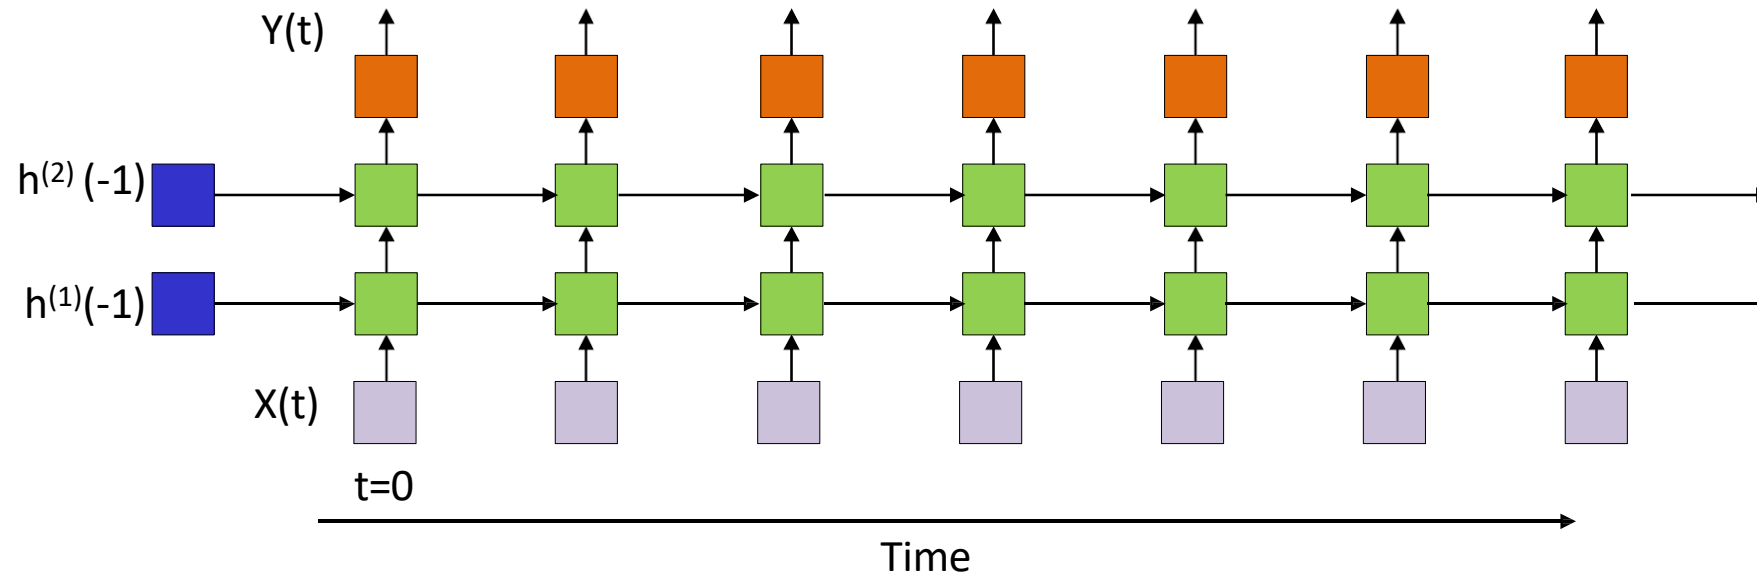


Single hidden layer RNN



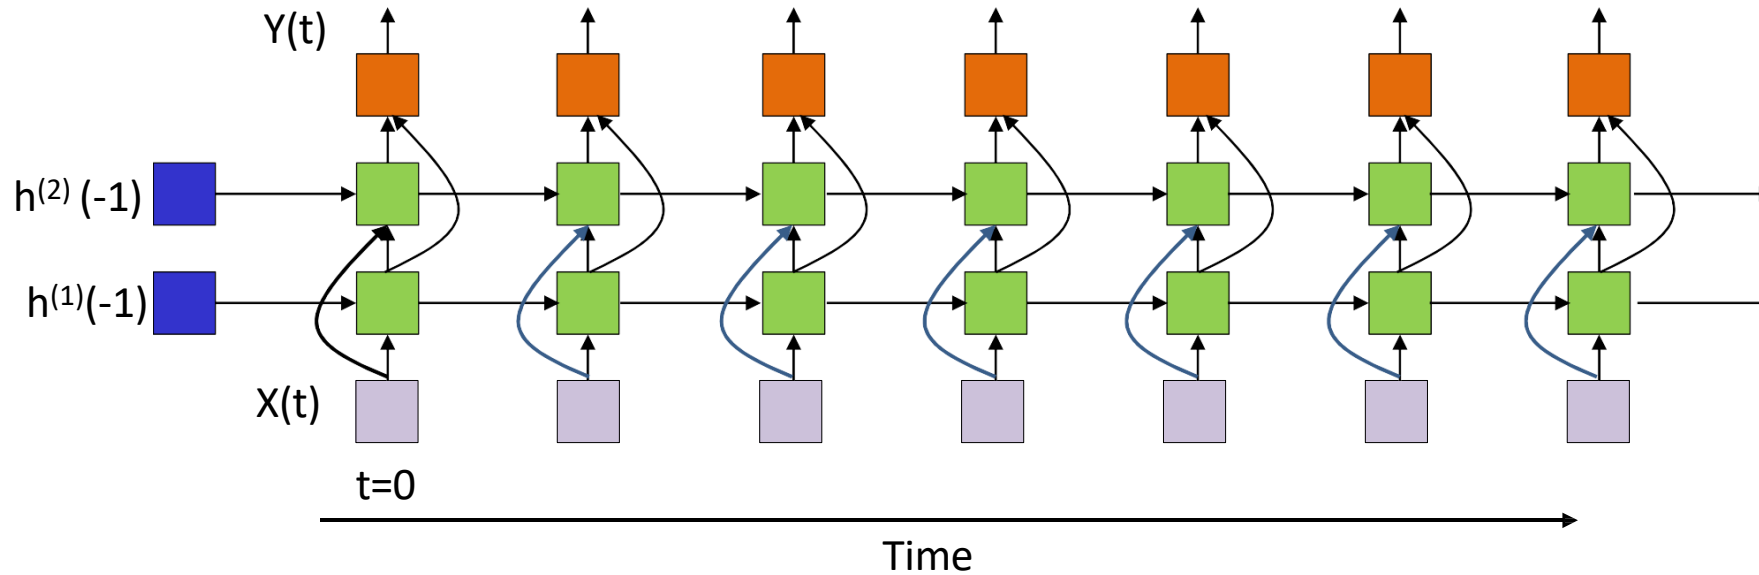
- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Multiple recurrent layer RNN



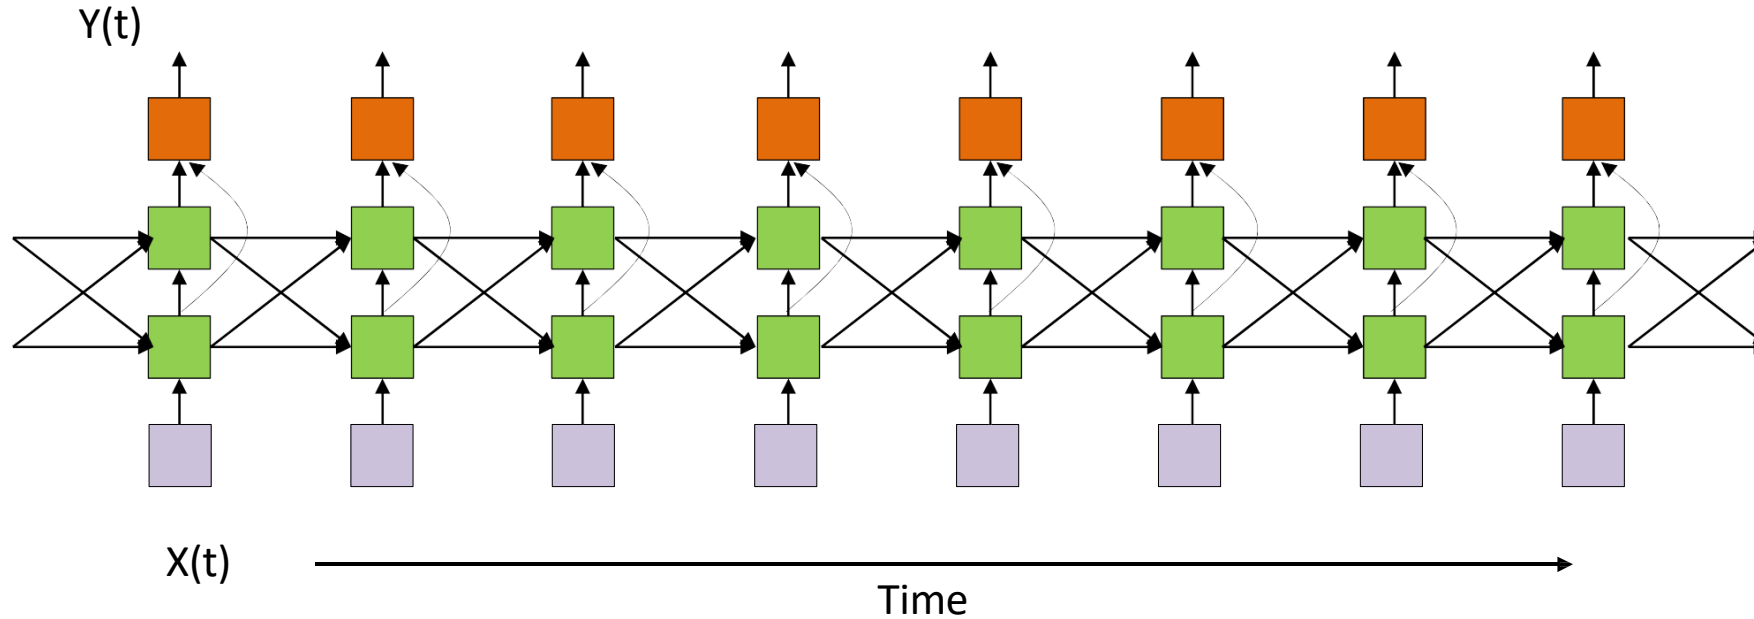
- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Multiple recurrent layer RNN



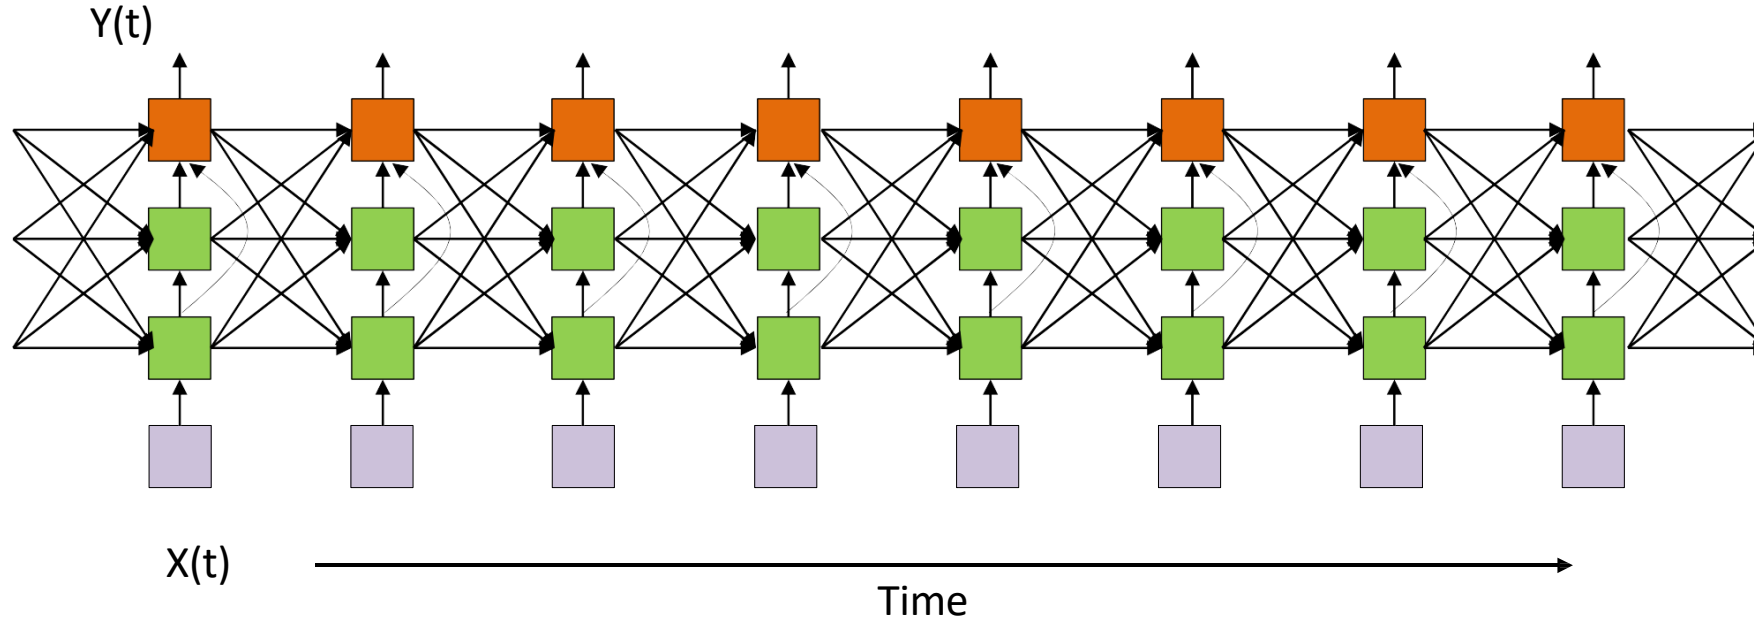
- We can also have skips..

A more complex state



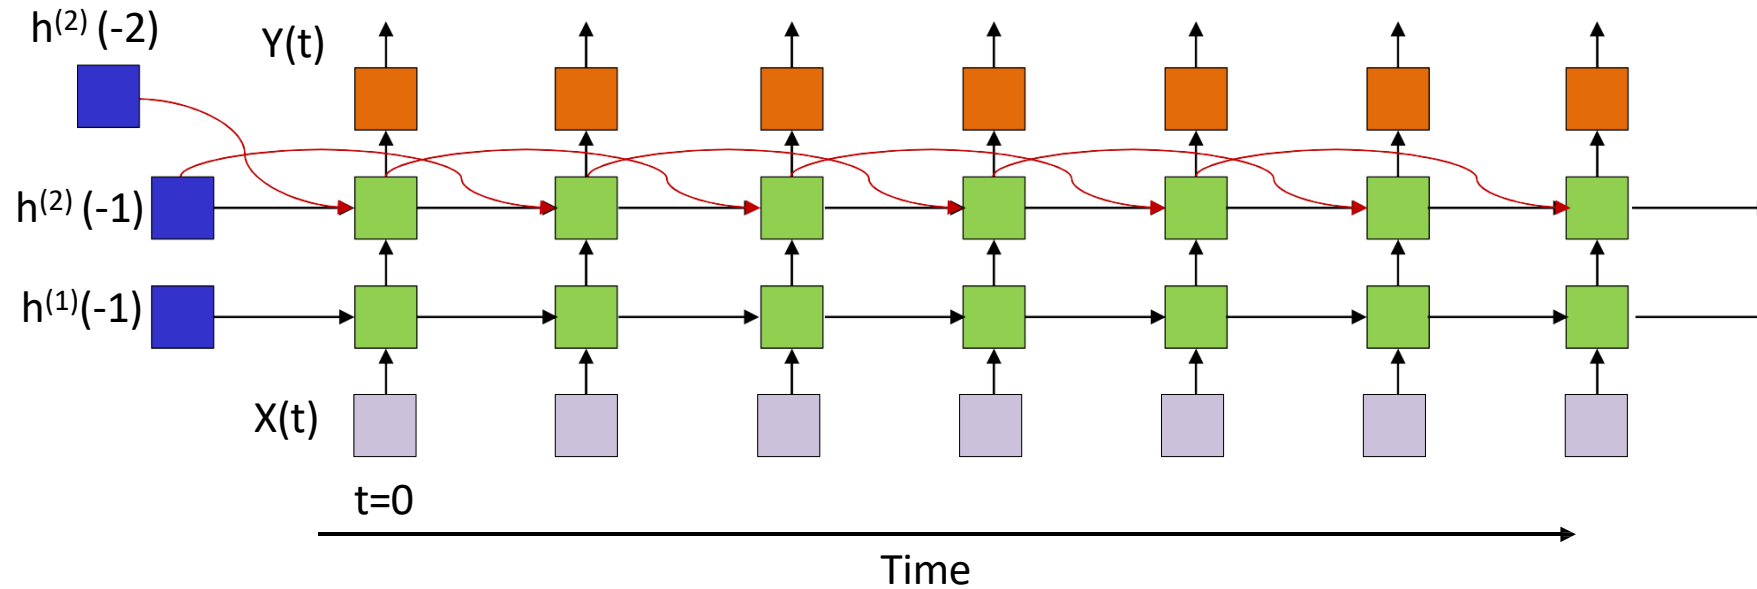
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Or the network may be even more complicated



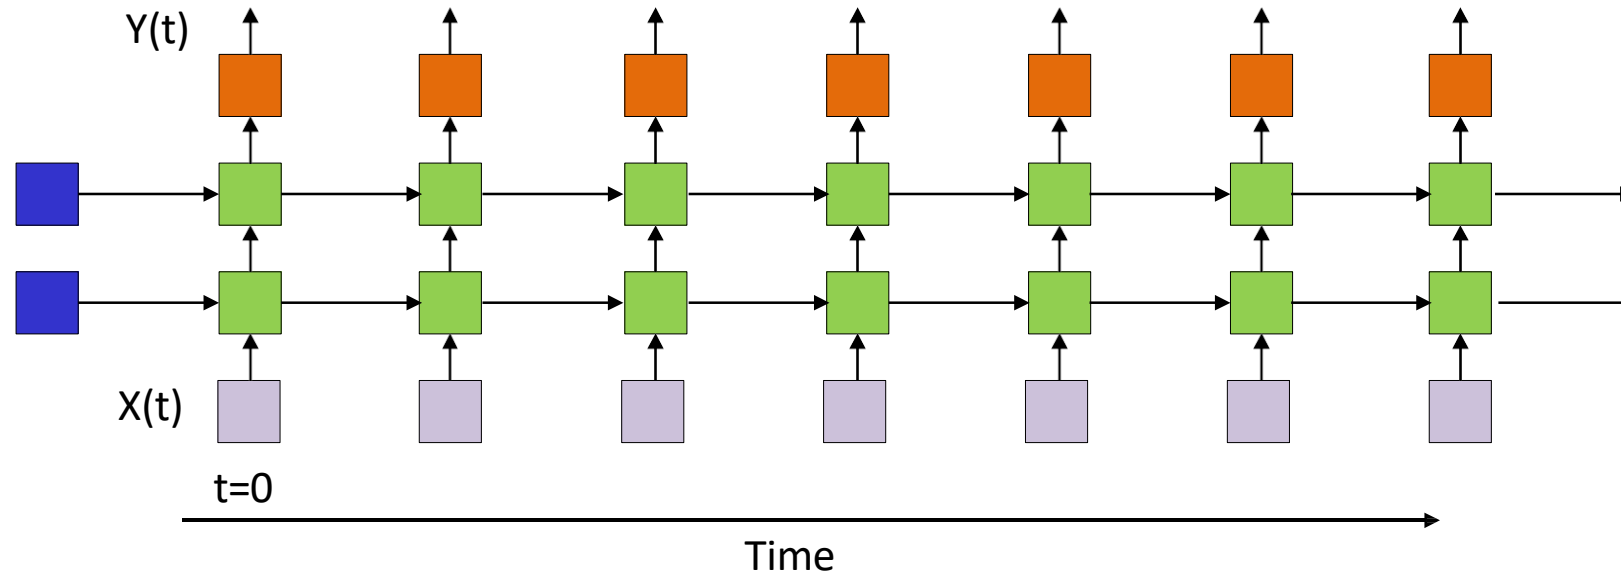
- All columns are identical
- *An input at $t=0$ affects outputs forever*

Generalization with other recurrences




- All columns (including incoming edges) are identical

The simplest structures are most popular




- Recurrent neural network
- All columns are identical
- *An input at $t=0$ affects outputs forever*

A Recurrent Neural Network




- Simplified models often drawn
- The loops imply recurrence


The detailed version of the simplified representation



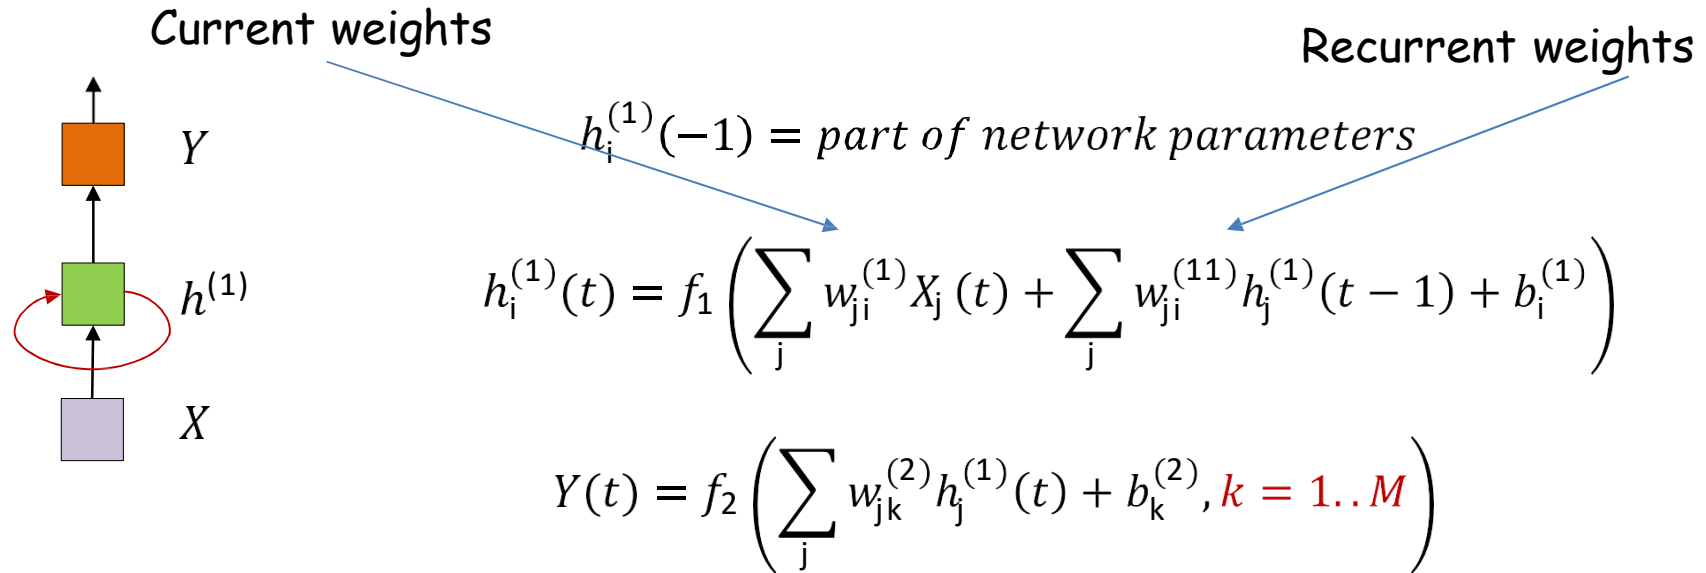
Multiple recurrent layer RNN



Multiple recurrent layer RNN

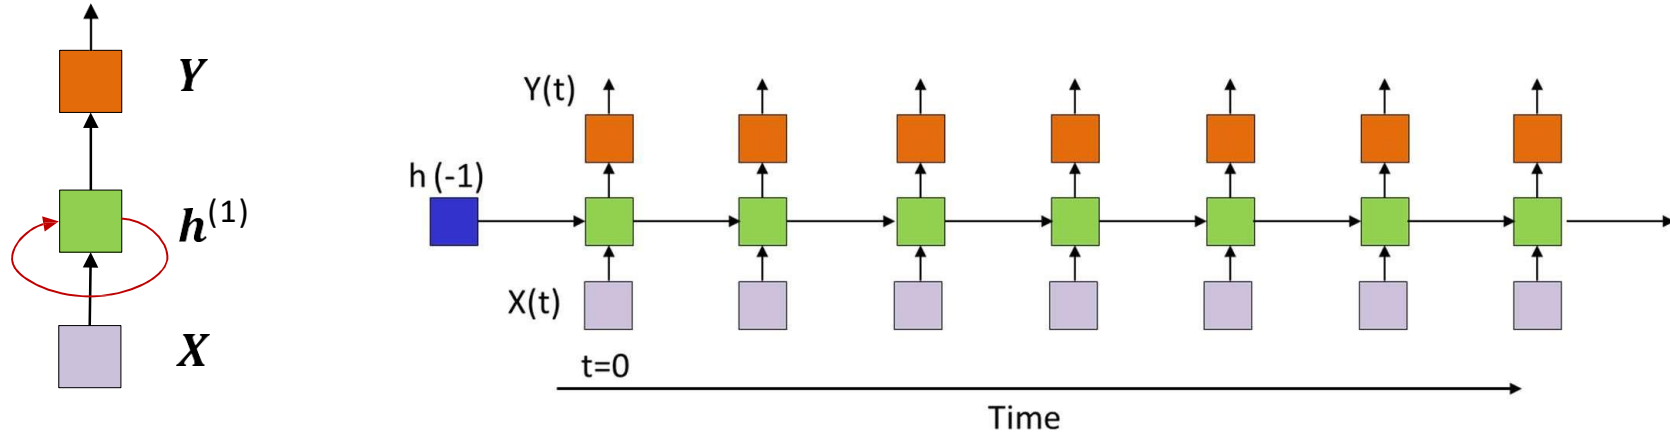


Equations



- Note superscript in indexing, which indicates layer of network from which inputs are obtained
- Assuming vector function at output, e.g. softmax
- The *state* node activation, $f_1()$ is typically $\tanh()$
- Every neuron also has a *bias* input

Equations



$h^{(1)}(-1) = \text{part of network parameters}$


- Computation:

$$h^{(1)}(t) = f_1(W^{(1)}X(t) + W^{(11)}h^{(1)}(t-1) + b^{(1)})$$

$$Y(t) = f_2(W^{(2)}h^{(1)}(t) + b^{(2)})$$

- The recurrent state activation $f_1()$ is typically $\tanh()$

Equations



$h_i^{(1)}(-1) = \text{part of network parameters}$

$h_i^{(2)}(-1) = \text{part of network parameters}$


$$h_i^{(1)}(t) = f_1 \left(\sum_j w_{ji}^{(1)} X_j(t) + \sum_j w_{ji}^{(11)} h_j^{(1)}(t-1) + b_i^{(1)} \right)$$

$$h_i^{(2)}(t) = f_2 \left(\sum_j w_{ji}^{(2)} h_j^{(1)}(t) + \sum_j w_{ji}^{(22)} h_j^{(2)}(t-1) + b_i^{(2)} \right)$$

$$Y(t) = f_3 \left(\sum_j w_{jk}^{(3)} h_j^{(2)}(t) + b_k^{(3)}, k = 1..M \right)$$

- Assuming vector function at output, e.g. softmax $f_3()$
- The *state* node activations, $f_k()$ are typically $\tanh()$
- Every neuron also has a *bias* input

Equations



$h^{(1)}(-1)$ and $h^{(2)}(-1)$ = part of network parameters

- Computation:

$$h^{(1)}(t) = f_1(W^{(1)}X(t) + W^{(11)}h^{(1)}(t-1) + b^{(1)})$$

$$h^{(2)}(t) = f_2(W^{(2)}h^{(1)}(t) + W^{(22)}h^{(2)}(t-1) + b^{(2)})$$


$$Y(t) = f_3(W^{(3)}h^{(2)}(t) + b^{(3)})$$

- The recurrent state activation is typically $\tanh()$

Equations

$\mathbf{h}^{(1)}(-1) = \text{part of network parameters}$

$\mathbf{h}^{(2)}(-1) = \text{part of network parameters}$




$$\mathbf{h}^{(1)}(t) = f_1(\mathbf{W}^{(01)} \mathbf{X}(t) + \mathbf{W}^{(11)} \mathbf{h}^{(1)}(t-1) + \mathbf{b}^{(1)})$$

$$\mathbf{h}^{(2)}(t) = f_2(\mathbf{W}^{(12)} \mathbf{h}^{(1)}(t) + \mathbf{W}^{(02)} \mathbf{X}(t) + \mathbf{W}^{(22)} \mathbf{h}^{(2)}(t-1) + \mathbf{b}^{(2)})$$


$$\mathbf{Y}(t) = f_3(\mathbf{W}^{(23)} \mathbf{h}^{(2)}(t) + \mathbf{W}^{(13)} \mathbf{h}^{(1)}(t) + \mathbf{b}^{(3)})$$

Variants on recurrent nets


one to one



one to many



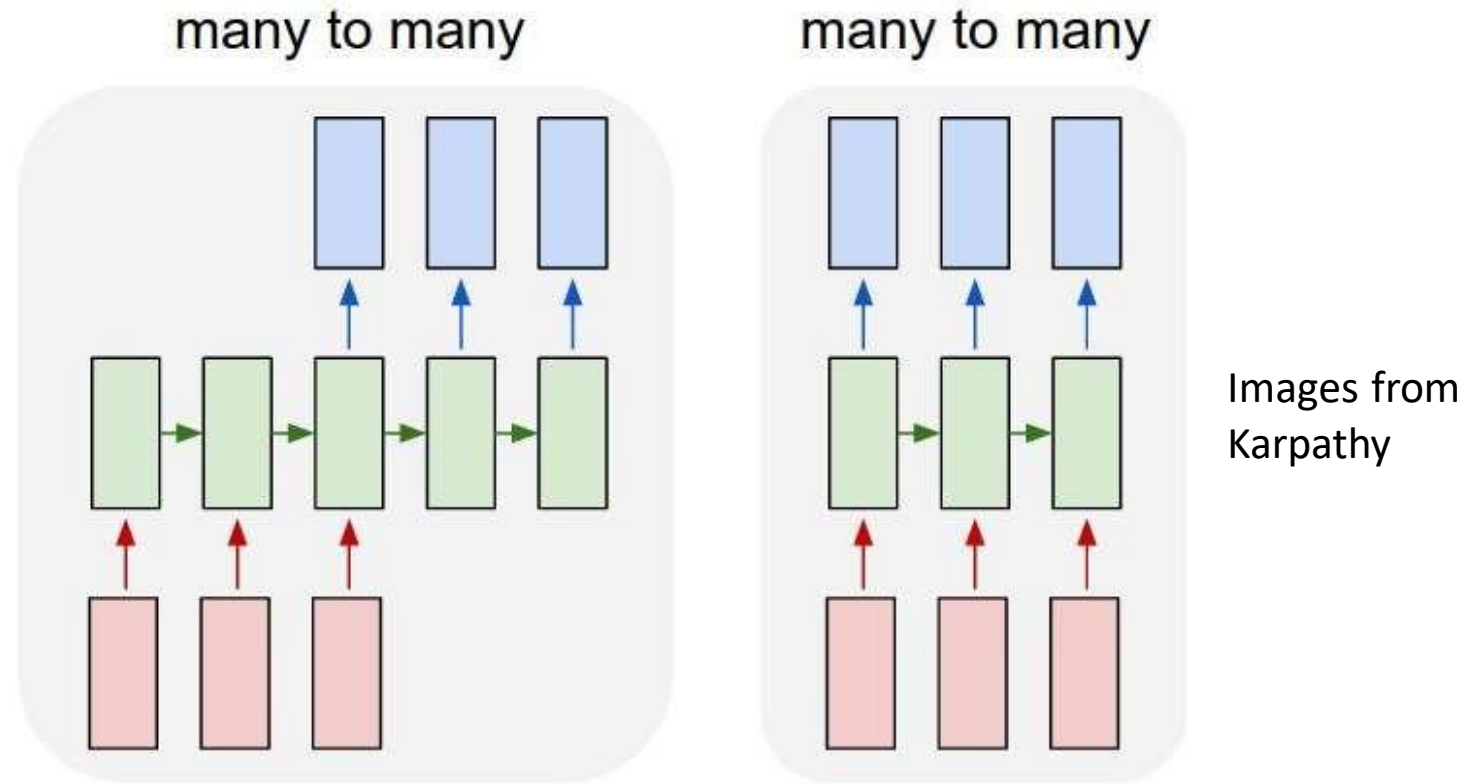
many to one



Images from
Karpathy

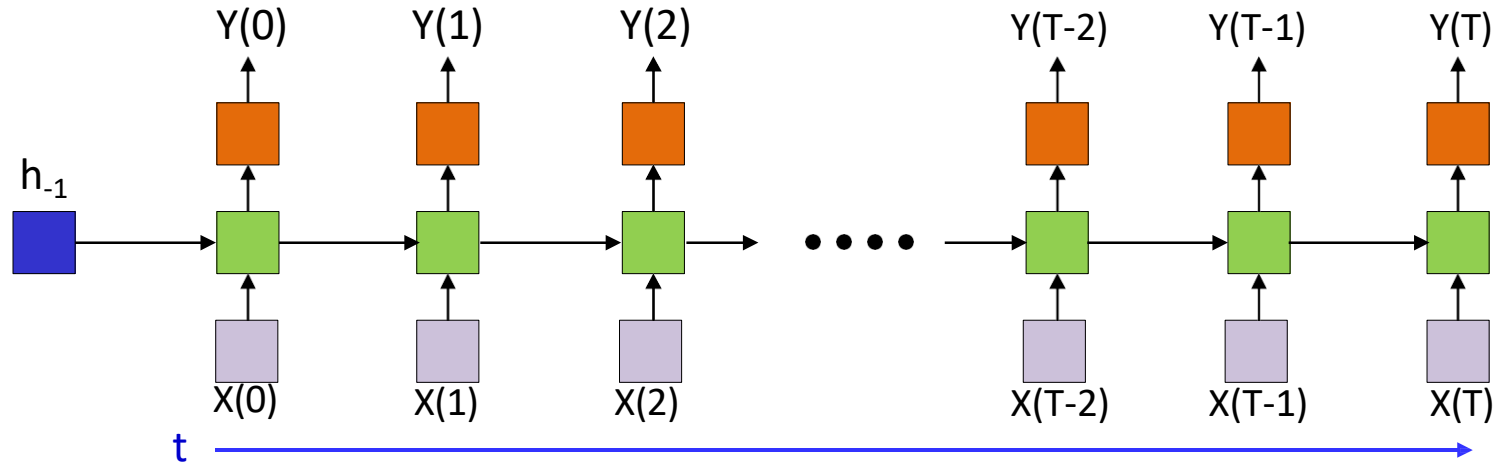
- 1: Conventional MLP
- 2: Sequence *generation*, e.g. image to caption
- 3: Sequence based *prediction or classification*, e.g. Speech recognition, text classification

Variants



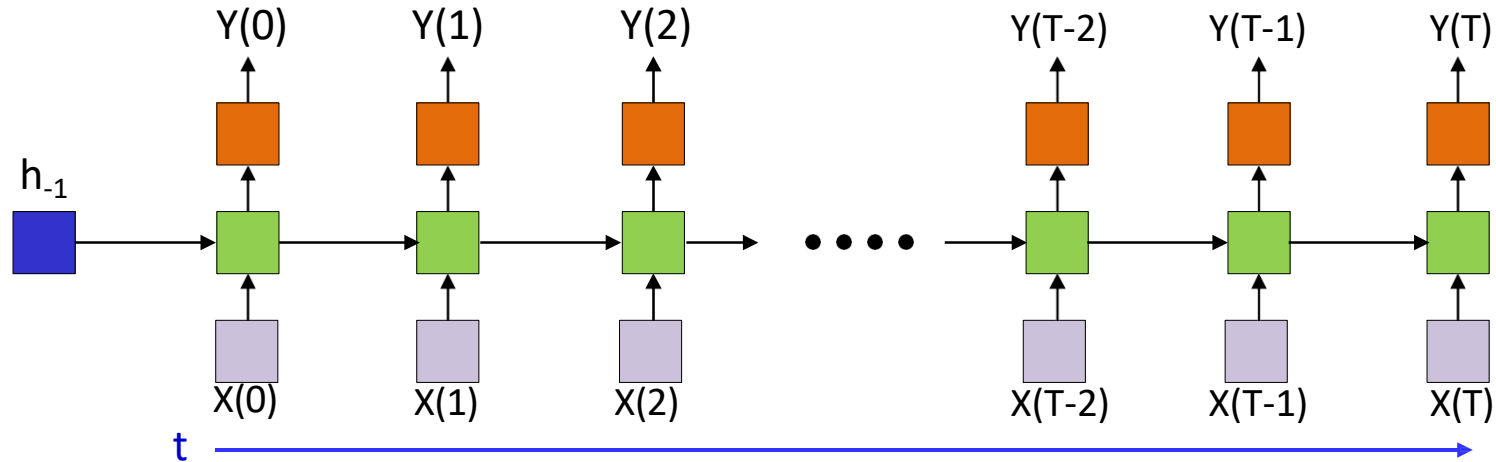
- 1: *Delayed* sequence to sequence, e.g. machine translation
- 2: Sequence to sequence, e.g. stock problem, label prediction
- Etc...

How do we *train* the network



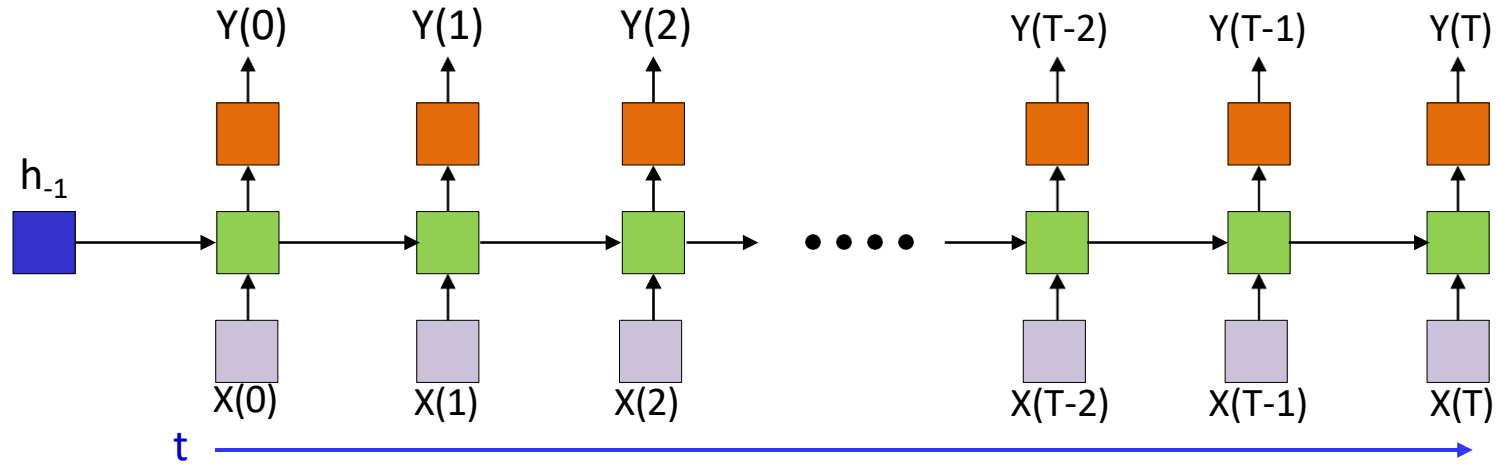
- Back propagation through time (BPTT)
- Given a collection of *sequence* inputs
 - $(\mathbf{X}_i, \mathbf{D}_i)$
 - $\mathbf{X}_i = X_{i,0}, \dots, X_{i,T}$
 - $\mathbf{D}_i = D_{i,0}, \dots, D_{i,T}$
- Train network parameters to minimize the error between the output of the network $\mathbf{Y}_i = Y_{i,0}, \dots, Y_{i,T}$ and the desired outputs
 - This is the most generic setting.

Training the RNN



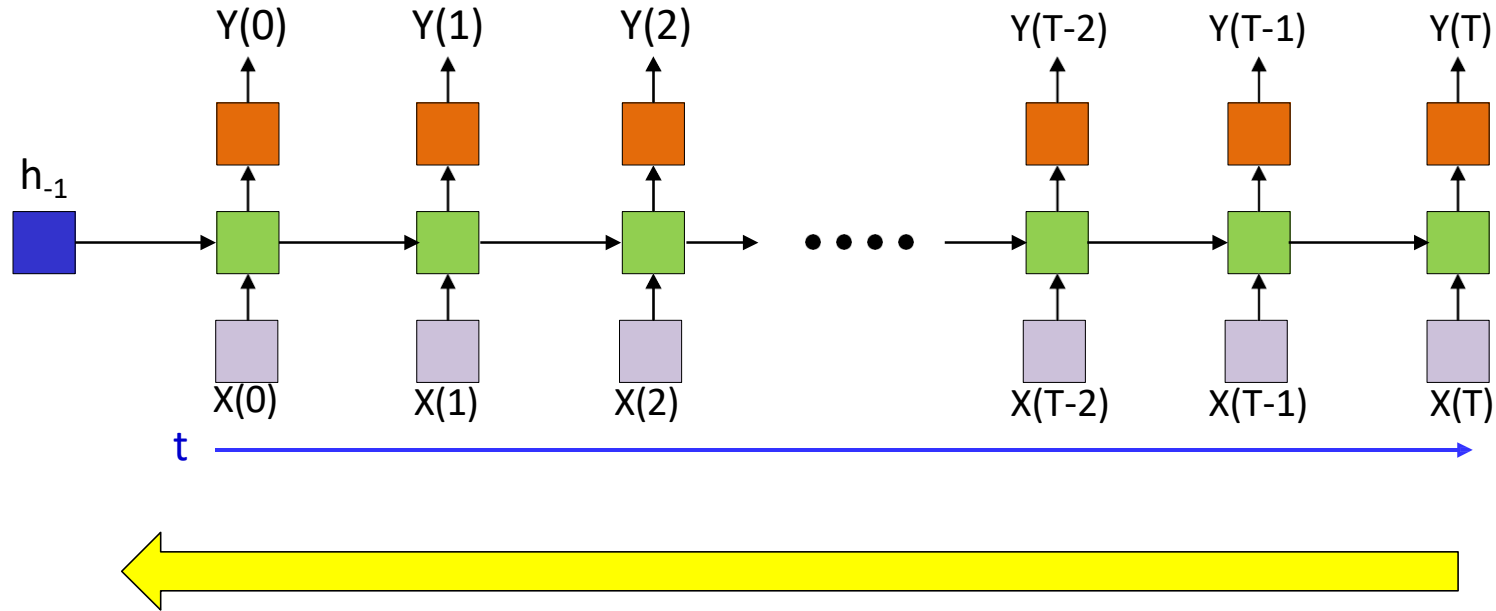
- The “unrolled” computation is just a giant shared-parameter neural network
 - All columns are identical and share parameters
- Network parameters can be trained via gradient-descent (or its variants) using shared-parameter gradient descent rules
 - Gradient computation requires a forward pass, back propagation, and pooling of gradients (for parameter sharing)

Training: Forward pass




- For each training input:
- Forward pass: pass the entire data sequence through the network, generate outputs

Training: Computing gradients



- For each training input:
- **Backward pass: Compute gradients via backpropagation**
 - *Back Propagation Through Time*


Back Propagation Through Time



Will only focus on *one* training instance


All subscripts represent *components* and not training instance index

Back Propagation Through Time




- The divergence computed is between the *sequence of outputs* by the network and the *desired sequence of outputs*
 - DIV is a scalar function of a series of vectors!
- This is *not* just the sum of the divergences at individual times
 - Unless we explicitly define it that way

Notation




- $Y(t)$ is the output at time t
 - $Y_i(t)$ is the i th output
- $Z^{(2)}(t)$ is the pre-activation value of the neurons at the output layer at time t
- $h(t)$ is the output of the hidden layer at time t
 - Assuming only one hidden layer in this example
- $Z^{(1)}(t)$ is the pre-activation value of the hidden layer at time t

Notation



- $W^{(1)} = [w_{ij}^{(1)}]$ is the matrix of *current* weights from the input to the hidden layer.
- $W^{(2)} = [w_{ij}^{(2)}]$ is the matrix of *current* weights from the hidden layer to the output layer
- $W^{(11)} = [w_{ij}^{(11)}]$ is the matrix of *recurrent* weights from the hidden layer to itself

Back Propagation Through Time




First step of backprop: Compute $\nabla_{Y(T)} DIV$ (Compute $\frac{dDIV}{dY_i(T)}$ for all i)

Note: DIV is a function of *all* outputs $Y(0) \dots Y(T)$

In general we will be required to compute $\frac{dDIV}{dY_i(t)}$ for all i and t . This can be a source of significant difficulty in many scenarios.

Back Propagation Through Time




Special case, when the overall divergence is a simple sum of local divergences at each time: $DIV = \sum_t Div(t)$

Will get $\nabla_{Y(t)} Div(t)$

$$\frac{\partial DIV}{\partial Y_i(t)} = \frac{\partial Div(t)}{\partial Y_i(t)}$$

Back Propagation Through Time




$$\nabla_{Z^{(2)}(T)} DIV = \nabla_{Y(T)} DIV \nabla_{Z^{(2)}(T)} Y(T)$$

Vector output activation


$$\frac{dDIV}{dZ_i^{(2)}(T)} = \frac{dDIV}{dY_i(T)} \frac{dY_i(T)}{dZ_i^{(2)}(T)} \quad \text{OR} \quad \frac{dDIV}{dZ_i^{(2)}(T)} = \sum_j \frac{dDIV}{dY_j(T)} \frac{dY_j(T)}{dZ_i^{(2)}(T)}$$

Back Propagation Through Time




$$\frac{dDIV}{dh_i(T)} = \sum_j \frac{dDIV}{dZ_j^{(2)}(T)} \frac{dZ_j^{(2)}(T)}{dh_i(T)} = \sum_j w_{ij}^{(2)} \frac{dDIV}{dZ_j^{(2)}(T)}$$

Back Propagation Through Time




$$\frac{dDIV}{dw_{ij}^{(2)}} = \frac{dDIV}{dZ_j^{(2)}(T)} h_i(T)$$

Back Propagation Through Time




$$\frac{dDIV}{dZ_i^{(1)}(T)} = \frac{dDIV}{dh_i(T)} \frac{dh_i(T)}{dZ_i^{(1)}(T)}$$

Back Propagation Through Time




$$\frac{dDIV}{dw_{ij}^{(1)}} = \frac{dDIV}{dZ_j^{(1)}(T)} X_i(T)$$

Back Propagation Through Time



$$\frac{dDIV}{dw_{ij}^{(11)}} = \frac{dDIV}{dZ_j^{(1)}(T)} h_i(T-1)$$


Back Propagation Through Time



Vector output activation

$$\frac{dDIV}{dZ_i^{(2)}(T-1)} = \frac{dDIV}{dY_i(T-1)} \frac{dY_i(T-1)}{dZ_i^{(2)}(T-1)} \quad \text{OR} \quad \frac{dDIV}{dZ_i^{(2)}(T-1)} = \sum_j \frac{dDIV}{dY_j(T-1)} \frac{dY_j(T-1)}{dZ_i^{(2)}(T-1)}$$

Back Propagation Through Time




$$\frac{dDIV}{dh_i(T-1)} = \sum_j w_{ij}^{(2)} \frac{dDIV}{dZ_j^{(2)}(T-1)} + \sum_j w_{ij}^{(11)} \frac{dDIV}{dZ_j^{(1)}(T)}$$

Note the addition $\Rightarrow \frac{dDIV}{dw_{ij}^{(2)}} += \frac{dDIV}{dZ_j^{(2)}(T-1)} h_i(T-1)$


Note the addition $\Rightarrow \nabla_{W^{(2)}} DIV += h(T-1) \nabla_{Z^{(2)}(T-1)} DIV$

Back Propagation Through Time



$$\frac{dDIV}{dZ_i^{(1)}(T-1)} = \frac{dDIV}{dh_i(T-1)} \frac{dh_i(T-1)}{dZ_i^{(1)}(T-1)}$$

Back Propagation Through Time




$$\frac{dDIV}{dZ_i^{(1)}(T-1)} = \frac{dDIV}{dh_i(T-1)} \frac{dh_i(T-1)}{dZ_i^{(1)}(T-1)}$$

$$\frac{dDIV}{dw_{ij}^{(1)}} += \frac{dDIV}{dZ_j^{(1)}(T-1)} X_i(T-1)$$


Note the addition

Back Propagation Through Time



Note the addition $\Rightarrow \frac{dDIV}{dw_{ij}^{(11)}} += \frac{dDIV}{dz_j^{(1)}(T-1)} h_i(T-2)$


Back Propagation Through Time



Continue computing derivatives going backward through time until..

$$\frac{dDIV}{dh_i(-1)} = \sum_j w_{ij}^{(11)} \frac{dDIV}{dZ_j^{(1)}(0)}$$

Back Propagation Through Time



Initialize all derivatives to 0

For $t = T$ downto 0


$$\begin{aligned}\nabla_{Z^{(2)}(t)} DIV &= \nabla_{F(t)} DIV \nabla_{Z^{(2)}(t)} Y(t) \\ \nabla_{h(t)} DIV &= \nabla_{Z^{(2)}(t)} DIV W^{(2)} + \nabla_{Z^{(1)}(t+1)} DIV W^{(11)} \\ \nabla_{Z^{(1)}(t)} DIV &= \nabla_{h(t)} DIV \nabla_{Z^{(1)}(t)} h(t)\end{aligned}$$

$$\begin{aligned}\nabla_{W^{(2)}} DIV &+= h(t) \nabla_{Z^{(2)}(t)} DIV \\ \nabla_{W^{(11)}} DIV &+= h(t-1) \nabla_{Z^{(1)}(t)} DIV \\ \nabla_{W^{(1)}} DIV &+= X(t) \nabla_{Z^{(1)}(t)} DIV\end{aligned}$$

$$\nabla_{h_{-1}} DIV = \nabla_{Z^{(1)}(0)} DIV W^{(11)}$$

Extensions to the RNN: *Bidirectional RNN*

Bidirectional RNN (BRNN)



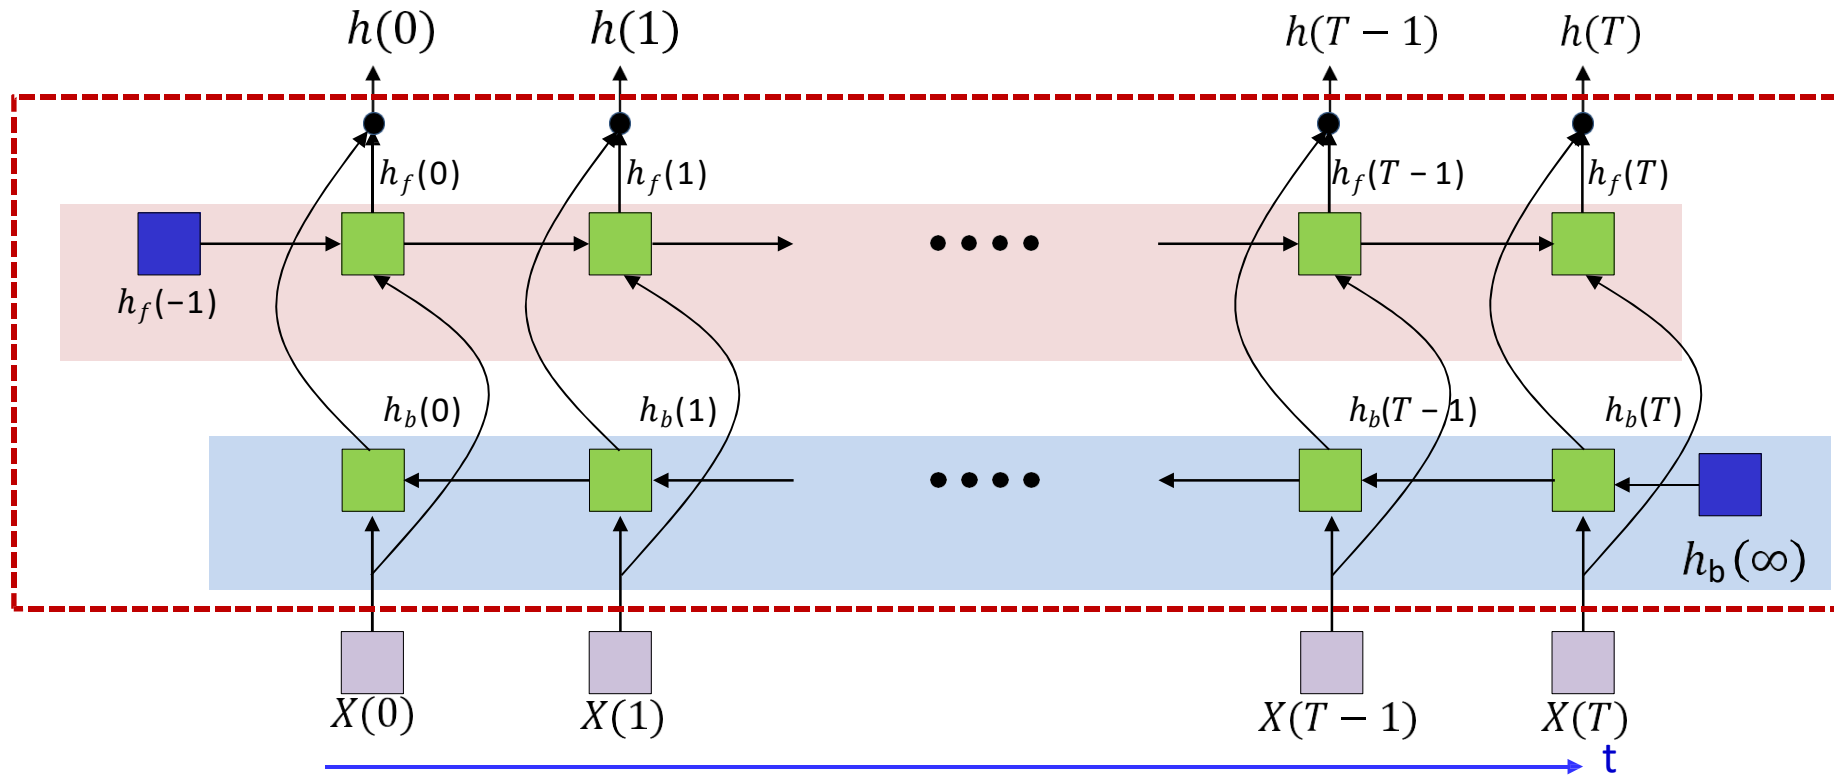
Must learn weights w_2 , w_3 , w_4 & w_5 ; in addition to w_1 & w_6 .

Proposed by Schuster and Paliwal 1997

Alex Graves, "[Supervised Sequence Labelling with Recurrent Neural Networks](#)"


- In problems where the entire input sequence is available, RNNs can be bidirectional
- RNN with both forward and backward recursion
- Explicitly models the fact that just as the future can be predicted from the past, the past can be deduced from the future

Bidirectional RNN




- “Block” performs bidirectional inference on input
 - “Input” could be input series $X(0) \dots X(T)$ or the output of a previous layer (or block)
- The Block has two components
 - A forward net process the data from $t=0$ to $t=T$
 - A backward net processes it backward from $t=T$ down to $t=0$

Bidirectional RNN block



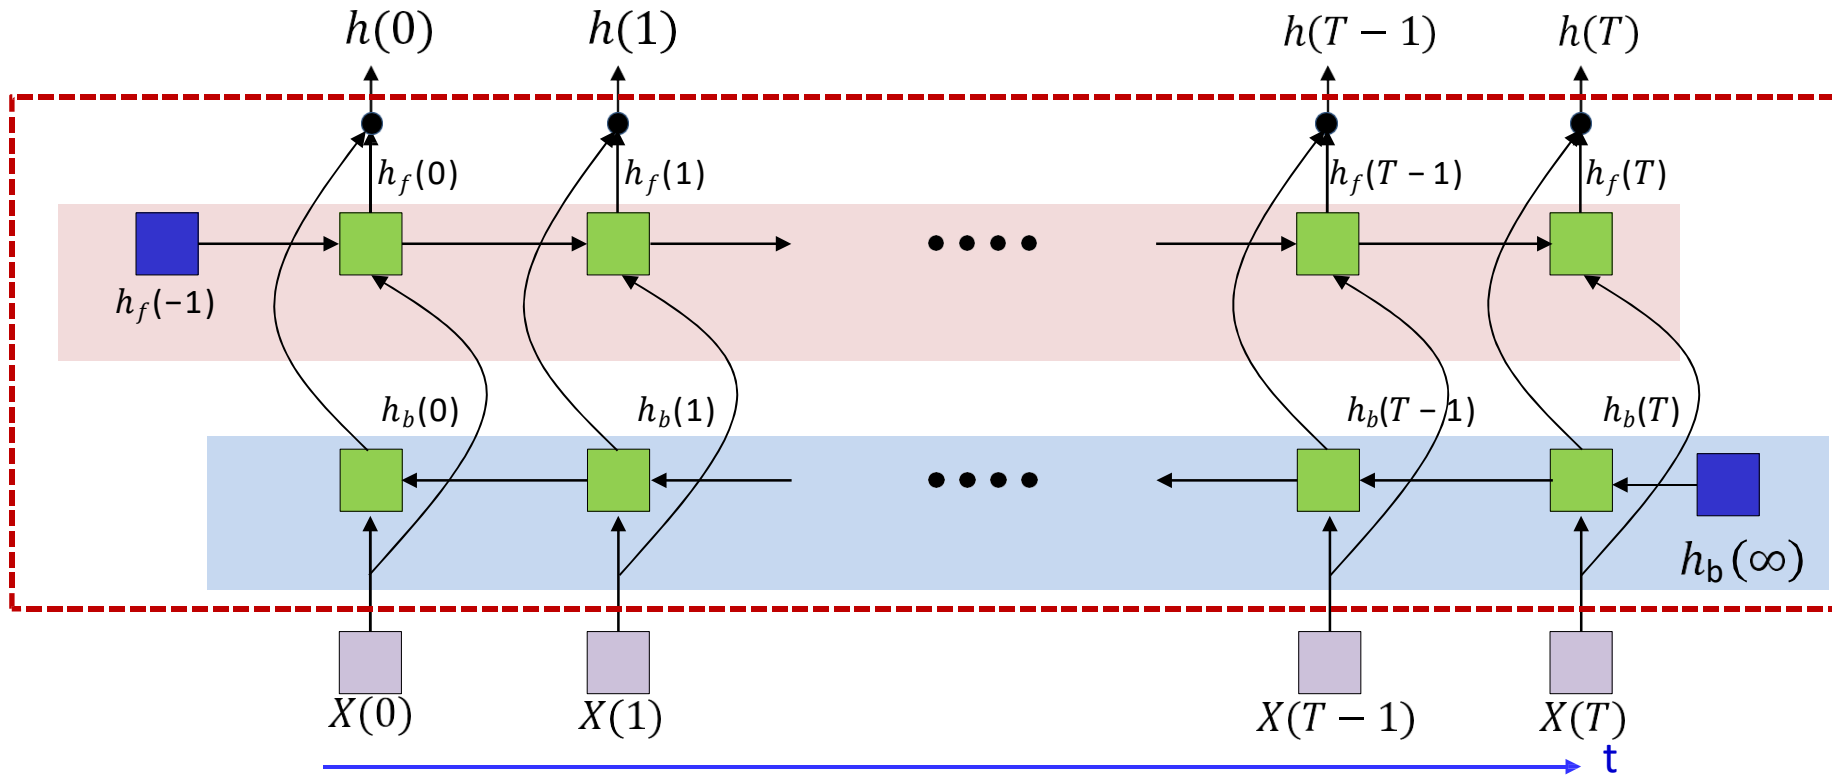
- The forward net process the data from $t=0$ to $t=T$
 - Only computing the hidden state values.

Bidirectional RNN block



- The backward nets processes the input data in *reverse time*, end to beginning
 - Initially only the hidden state values are computed
 - Clearly, this is not an online process and requires the *entire* input data
 - Note: *This is not the backward pass of backprop.*

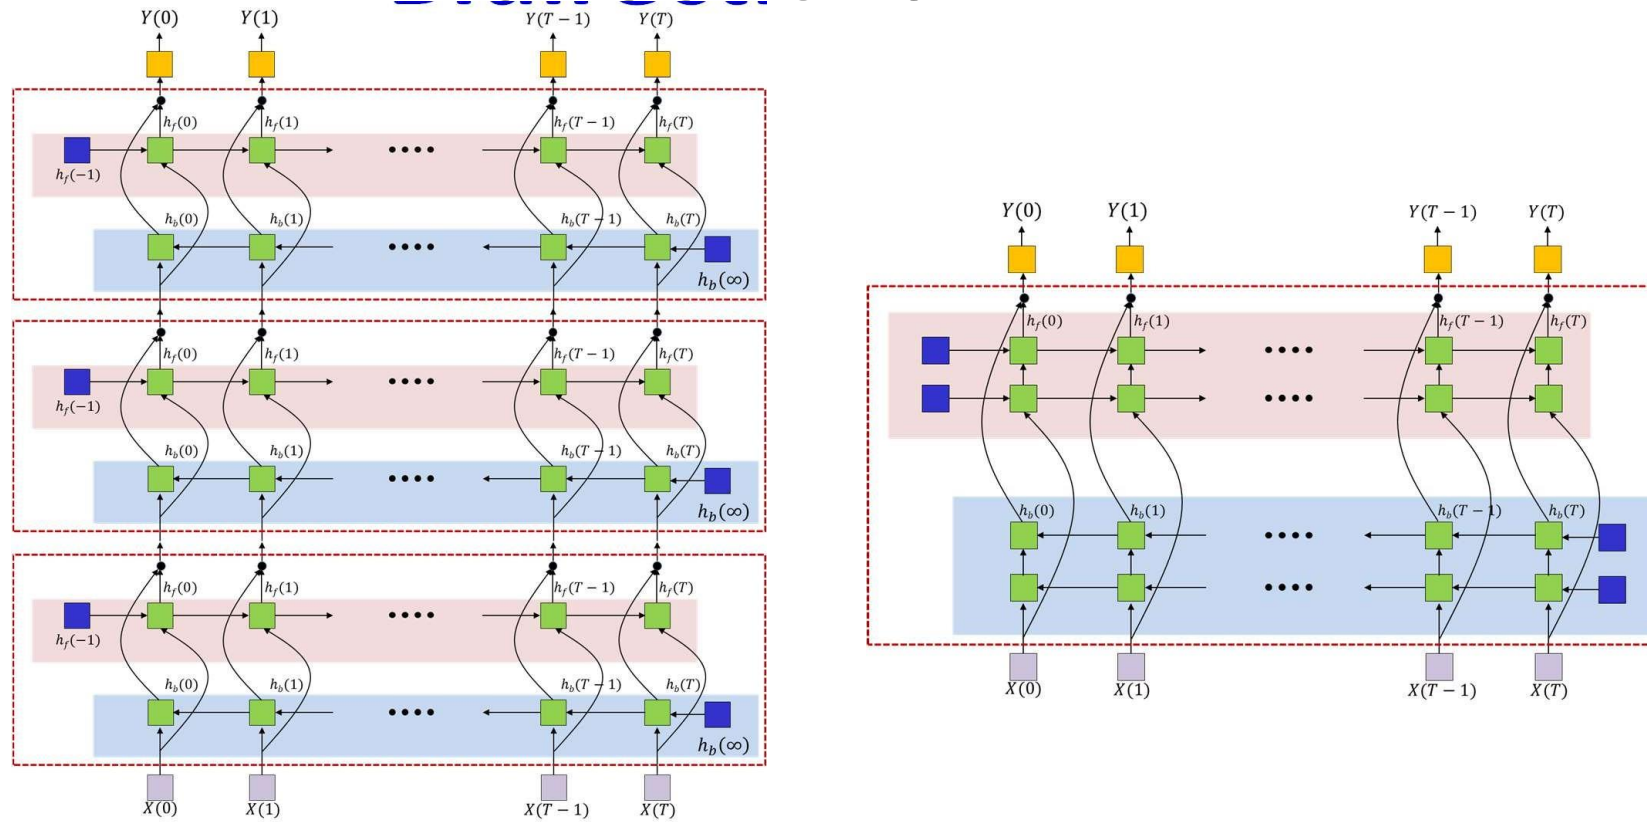
Bidirectional RNN block



- The computed states of both networks are combined to give you the output of the bidirectional block
 - Typically just concatenate them

$$h(t) = [h_f(t); h_b(t)]$$

Bidirectional RNN



- Actual network may be formed by stacking many independent bidirectional blocks followed by an output layer
 - Forward and backward nets in each block are a single layer
- Or by a single bidirectional block followed by an output layer
 - The forward and backward nets may have several layers
- In either case, it's sufficient to understand forward inference and backprop rules for a single block
 - Full forward or backprop computation simply requires repeated application of these rules