

# Generative modeling 2

CSE 849 Deep Learning  
Spring 2025

Zijun Cui

# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

Input  
Data

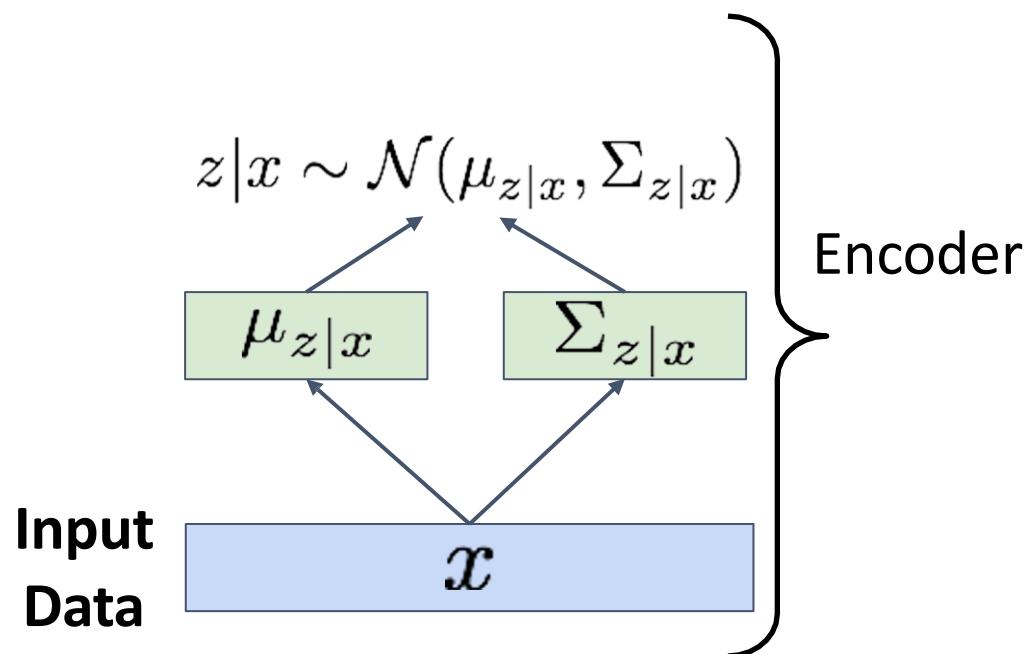
$x$

# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes

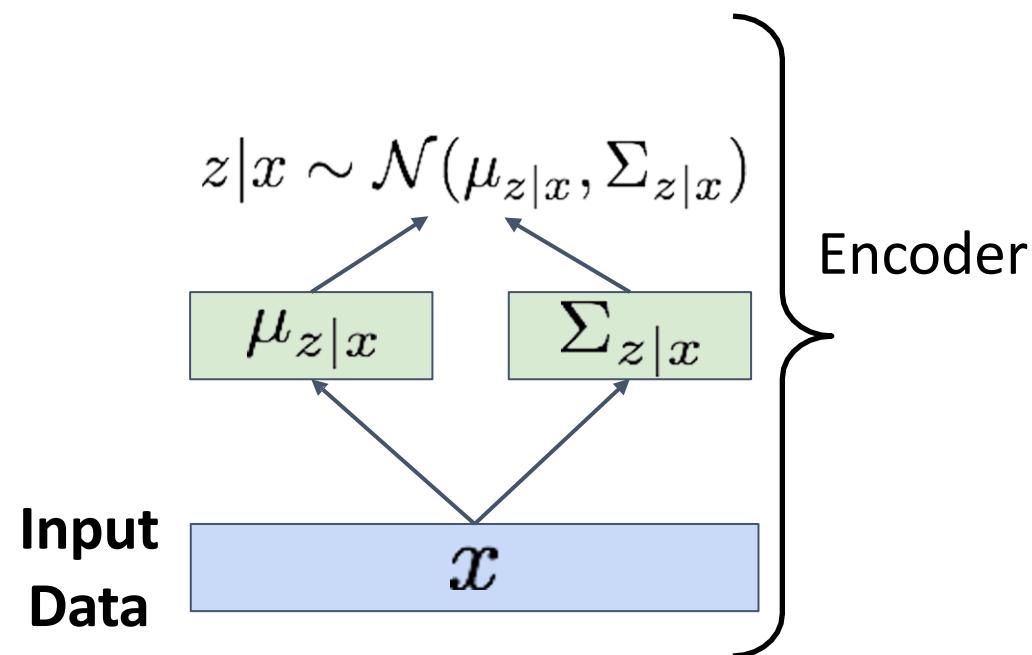


# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$**



# Variational Autoencoders

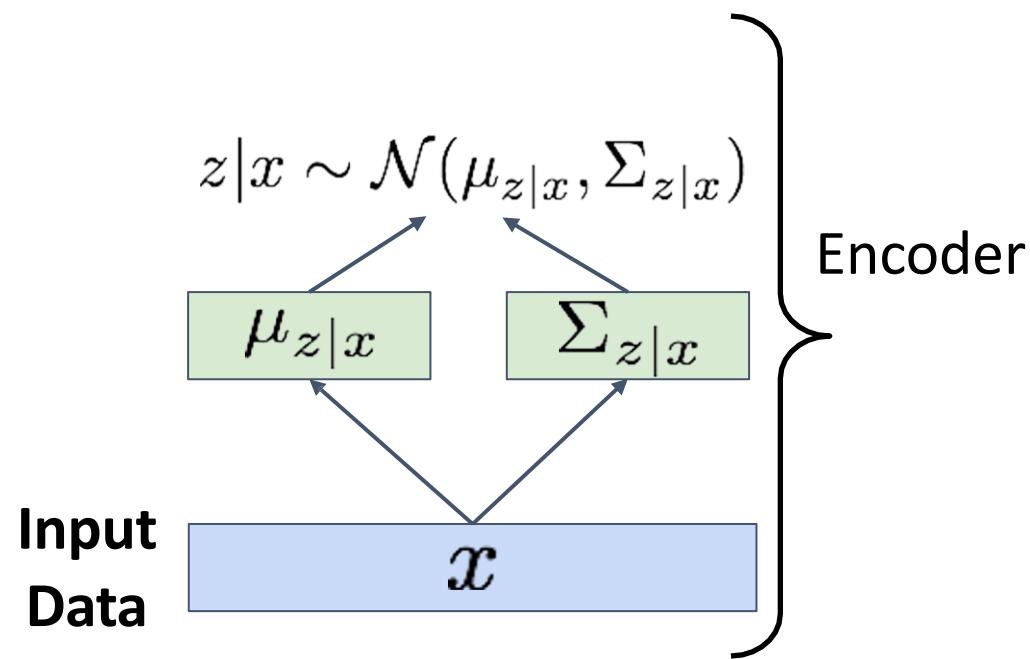
Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$**

$$\begin{aligned} -D_{KL} (q_\phi(z|x), p(z)) &= \int_Z q_\phi(z|x) \log \frac{p(z)}{q_\phi(z|x)} dz \\ &= \int_Z N(z; \mu_{z|x}, \Sigma_{z|x}) \log \frac{N(z; 0, I)}{N(z; \mu_{z|x}, \Sigma_{z|x})} dz \\ &= \frac{1}{2} \sum_{j=1}^J \left( 1 + \log \left( (\Sigma_{z|x})_j^2 \right) - (\mu_{z|x})_j^2 - (\Sigma_{z|x})_j^2 \right) \end{aligned}$$

Closed form solution when  $q_\phi$  is diagonal Gaussian and  $p$  is unit Gaussian (Assume  $z$  has dimension  $J$ )

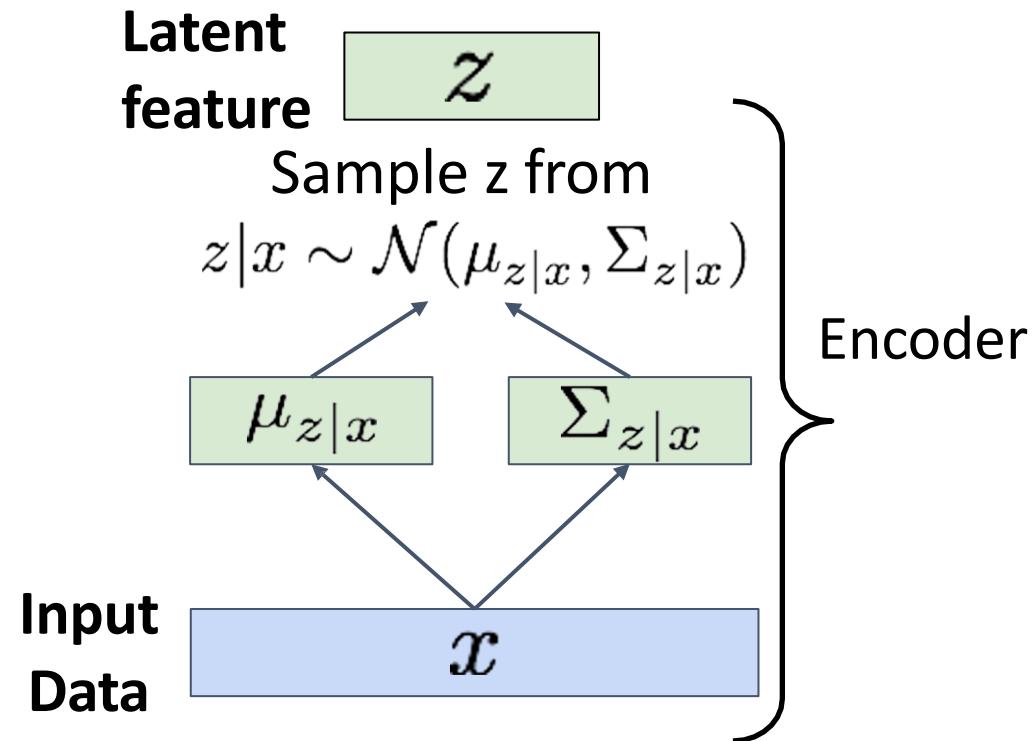


# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$**
3. Sample latent feature  $z$  from encoder output

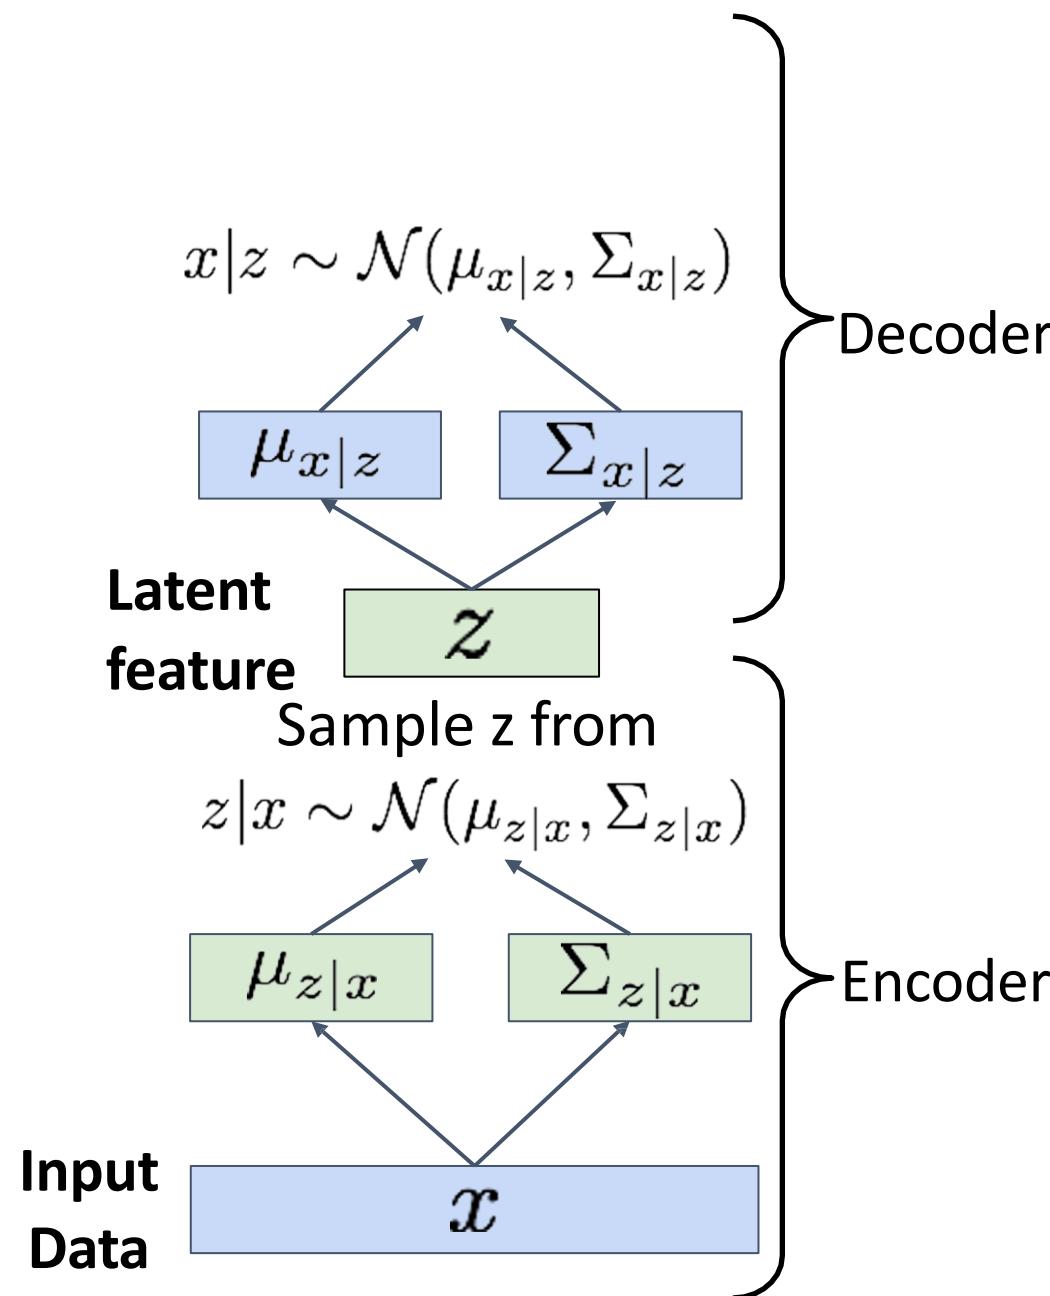


# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)} [\log p_\theta(x|z)] - D_{KL} (q_\phi(z|x), p(z))$$

1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$**
3. Sample latent feature  $z$  from encoder output
4. Run sampled  $z$  through **decoder** to get a distribution over data samples

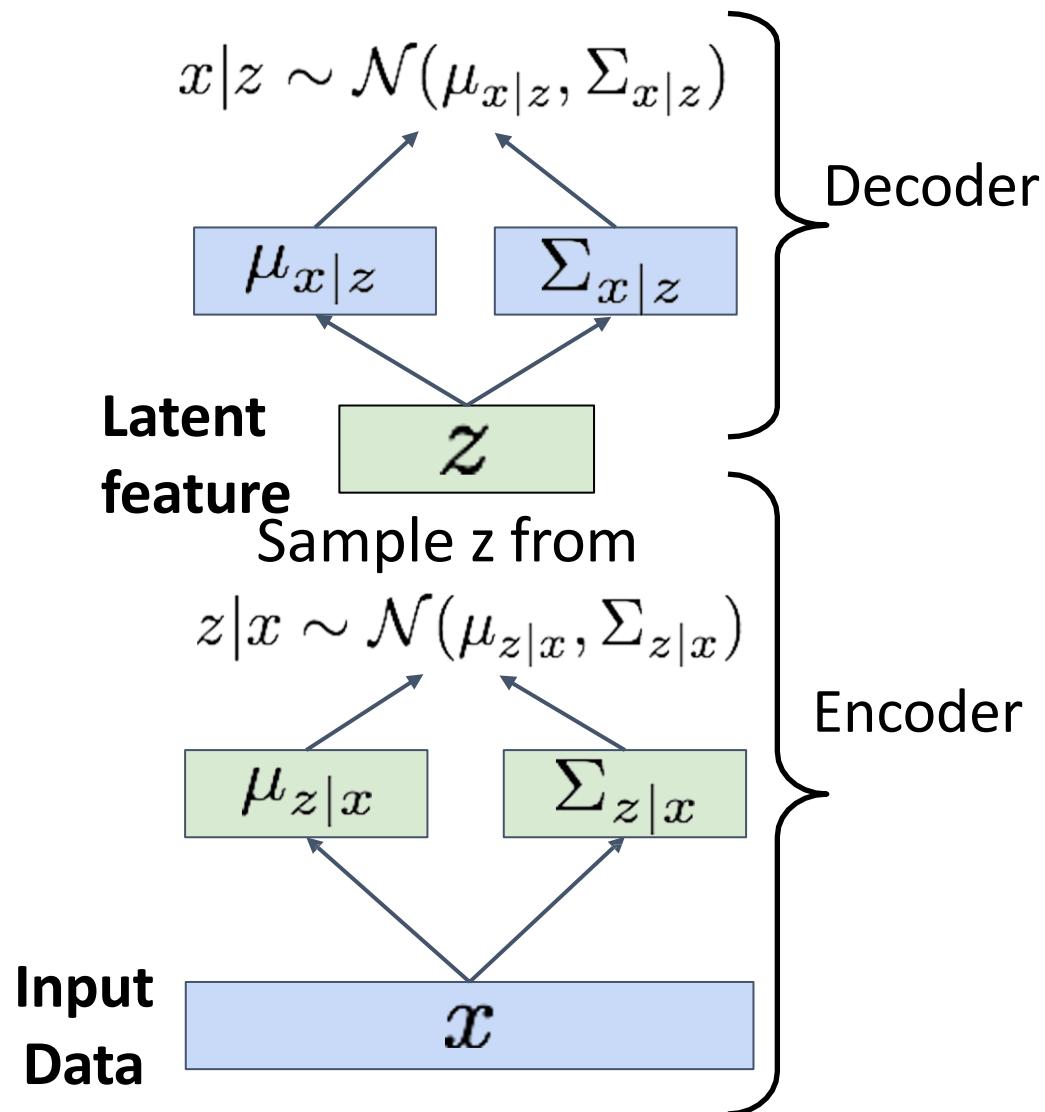


# Variational Autoencoders

Train by maximizing the  
**variational lower bound**

$$E_{z \sim q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x), p(z))$$

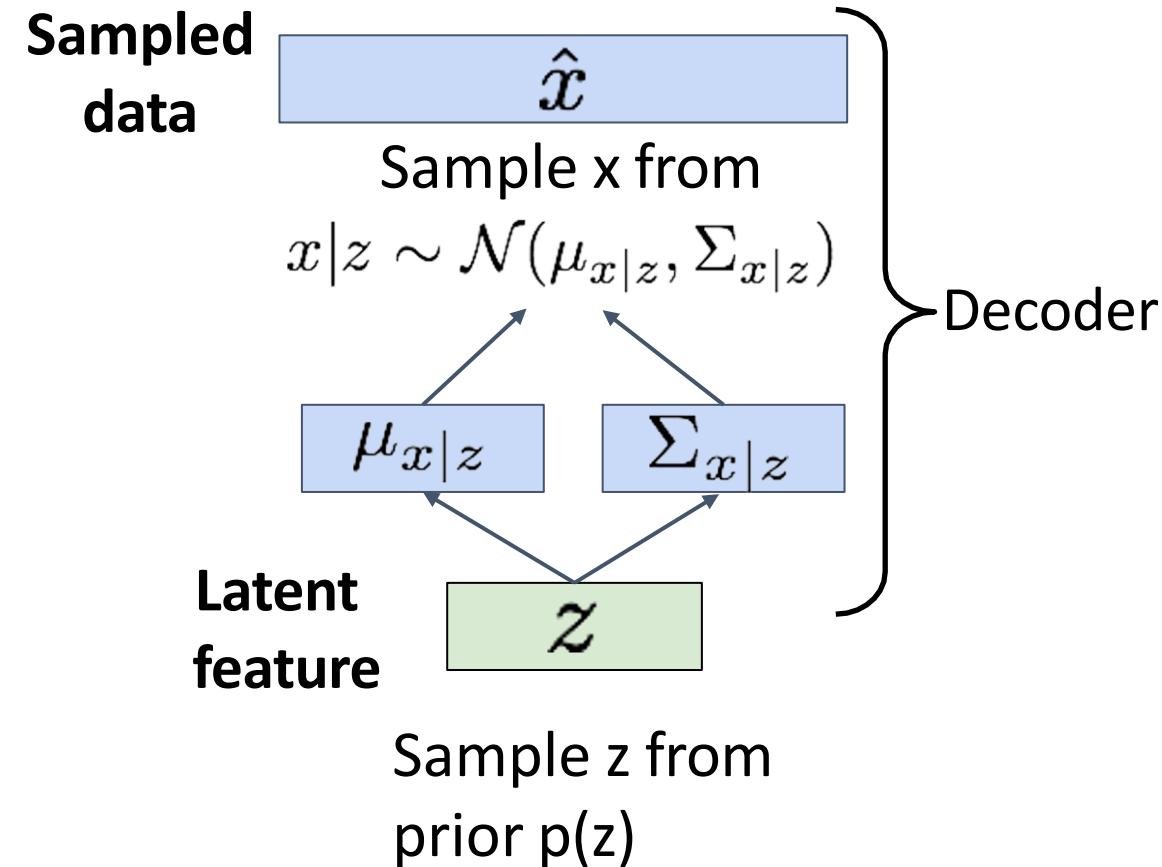
1. Run input data through **encoder** to get a distribution over latent codes
2. **Encoder output should match the prior  $p(z)$**
3. Sample feature  $z$  from encoder output
4. Run sampled  $z$  through **decoder** to get a distribution over data samples
5. **Original input data should be likely under the distribution output from (4)**



# Variational Autoencoders: Generating Data

After training we can generate new data!

1. Sample z from prior  $p(z)$
2. Run sampled z through decoder to get distribution over data x
3. Sample from distribution in (2) to generate data



# Variational Autoencoders: Generating Data

32x32 CIFAR-10



Labeled Faces in the Wild

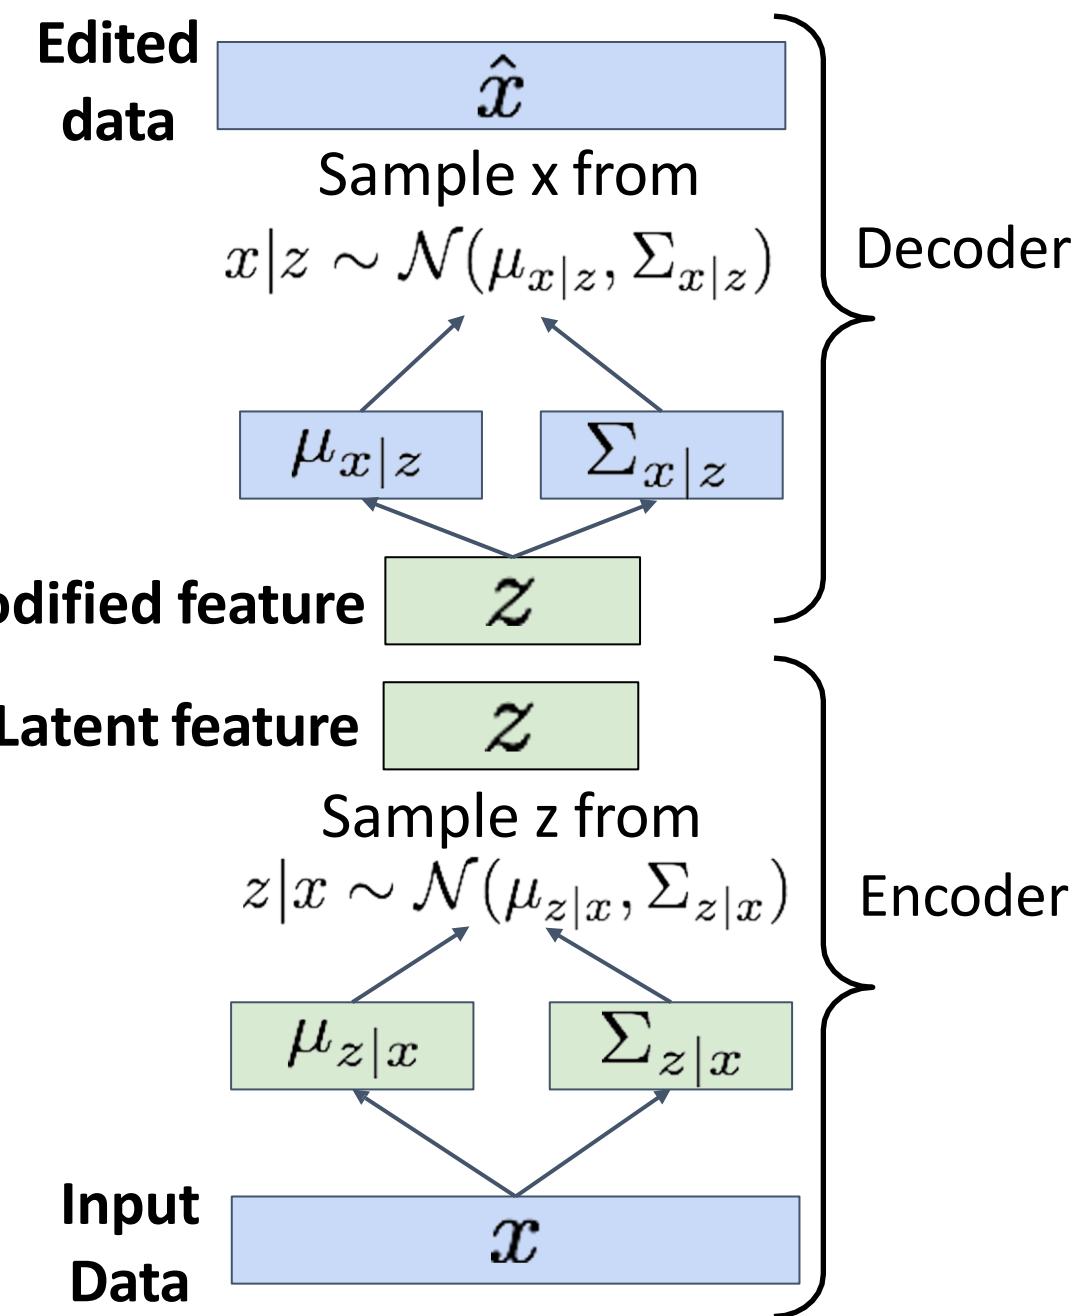


Figures from (L) Dirk Kingma et al. 2016; (R) Anders Larsen et al. 2017.

# Variational Autoencoders

After training we can **edit images**

1. Run input data through **encoder** to get a distribution over latent codes
2. Sample feature  $z$  from encoder output
3. Modify some dimensions of sampled feature
4. Run modified  $z$  through **decoder** to get a distribution over data samples
5. Sample new data from (4)



# Variational Autoencoders

The diagonal prior on  $p(z)$  causes dimensions of  $z$  to be independent

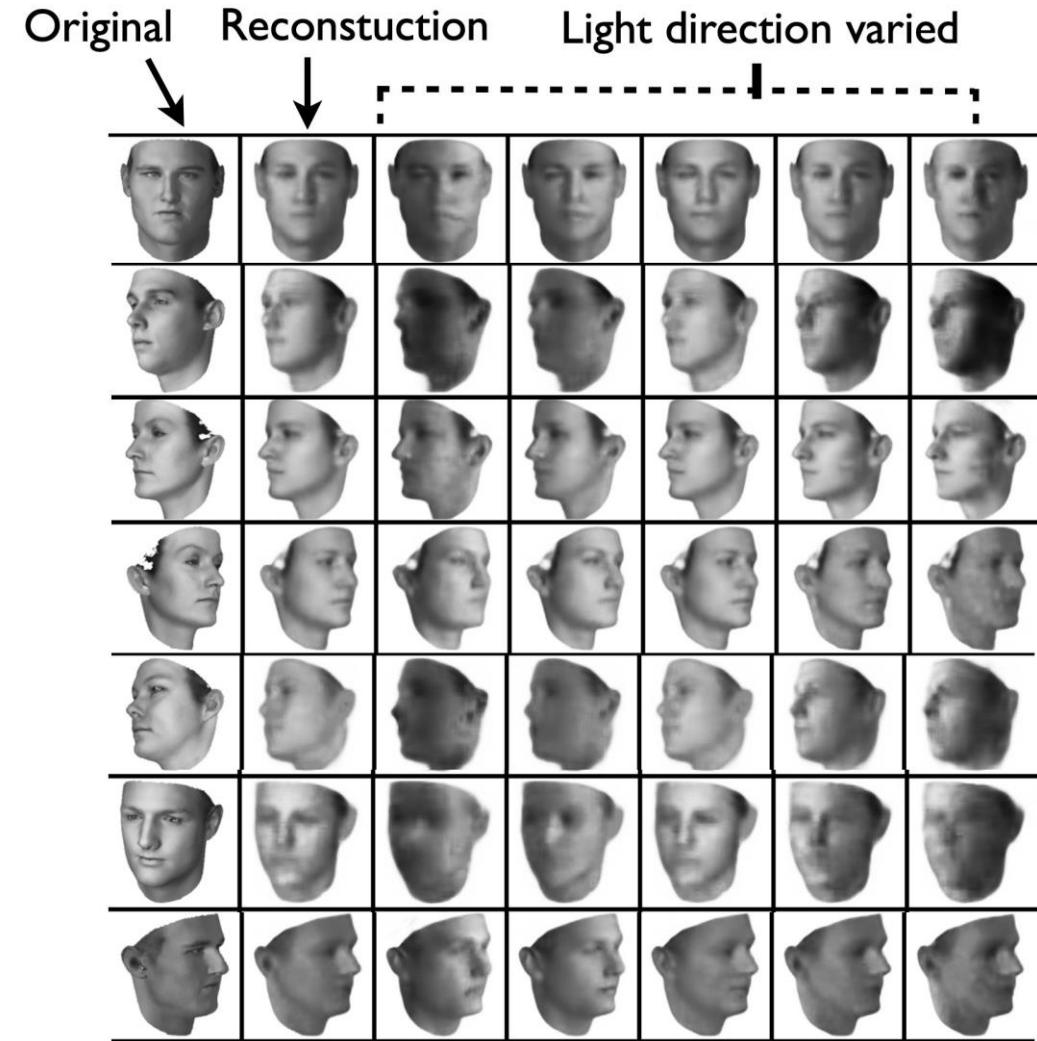
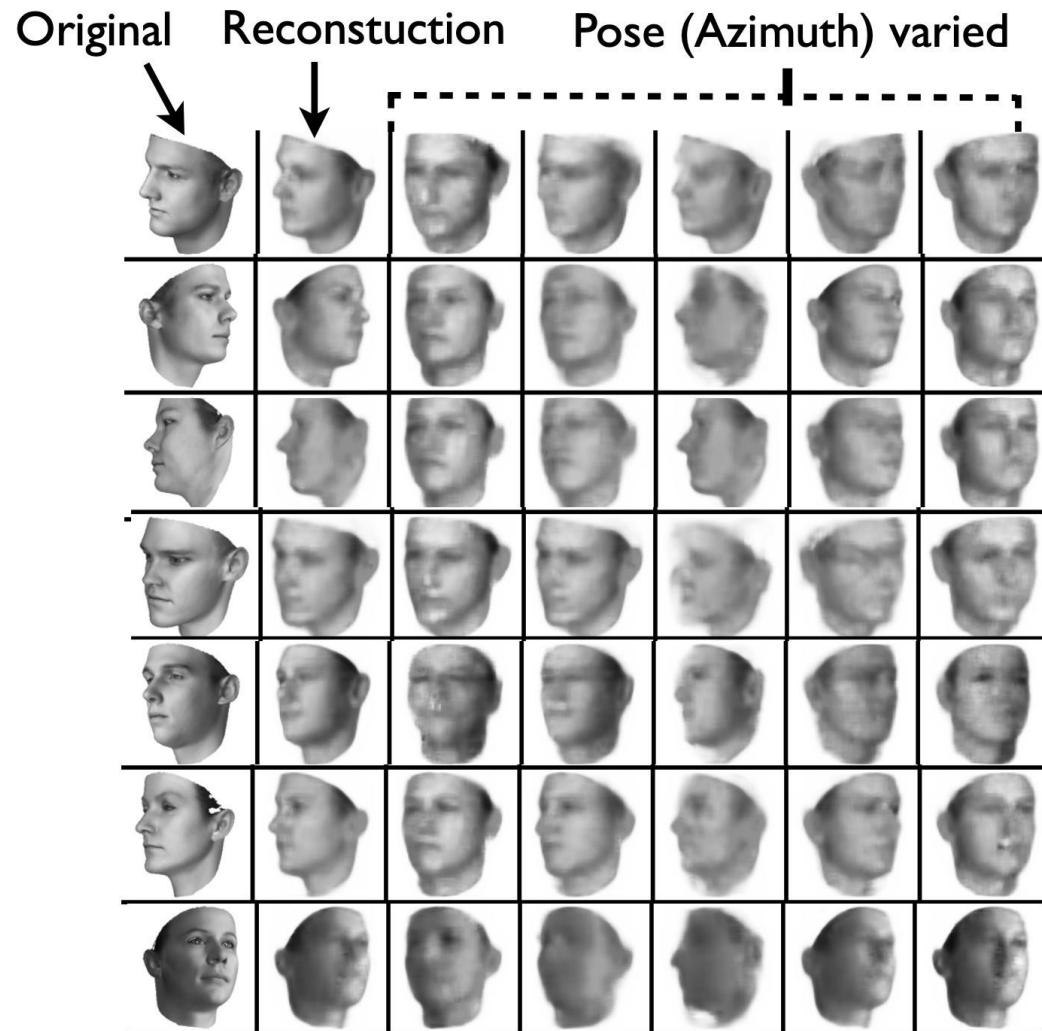
“Disentangling factors of variation”

Vary  $z_1$   
Degree of smile

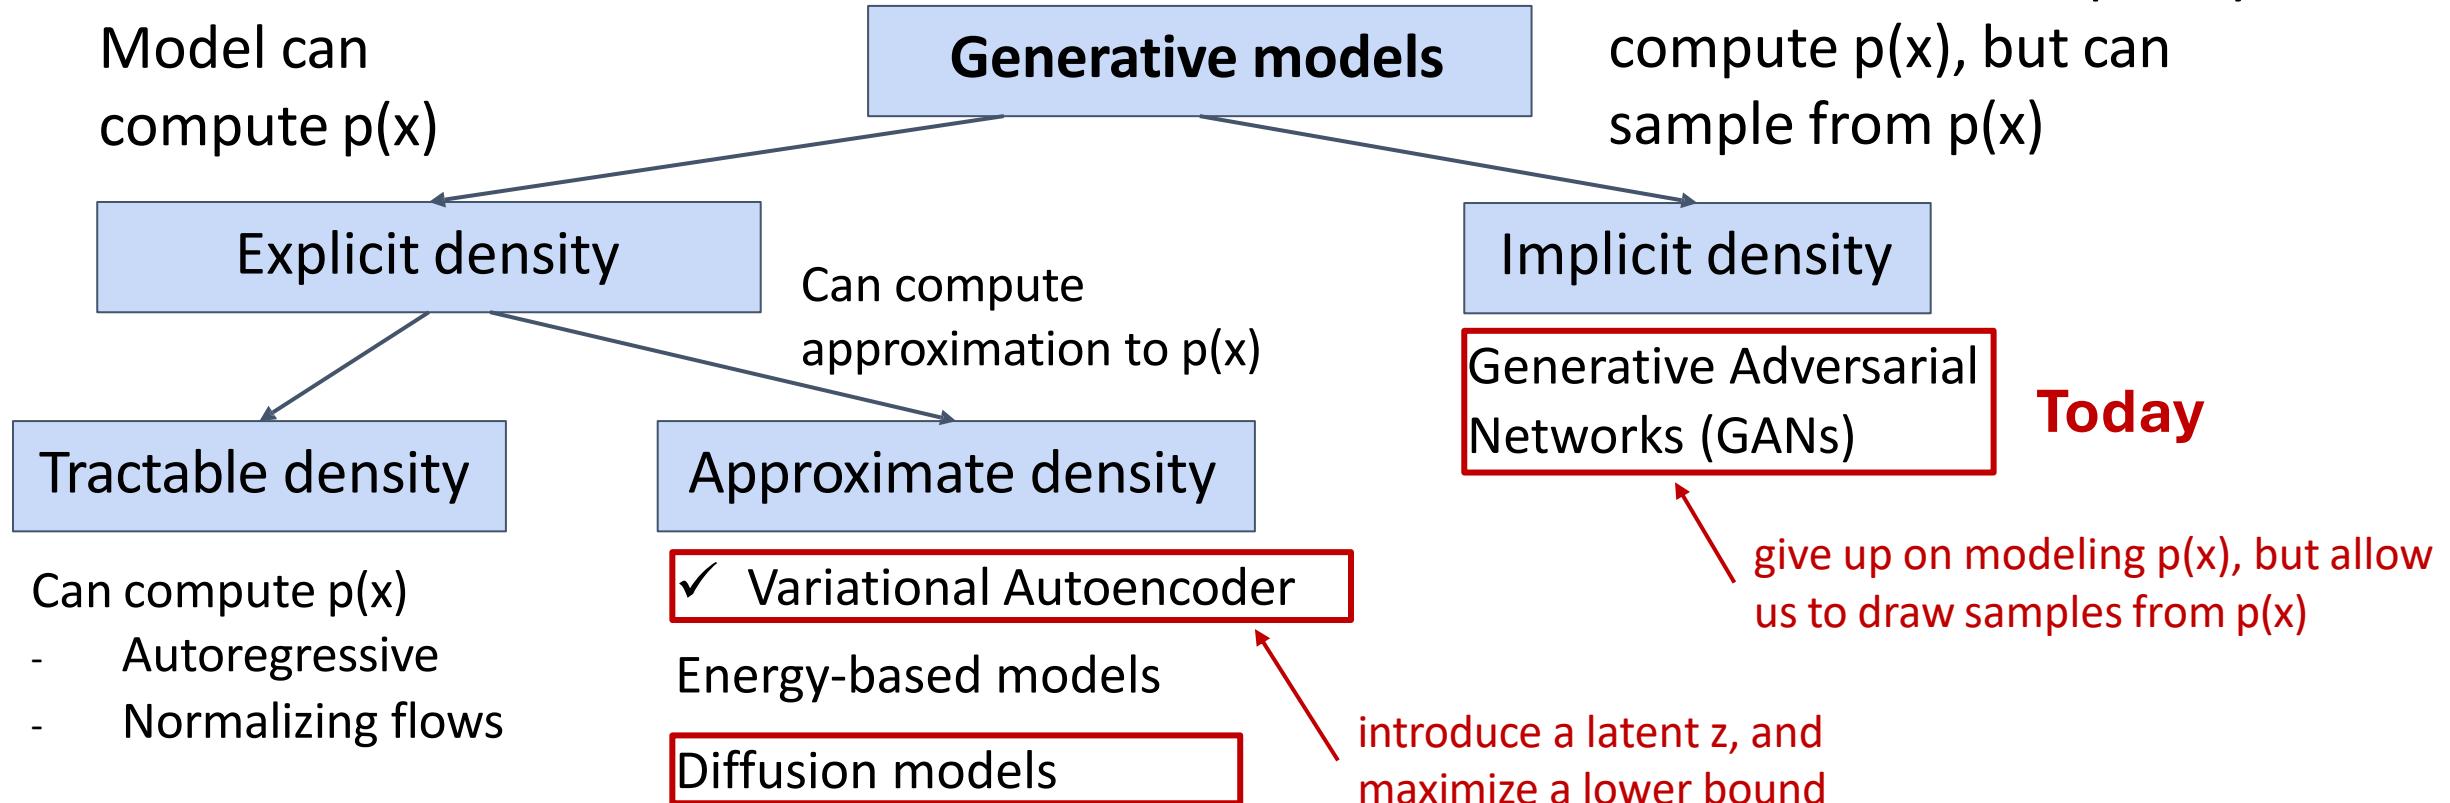


Vary  $z_2$   
Head pose

# Variational Autoencoders: Image Editing



# Generative Models So Far:



# Generative Adversarial Networks

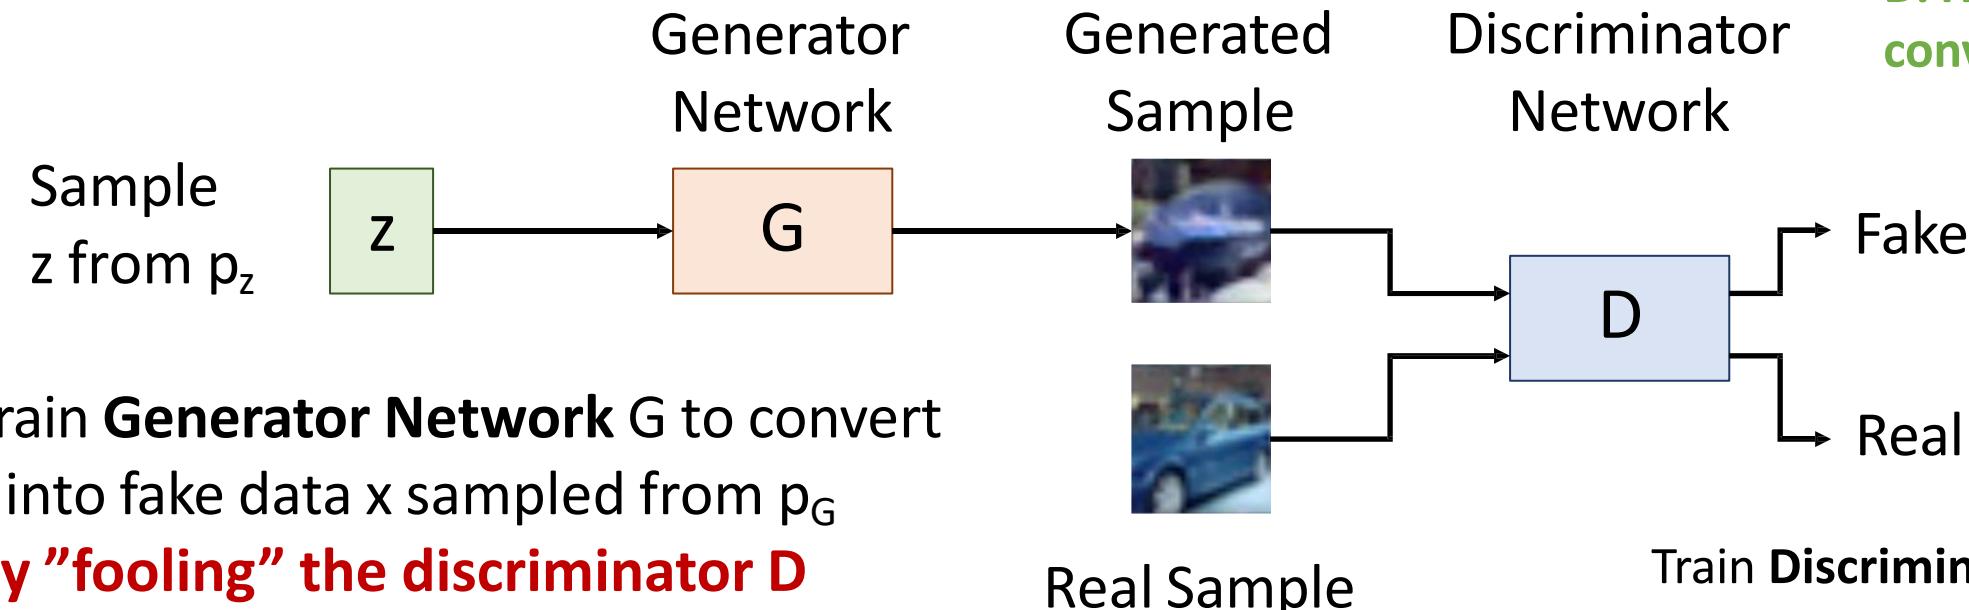
**Setup:** Assume we have data  $x_i$  drawn from distribution  $p_{\text{data}}(x)$ . Want to sample from  $p_{\text{data}}$ .

**Idea:** Introduce a latent variable  $z$  with simple prior  $p(z)$ .

Sample  $z \sim p(z)$  and pass to a **Generator Network**  $x = G(z)$

Then  $x$  is a sample from the **Generator distribution**  $p_G$ . Want  $p_G = p_{\text{data}}$ !

Jointly train G and D. Hopefully  $p_G$  converges to  $p_{\text{data}}$ !



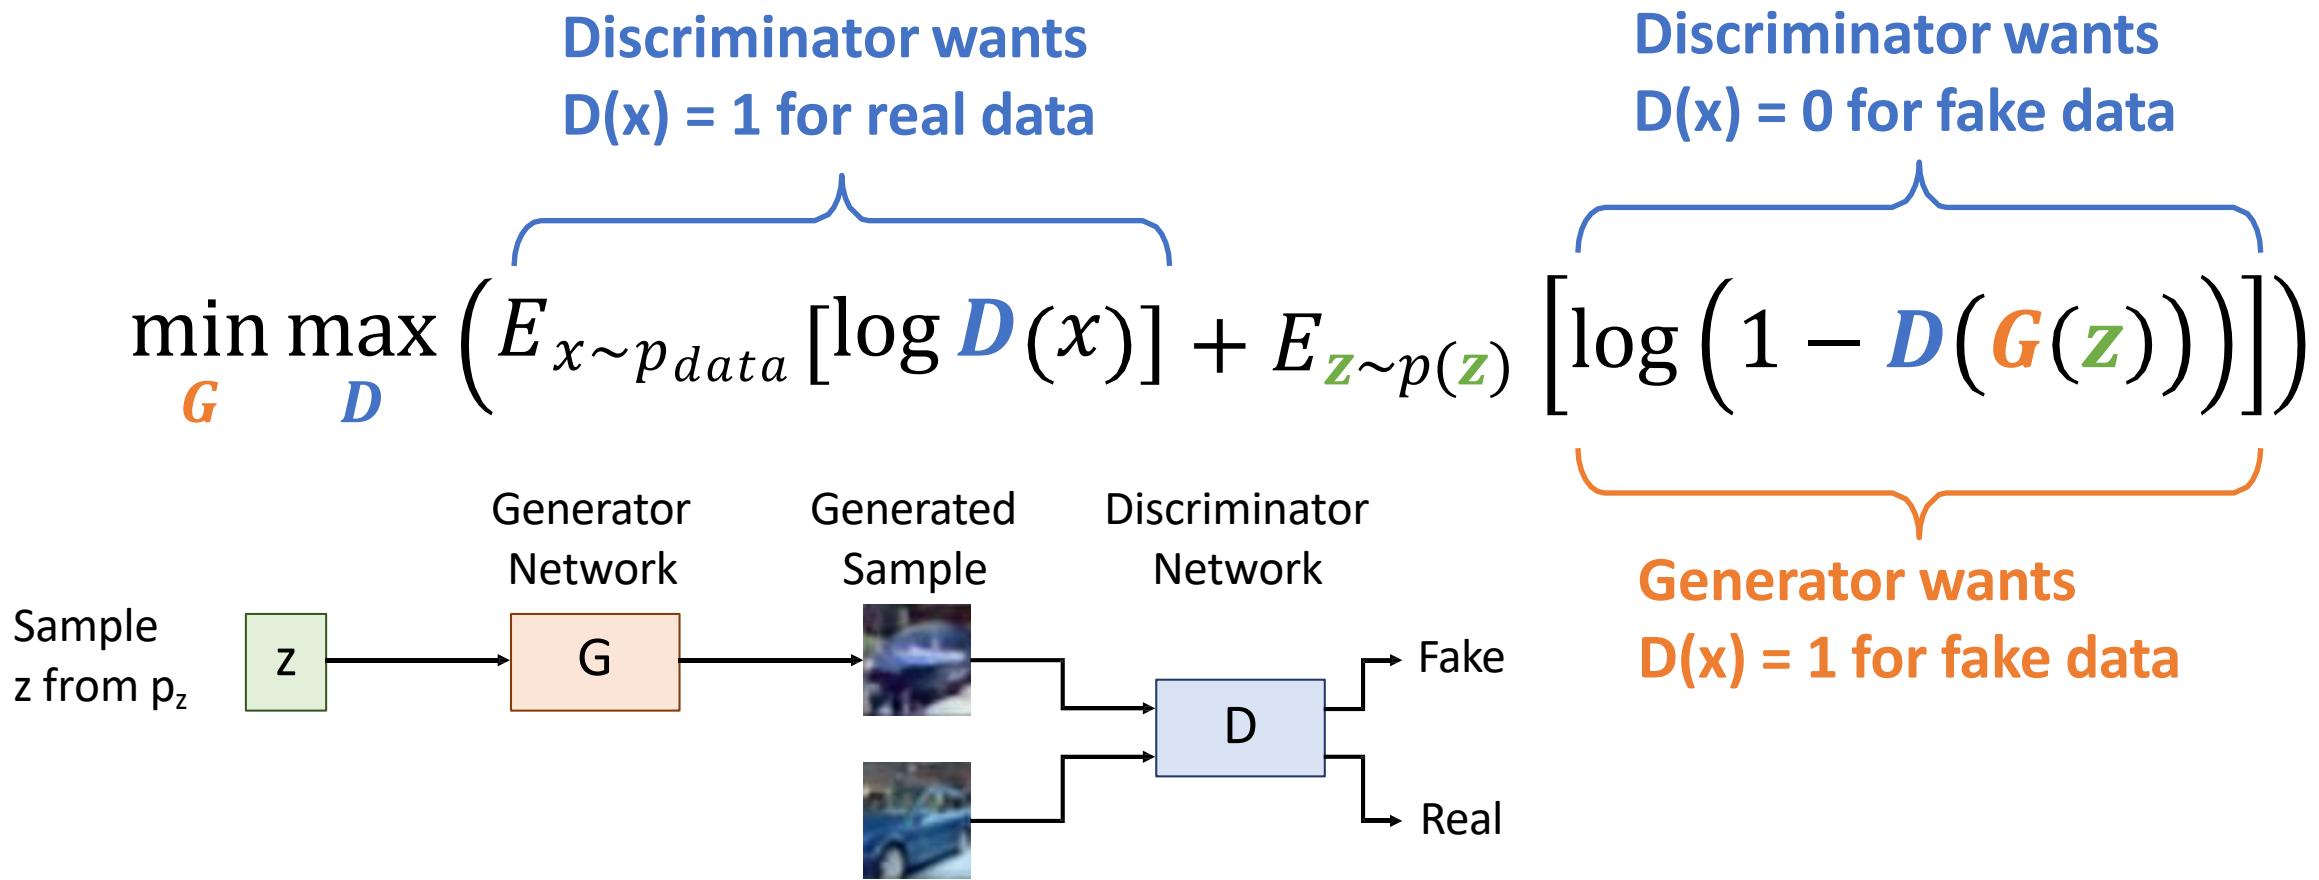
# Generative Adversarial Networks: Training Objective

Jointly train generator G and discriminator D with a **minimax game**

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

# Generative Adversarial Networks: Training Objective

Jointly train generator G and discriminator D with a **minimax game**



# Generative Adversarial Networks: Training Objective

Jointly train generator G and discriminator D with a **minimax game**

Train G and D using alternating gradient updates

$$\begin{aligned} & \min_{\mathbf{G}} \max_{\mathbf{D}} \left( E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - \mathbf{D}(\mathbf{G}(\mathbf{z})))] \right) \\ &= \min_{\mathbf{G}} \max_{\mathbf{D}} \mathbf{V}(\mathbf{G}, \mathbf{D}) \end{aligned}$$

We are not minimizing any overall loss.

No training curves to look at.

For t in 1, ... T:

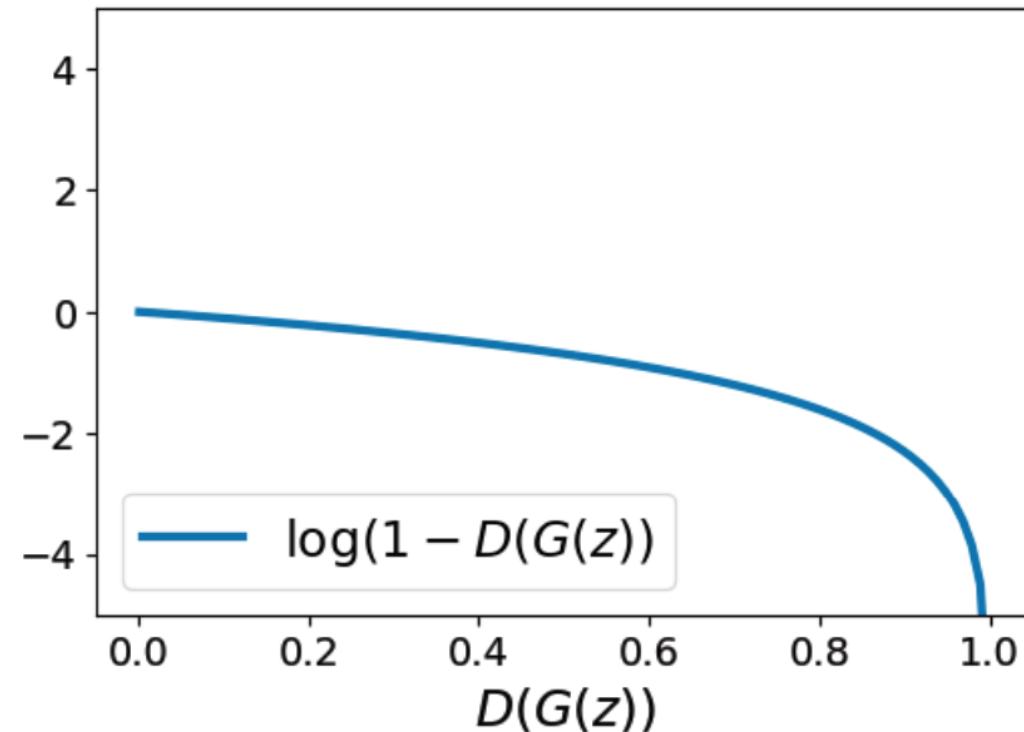
1. (Update D)  $\mathbf{D} = \mathbf{D} + \alpha_{\mathbf{D}} \frac{\partial \mathbf{V}}{\partial \mathbf{D}}$
2. (Update G)  $\mathbf{G} = \mathbf{G} - \alpha_{\mathbf{G}} \frac{\partial \mathbf{V}}{\partial \mathbf{G}}$

# Generative Adversarial Networks: Training Objective

Jointly train generator G and discriminator D with a **minimax game**

$$\min_{\mathbf{G}} \max_{\mathbf{D}} \left( E_{x \sim p_{data}} [\log \mathbf{D}(x)] + E_{\mathbf{z} \sim p(\mathbf{z})} [\log (1 - \mathbf{D}(\mathbf{G}(\mathbf{z})))] \right)$$

At start of training, generator is very bad and discriminator can easily tell apart real/fake, so  $D(G(z))$  close to 0



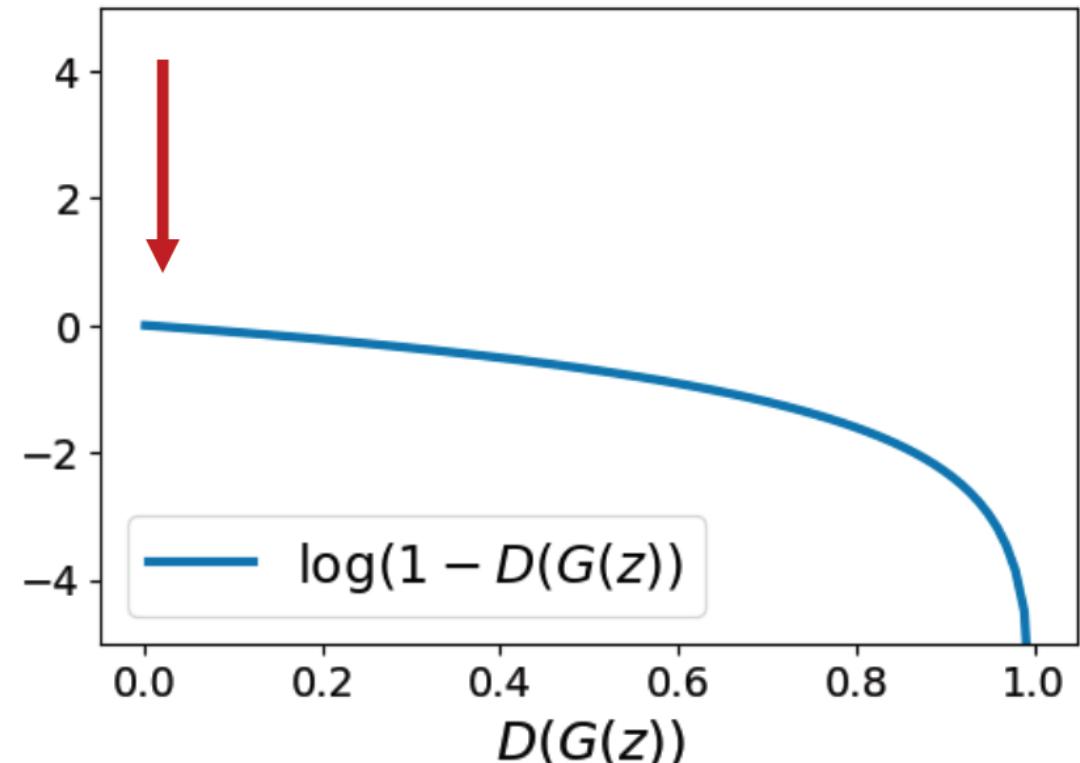
# Generative Adversarial Networks: Training Objective

Jointly train generator G and discriminator D with a **minimax game**

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

At start of training, generator is very bad and discriminator can easily tell apart real/fake, so  $D(G(z))$  close to 0

**Problem:** Vanishing gradients for G



# Generative Adversarial Networks: Training Objective

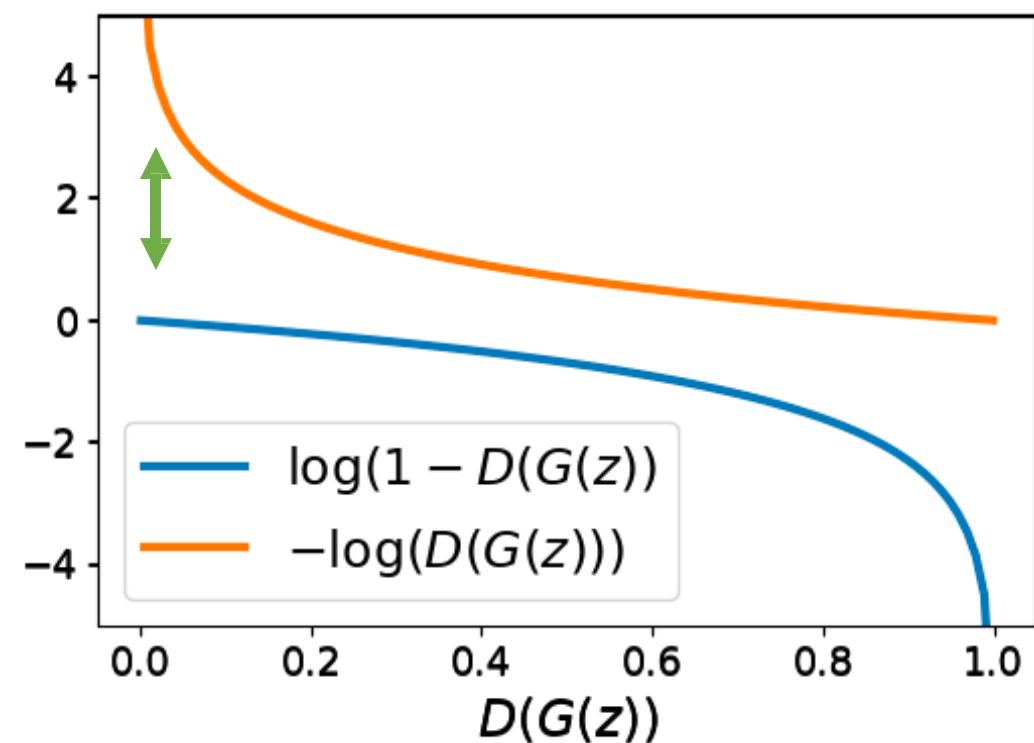
Jointly train generator G and discriminator D with a **minimax game**

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

At start of training, generator is very bad and discriminator can easily tell apart real/fake, so  $D(G(z))$  close to 0

**Problem:** Vanishing gradients for G

**Solution:** Right now G is trained to minimize  $\log(1 - D(G(z)))$ . Instead, train G to minimize  $-\log(D(G(z)))$ . Then G gets strong gradients at start of training!



# Generative Adversarial Networks: Optimality

Jointly train generator G and discriminator D with a **minimax game**

Why is this particular objective a good idea?

$$\min_{\textcolor{brown}{G}} \max_{\textcolor{blue}{D}} \left( E_{x \sim p_{data}} [\log \textcolor{blue}{D}(x)] + E_{\textcolor{green}{z} \sim p(\textcolor{green}{z})} [\log (1 - \textcolor{blue}{D}(\textcolor{brown}{G}(\textcolor{green}{z})))] \right)$$

This minimax game achieves its global minimum when  $p_G = p_{data}$ !

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

(Our objective so far)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

(Change of variables on second term)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \max_D \int_X (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

(Definition of expectation)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X \max_D (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

(Push  $\max_D$  inside integral)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X \max_D (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

$$f(y) = a \log y + b \log (1 - y)$$

(Side computation to compute max)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X \max_D (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

$$f(y) = a \log y + b \log(1 - y)$$

$$f'(y) = \frac{a}{y} - \frac{b}{1 - y}$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X \max_D (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

$$f(y) = a \log y + b \log(1 - y) \quad f'(y) = 0 \Leftrightarrow y = \frac{a}{a+b} \text{ (local max)}$$

$$f'(y) = \frac{a}{y} - \frac{b}{1-y}$$

**Optimal Discriminator:**  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X \max_D (p_{data}(x) \log D(x) + p_G(x) \log (1 - D(x))) dx$$

**Optimal Discriminator:**  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log (1 - D_G^*(x))) dx$$

**Optimal Discriminator:**  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{x \sim p_G} [\log (1 - D(x))] \right)$$

$$= \min_G \int_X (p_{data}(x) \log D_G^*(x) + p_G(x) \log (1 - D_G^*(x))) dx$$

$$= \min_G \int_X \left( p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) dx$$

**Optimal Discriminator:**  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$

# Generative Adversarial Networks: Optimality

$$\begin{aligned} & \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right) \\ &= \min_G \int_X \left( p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) dx \end{aligned}$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \int_X \left( p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) dx$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right] \right)$$

(Definition of expectation)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \int_X \left( p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) dx$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2}{2} \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2}{2} \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right] \right)$$

(Multiply by a constant)

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \int_X \left( p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} + p_G(x) \log \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right) dx$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2}{2} \frac{p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2}{2} \frac{p_G(x)}{p_{data}(x) + p_G(x)} \right] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

**Kullback-Leibler Divergence:**

$$KL(p, q) = E_{x \sim p} \left[ \log \frac{p(x)}{q(x)} \right]$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

$$= \min_G \left( KL \left( p_{data}, \frac{p_{data} + p_G}{2} \right) + KL \left( p_G, \frac{p_{data} + p_G}{2} \right) - \log 4 \right)$$

**Kullback-Leibler Divergence:**

$$KL(p, q) = E_{x \sim p} \left[ \log \frac{p(x)}{q(x)} \right]$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

$$= \min_G \left( KL \left( p_{data}, \frac{p_{data} + p_G}{2} \right) + KL \left( p_G, \frac{p_{data} + p_G}{2} \right) - \log 4 \right)$$

**Jensen-Shannon Divergence:**

$$JSD(p, q) = \frac{1}{2} KL \left( p, \frac{p + q}{2} \right) + \frac{1}{2} KL \left( q, \frac{p + q}{2} \right)$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

$$= \min_G \left( KL \left( p_{data}, \frac{p_{data} + p_G}{2} \right) + KL \left( p_G, \frac{p_{data} + p_G}{2} \right) - \log 4 \right)$$

$$= \min_G (2 * JSD(p_{data}, p_G) - \log 4)$$

**Jensen-Shannon Divergence:**

$$JSD(p, q) = \frac{1}{2} KL \left( p, \frac{p + q}{2} \right) + \frac{1}{2} KL \left( q, \frac{p + q}{2} \right)$$

# Generative Adversarial Networks: Optimality

$$\min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right)$$

$$= \min_G \left( E_{x \sim p_{data}} \left[ \log \frac{2 * p_{data}(x)}{p_{data}(x) + p_G(x)} \right] + E_{x \sim p_G} \left[ \log \frac{2 * p_G(x)}{p_{data}(x) + p_G(x)} \right] - \log 4 \right)$$

$$= \min_G \left( KL \left( p_{data}, \frac{p_{data} + p_G}{2} \right) + KL \left( p_G, \frac{p_{data} + p_G}{2} \right) - \log 4 \right)$$

$$= \min_G (2 * JSD(p_{data}, p_G) - \log 4)$$

JSD is always nonnegative, and zero only when the two distributions are equal!  
Thus  $p_{data} = p_G$  is the global min

**Jensen-Shannon Divergence:**

$$JSD(p, q) = \frac{1}{2} KL \left( p, \frac{p + q}{2} \right) + \frac{1}{2} KL \left( q, \frac{p + q}{2} \right)$$

# Generative Adversarial Networks: Optimality

$$\begin{aligned} & \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right) \\ &= \min_G (2 * JSD(p_{data}, p_G) - \log 4) \end{aligned}$$

# Generative Adversarial Networks: Optimality

$$\begin{aligned} & \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right) \\ &= \min_G (2 * JSD(p_{data}, p_G) - \log 4) \end{aligned}$$

**Summary:** The global minimum of the minimax game happens when:

1.  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$  (Optimal discriminator for any G)
2.  $p_G(x) = p_{data}(x)$  (Optimal generator for optimal D)

# Generative Adversarial Networks: Optimality

$$\begin{aligned} & \min_G \max_D \left( E_{x \sim p_{data}} [\log D(x)] + E_{z \sim p(z)} [\log (1 - D(G(z)))] \right) \\ &= \min_G (2 * JSD(p_{data}, p_G) - \log 4) \end{aligned}$$

**Summary:** The global minimum of the minimax game happens when:

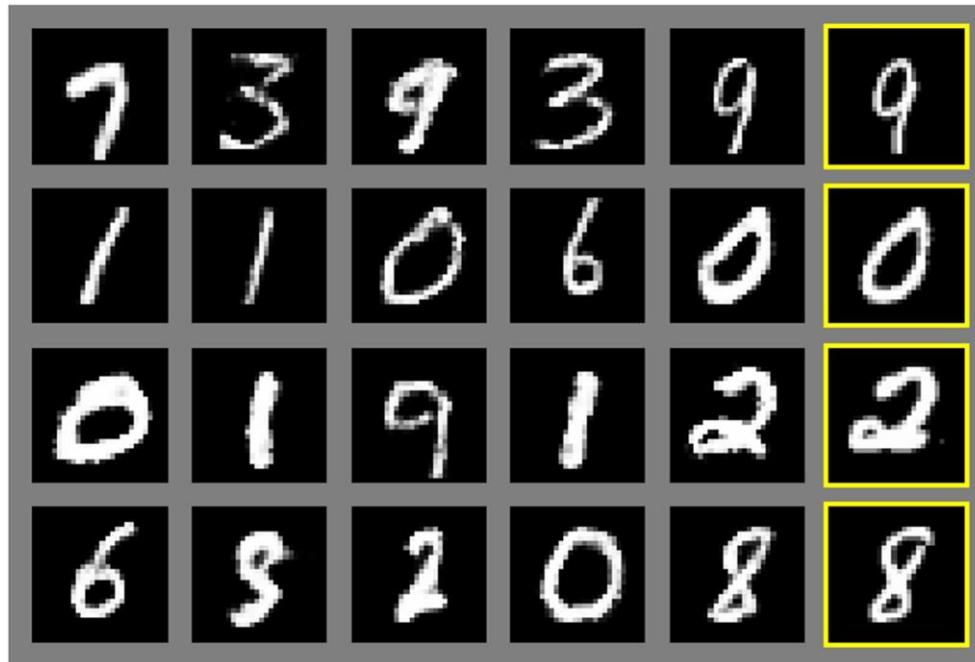
1.  $D_G^*(x) = \frac{p_{data}(x)}{p_{data}(x) + p_G(x)}$  (Optimal discriminator for any G)
2.  $p_G(x) = p_{data}(x)$  (Optimal generator for optimal D)

## Caveats:

1. G and D are neural nets with fixed architecture. We don't know whether they can actually represent the optimal D and G.
2. This tells us nothing about convergence to the optimal solution

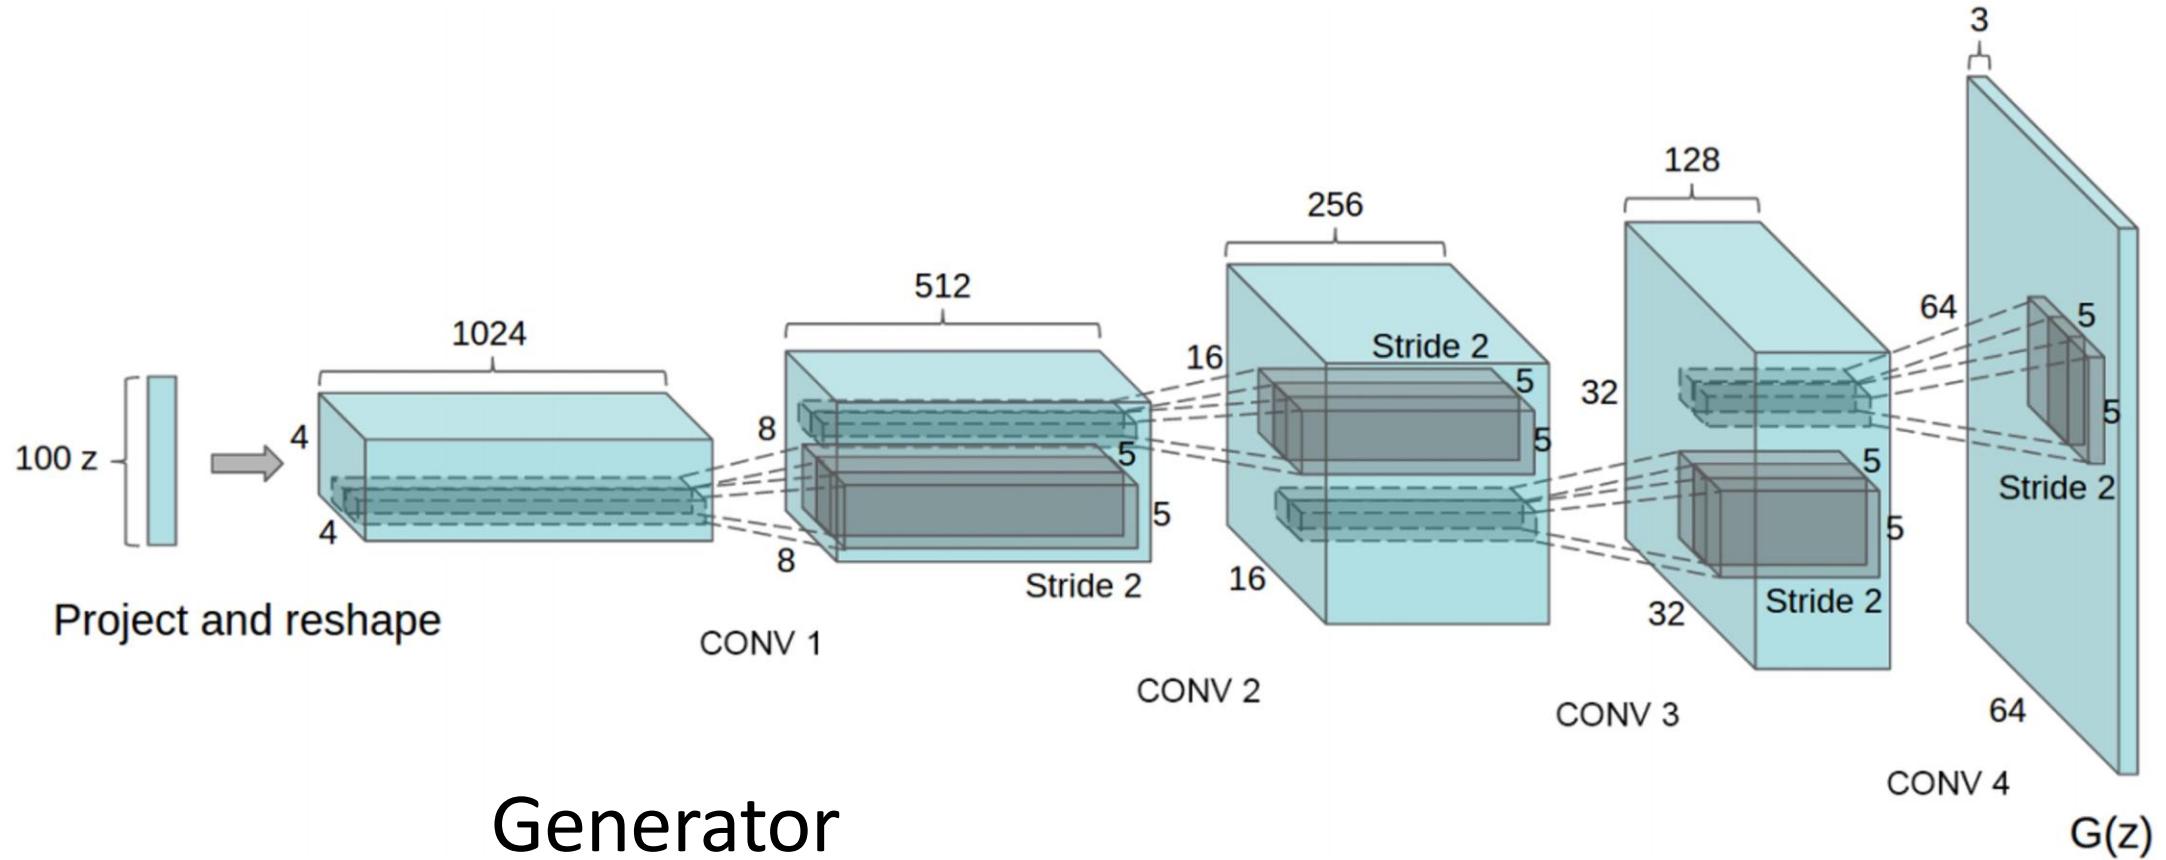
# Generative Adversarial Networks: Results

Generated samples



Nearest neighbor from training set

# Generative Adversarial Networks: DC-GAN



Generator

# Generative Adversarial Networks: DC-GAN

Samples  
from the  
model  
look  
much  
better!

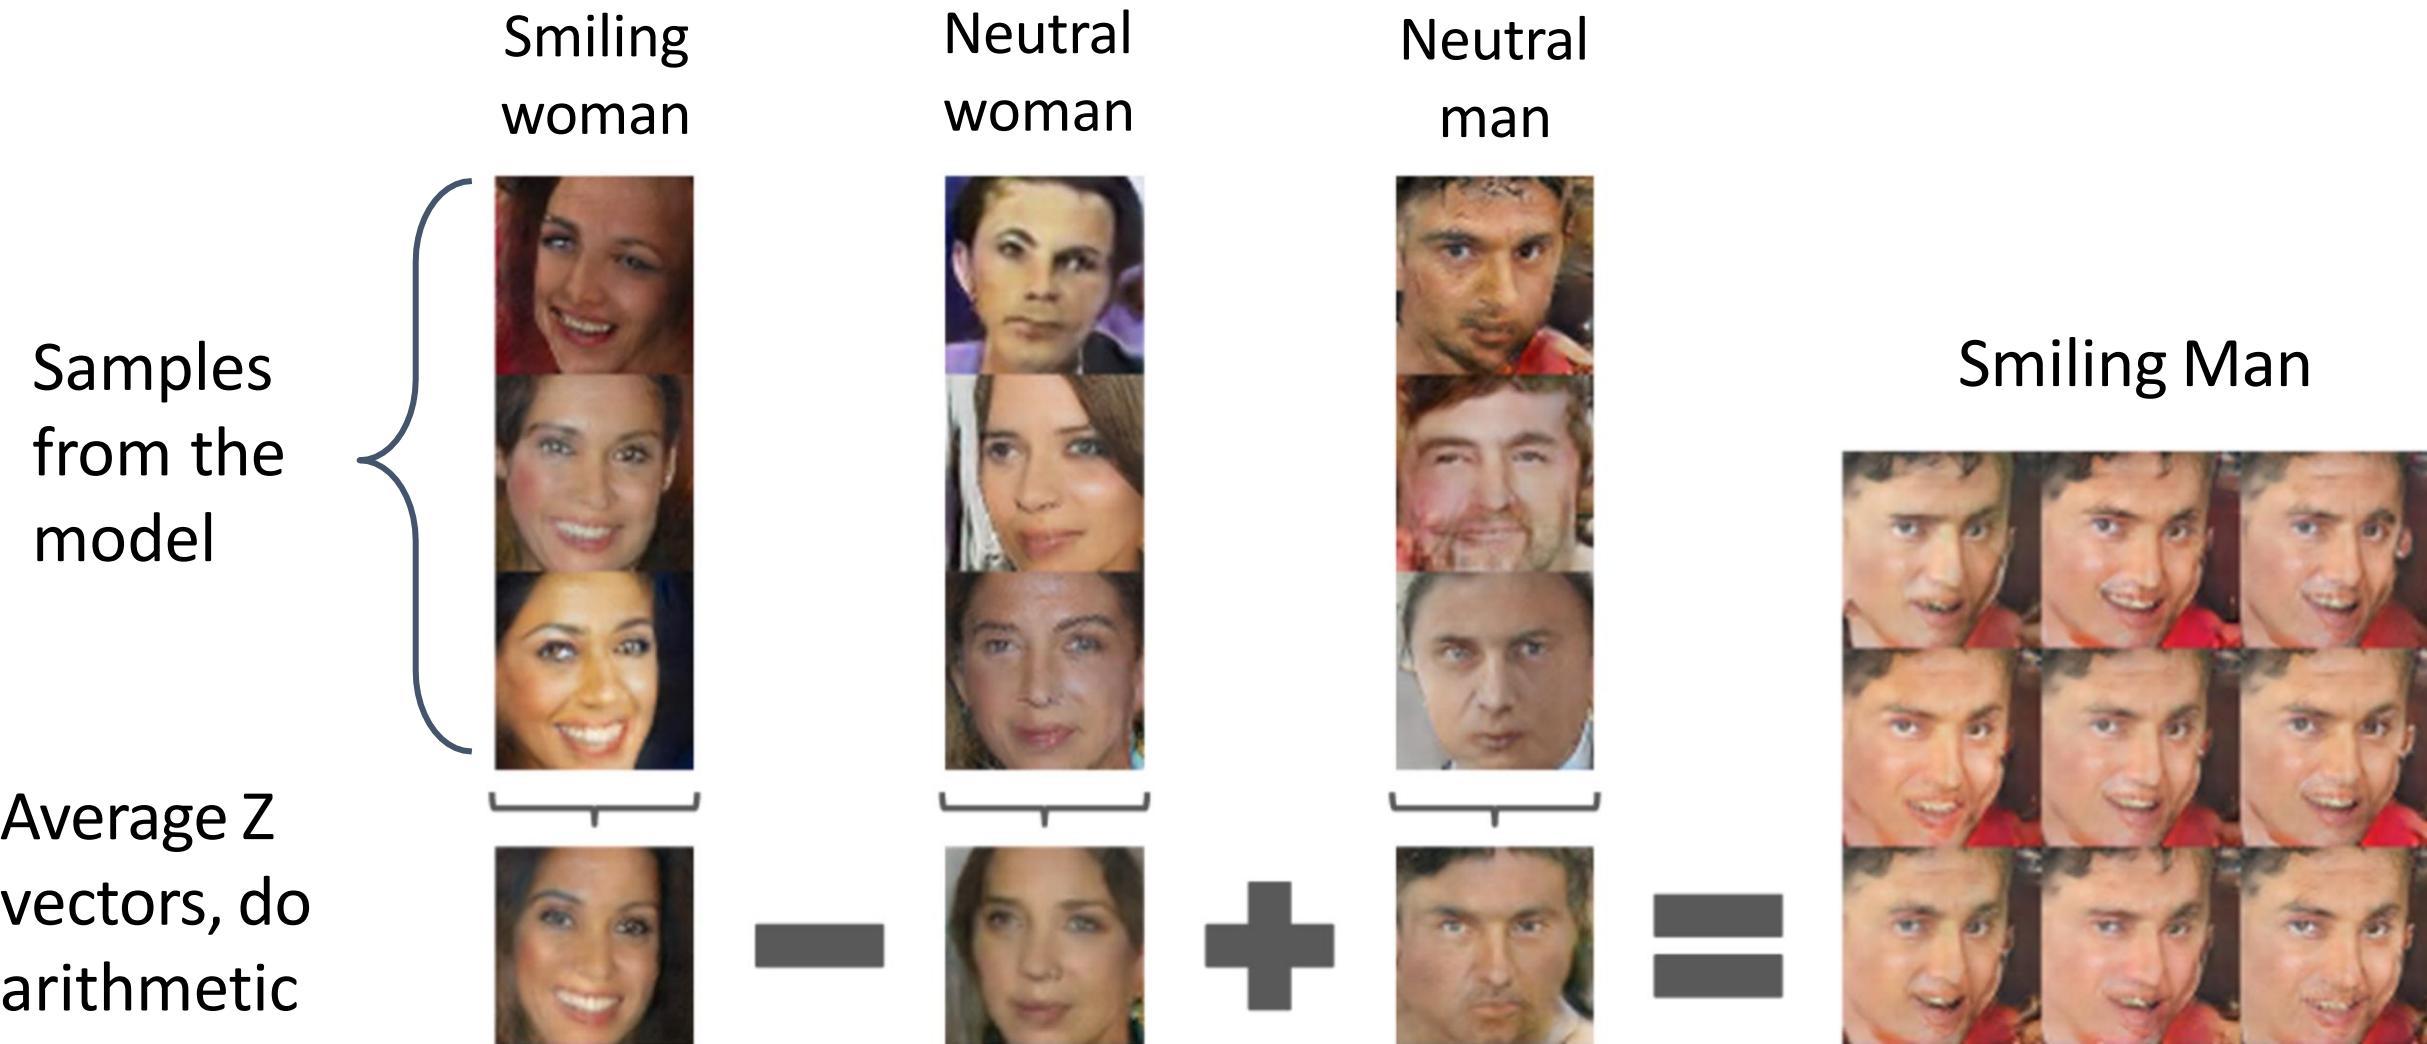


# Generative Adversarial Networks: Interpolation

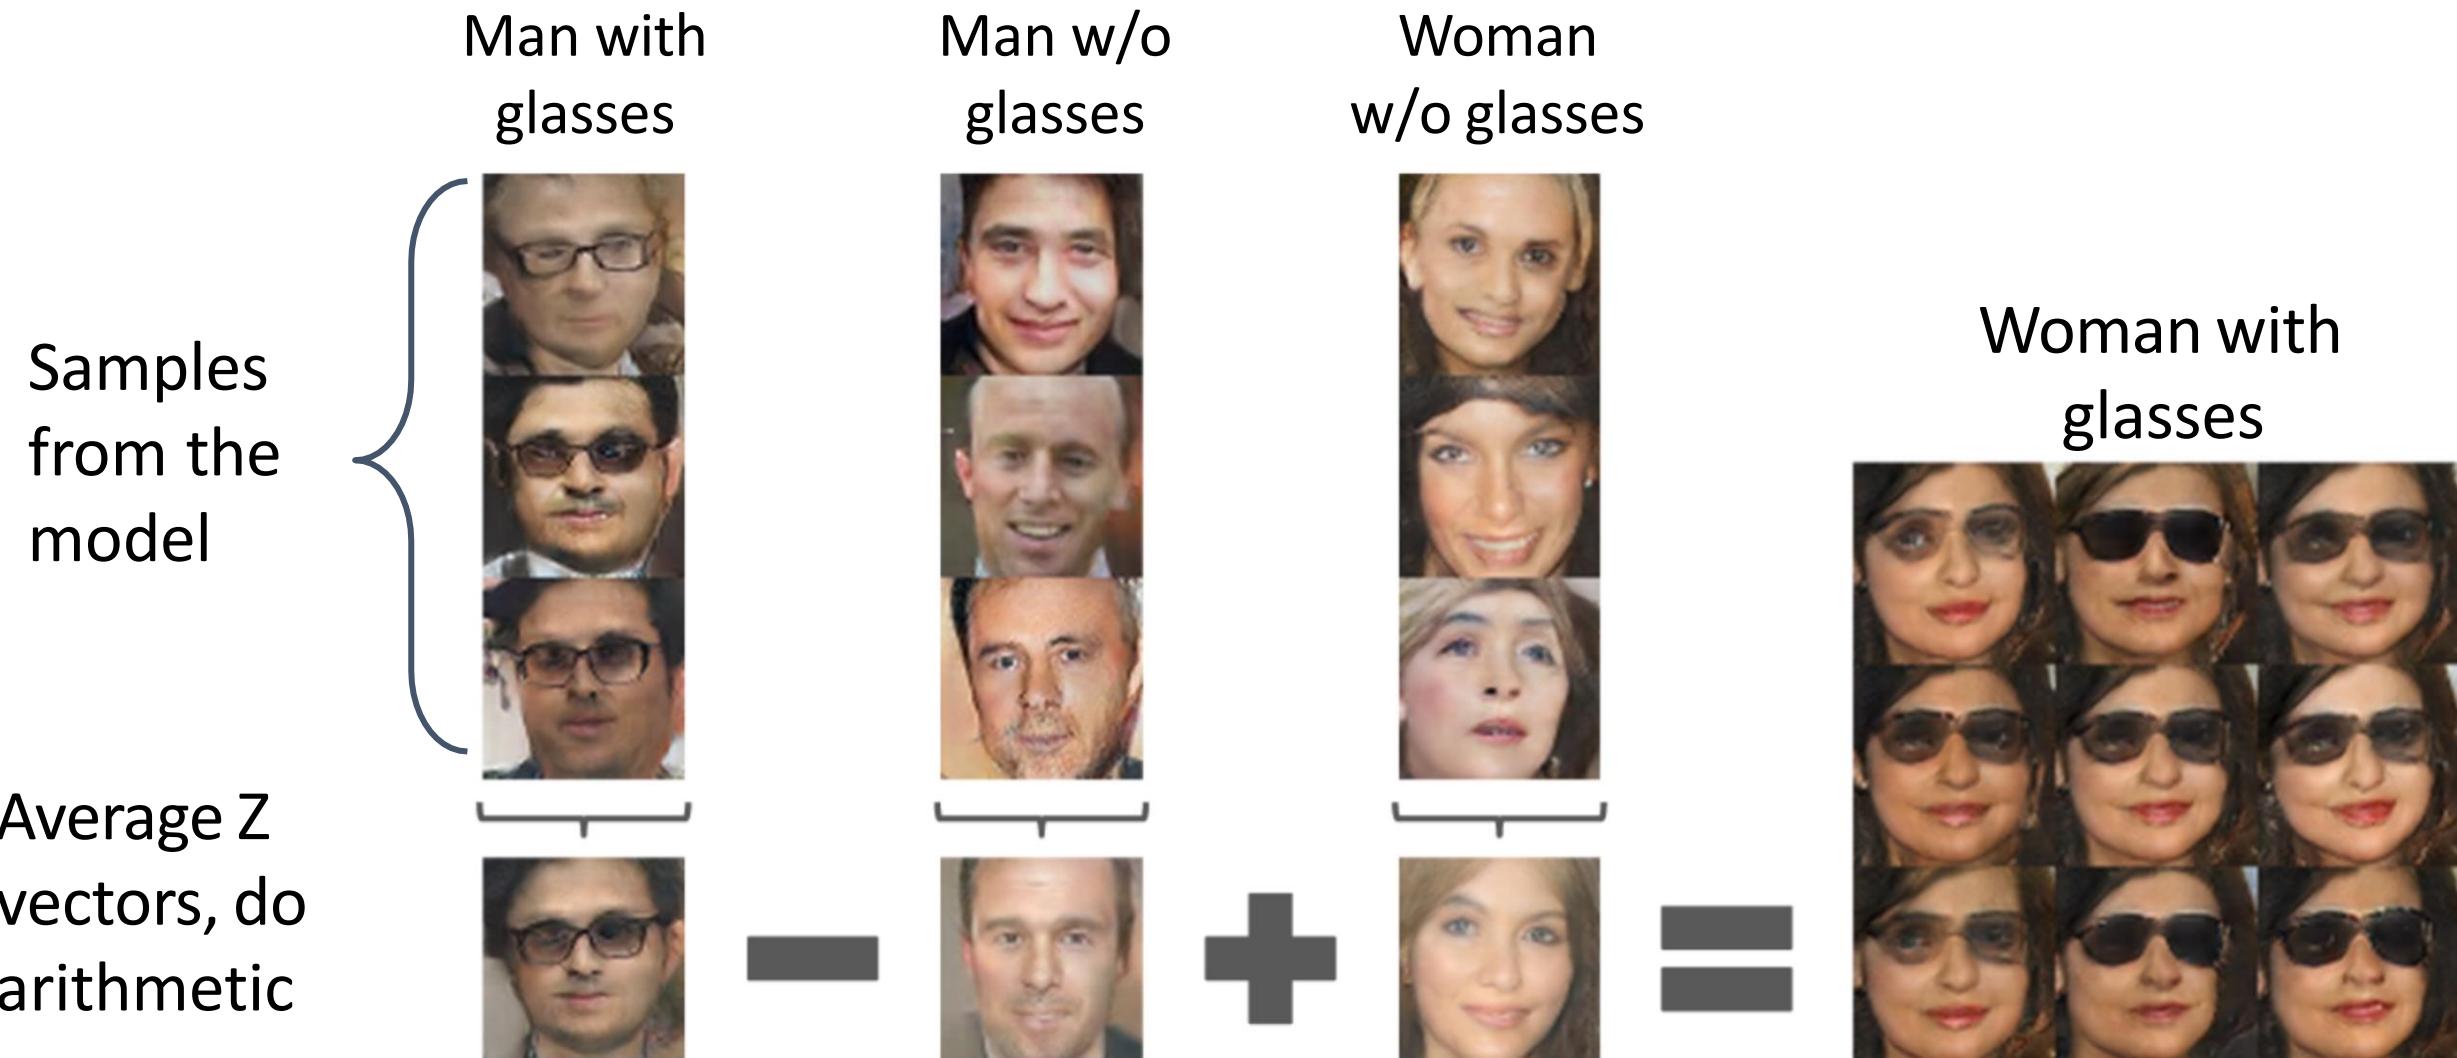
Interpolating  
between  
points in  
latent  $z$   
space



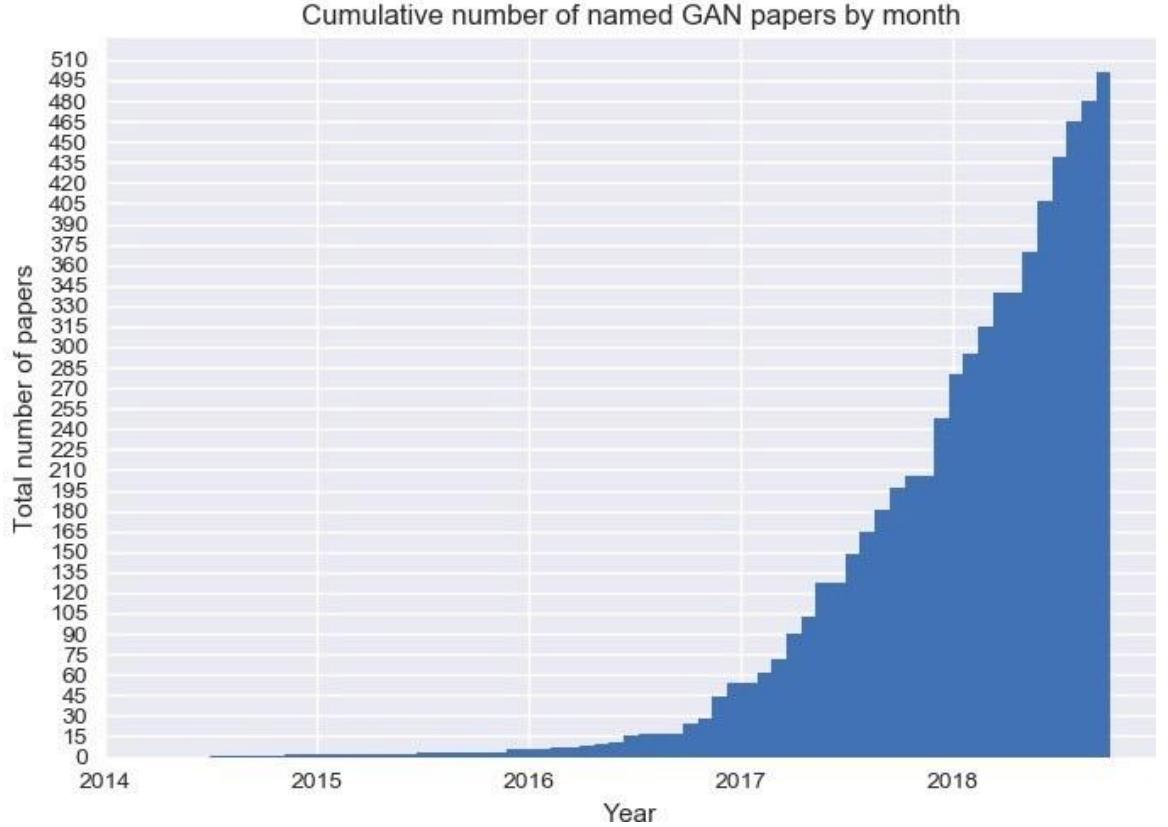
# Generative Adversarial Networks: Vector Math



# Generative Adversarial Networks: Vector Math



# 2017 to present: Explosion of GANs

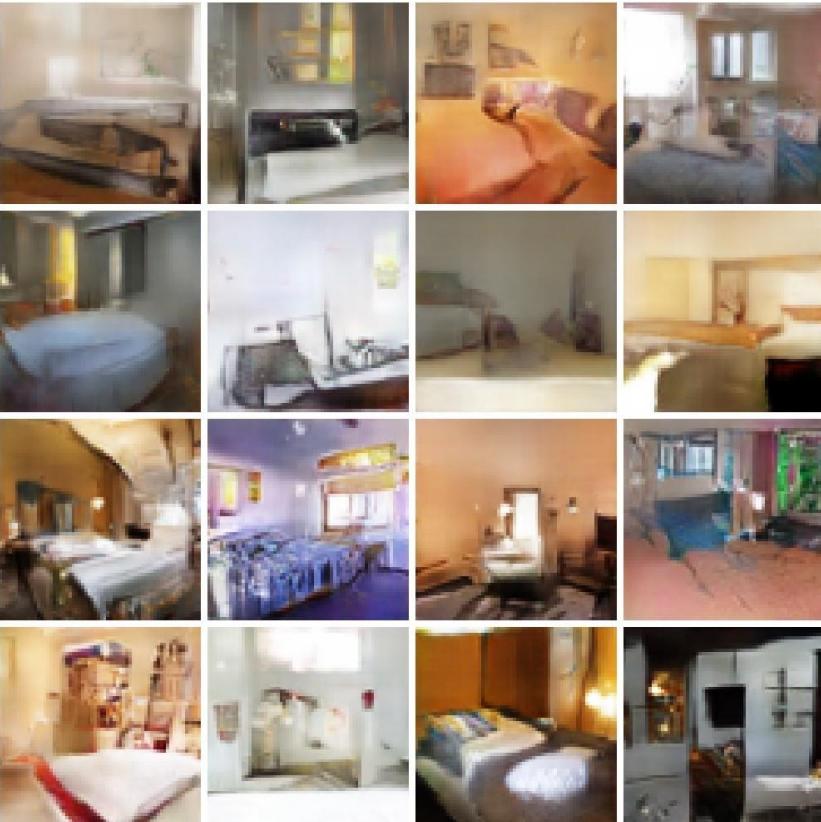


<https://github.com/hindupuravinash/the-gan-zoo>

- BGAN - Binary Generative Adversarial Networks for Image Retrieval (github)
- BiGAN - Autonomously and Simultaneously Refining Deep Neural Network Parameters by a Bi-Generative Adversarial Network Aided Genetic Algorithm
- BicycleGAN - Toward Multimodal Image-to-Image Translation (github)
- BIGAN - Adversarial Feature Learning
- BinGAN - BinGAN: Learning Compact Binary Descriptors with a Regularized GAN
- BourGAN - BourGAN: Generative Networks with Metric Embeddings
- BranchGAN - Branched Generative Adversarial Networks for Multi-Scale Image Manifold Learning
- BRE - Improving GAN Training via Binarized Representation Entropy (BRE) Regularization (github)
- BridgeGAN - Generative Adversarial Frontal View to Bird View Synthesis
- BS-GAN - Boundary-Seeking Generative Adversarial Networks
- BubGAN - BubGAN: Bubble Generative Adversarial Networks for Synthesizing Realistic Bubbly Flow Images
- BWGAN - Banach Wasserstein GAN
- C-GAN - Face Aging with Contextual Generative Adversarial Nets
- C-RNN-GAN - C-RNN-GAN: Continuous recurrent neural networks with adversarial training (github)
- CA-GAN - Composition-aided Sketch-realistic Portrait Generation
- CaloGAN - CaloGAN: Simulating 3D High Energy Particle Showers in Multi-Layer Electromagnetic Calorimeters with Generative Adversarial Networks (github)
- CAN - CAN: Creative Adversarial Networks, Generating Art by Learning About Styles and Deviating from Style Norms
- CapsGAN - CapsGAN: Using Dynamic Routing for Generative Adversarial Networks
- CapsuleGAN - CapsuleGAN: Generative Adversarial Capsule Network
- CatGAN - Unsupervised and Semi-supervised Learning with Categorical Generative Adversarial Networks
- CatGAN - CatGAN: Coupled Adversarial Transfer for Domain Generation
- CausalGAN - CausalGAN: Learning Causal Implicit Generative Models with Adversarial Training
- CC-GAN - Semi-Supervised Learning with Context-Conditional Generative Adversarial Networks (github)
- cd-GAN - Conditional Image-to-Image Translation
- CDcGAN - Simultaneously Color-Depth Super-Resolution with Conditional Generative Adversarial Network
- CE-GAN - Deep Learning for Imbalance Data Classification using Class Expert Generative Adversarial Network
- CFG-GAN - Composite Functional Gradient Learning of Generative Adversarial Models
- CGAN - Conditional Generative Adversarial Nets
- CGAN - Controllable Generative Adversarial Networks
- Chekhov GAN - An Online Learning Approach to Generative Adversarial Networks
- cIGAN - Conditional Infilling GANs for Data Augmentation in Mammogram Classification
- CinCGAN - Unsupervised Image Super-Resolution using Cycle-in-Cycle Generative Adversarial Networks
- CipherGAN - Unsupervised Cipher Cracking Using Discrete GANs
- ClusterGAN - ClusterGAN: Latent Space Clustering in Generative Adversarial Networks
- CM-GAN - CM-GANs: Cross-modal Generative Adversarial Networks for Common Representation Learning
- CoAtt-GAN - Are You Talking to Me? Reasoned Visual Dialog Generation through Adversarial Learning
- CoGAN - Coupled Generative Adversarial Networks
- ComboGAN - ComboGAN: Unrestrained Scalability for Image Domain Translation (github)
- ConceptGAN - Learning Compositional Visual Concepts with Mutual Consistency
- Conditional cycleGAN - Conditional CycleGAN for Attribute Guided Face Image Generation
- contrast-GAN - Generative Semantic Manipulation with Contrasting GAN
- Context-RNN-GAN - Contextual RNN-GANs for Abstract Reasoning Diagram Generation
- CorrGAN - Correlated discrete data generation using adversarial training
- Coulomb GAN - Coulomb GANs: Provably Optimal Nash Equilibria via Potential Fields
- Cover-GAN - Generative Steganography with Kerckhoff's Principle based on Generative Adversarial Networks
- cowboy - Defending Against Adversarial Attacks by Leveraging an Entire GAN
- CR-GAN - CR-GAN: Learning Complete Representations for Multi-view Generation
- Cramér GAN - The Cramer Distance as a Solution to Biased Wasserstein Gradients
- Cross-GAN - Crossing Generative Adversarial Networks for Cross-View Person Re-identification
- crVAE-GAN - crVAE-GAN: Learning Compositional Visual Concepts with Mutual Consistency
- CS-GAN - Improving Neural Machine Translation with Conditional Sequence Generative Adversarial Nets
- CSG - Speech-Driven Expressive Talking Lips with Conditional Sequential Generative Adversarial Networks
- CT-GAN - CT-GAN: Conditional Transformation Generative Adversarial Network for Image Attribute Modification
- CVAE-GAN - CVAE-GAN: Fine-Grained Image Generation through Asymmetric Training
- CycleGAN - Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks

# GAN Improvements: Improved Loss Functions

Wasserstein GAN (WGAN)



Arjovsky, Chintala, and Bottou, "Wasserstein GAN", 2017

WGAN with Gradient Penalty  
(WGAN-GP)



Gulrajani et al, "Improved Training of  
Wasserstein GANs", NeurIPS 2017

# GAN Improvements: Higher Resolution

256 x 256 bedrooms

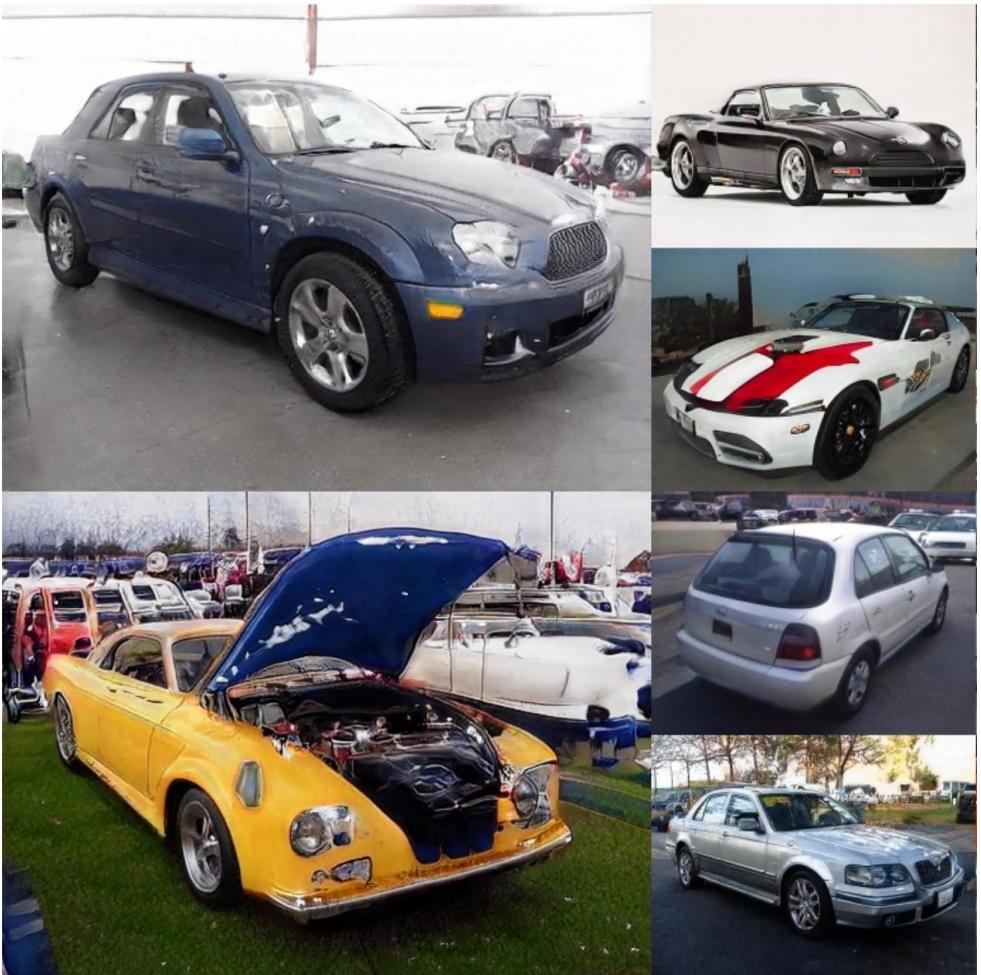


1024 x 1024 faces



# GAN Improvements: Higher Resolution

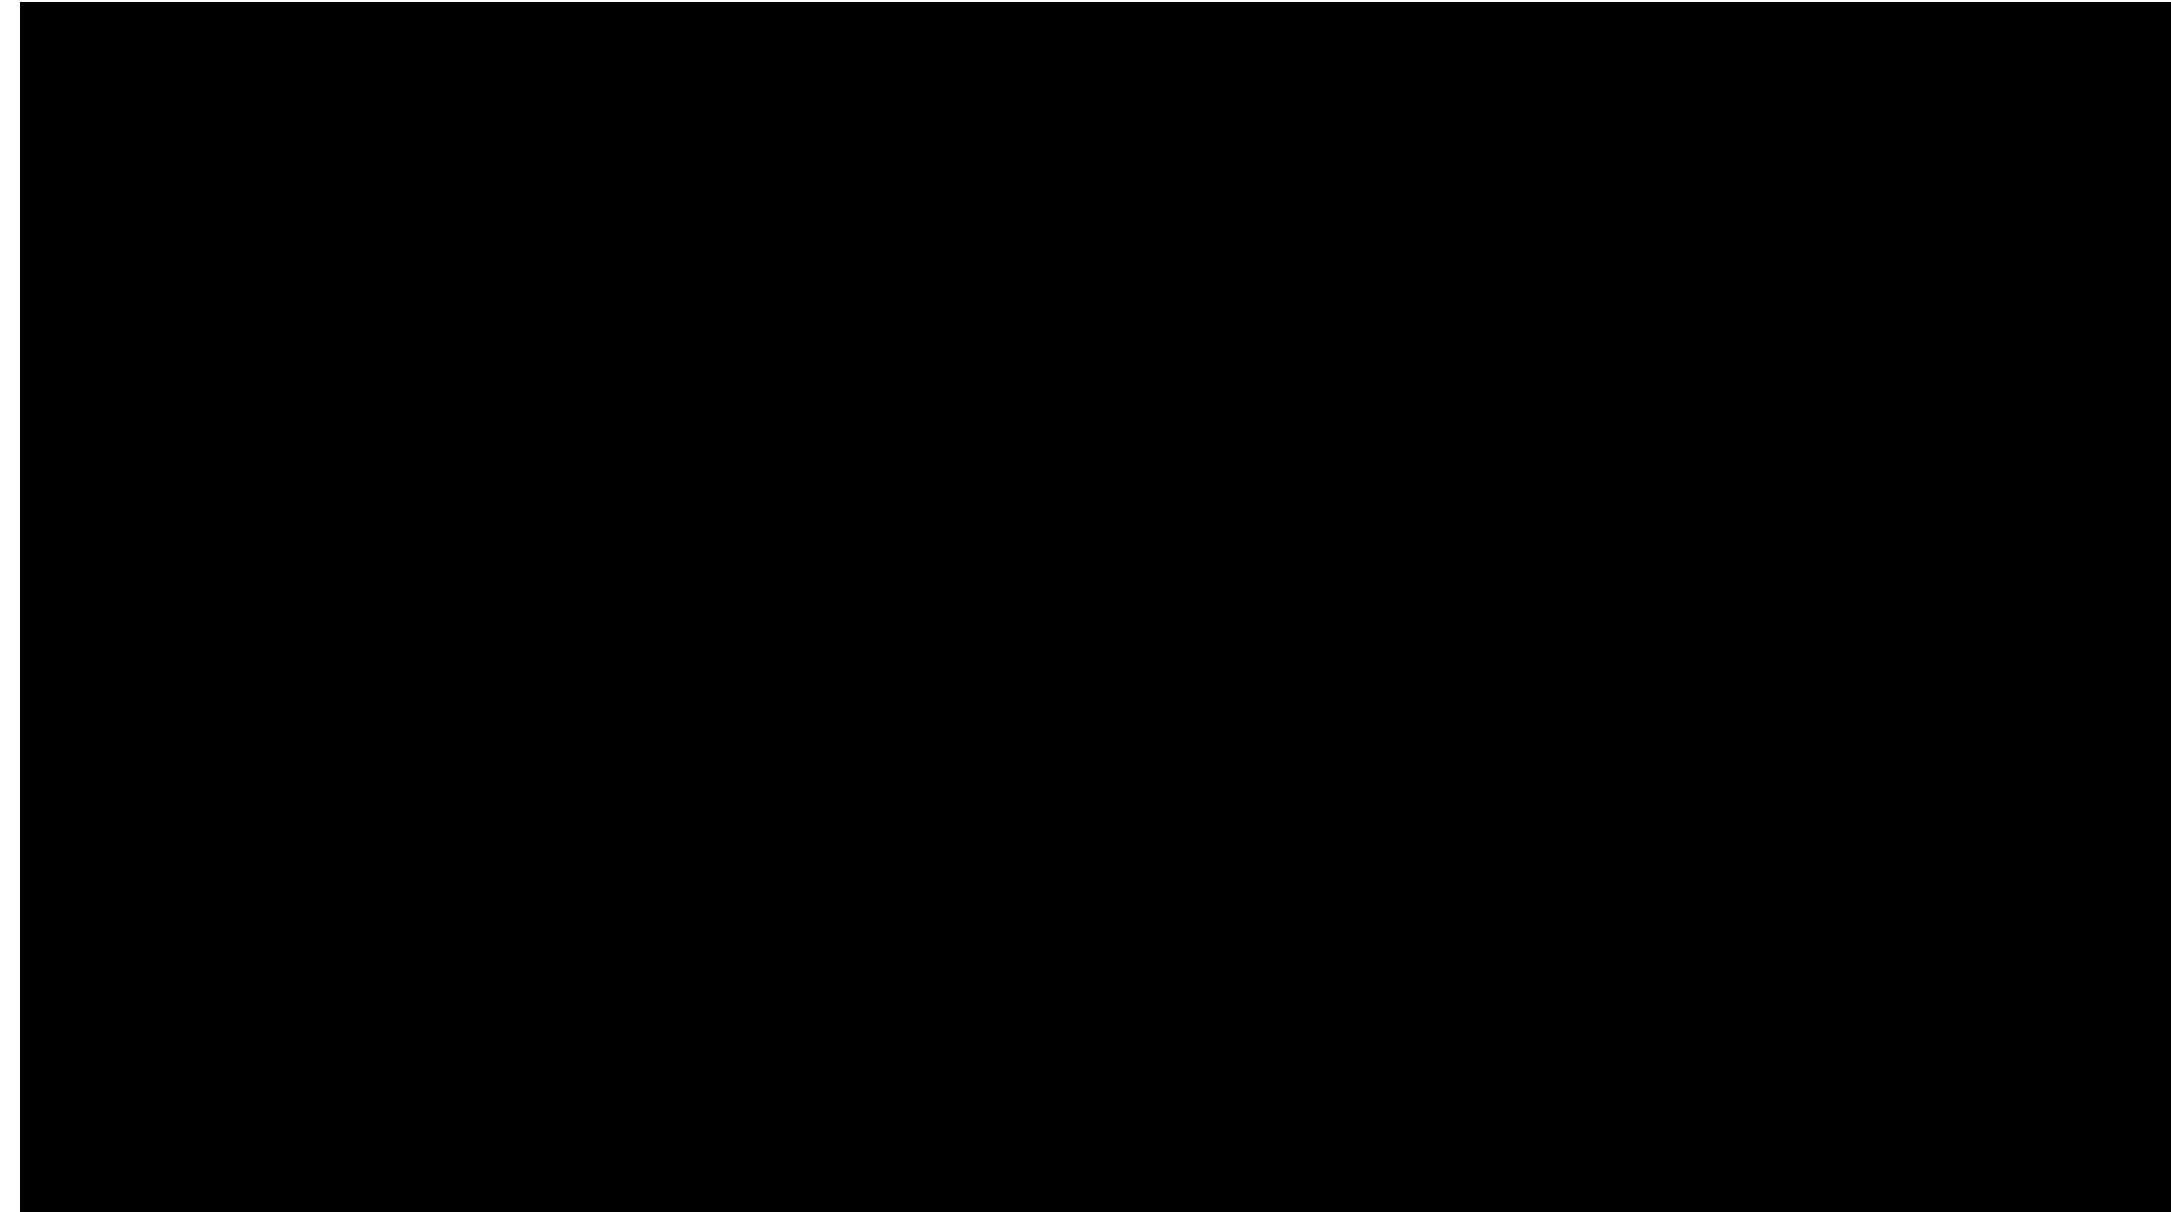
512 x 384 cars



1024 x 1024 faces



[Images](#) are licensed under [CC BY-NC 4.0](#)



Karras et al, “A Style-Based Generator Architecture for Generative Adversarial Networks”, CVPR 2019

Video is licensed under [CC BY-NC 4.0](#).

Source: [https://drive.google.com/drive/folders/1NFO7\\_vH0t98J13ckJYFd7kuaTkyeRJ86](https://drive.google.com/drive/folders/1NFO7_vH0t98J13ckJYFd7kuaTkyeRJ86)

# StyleGAN2

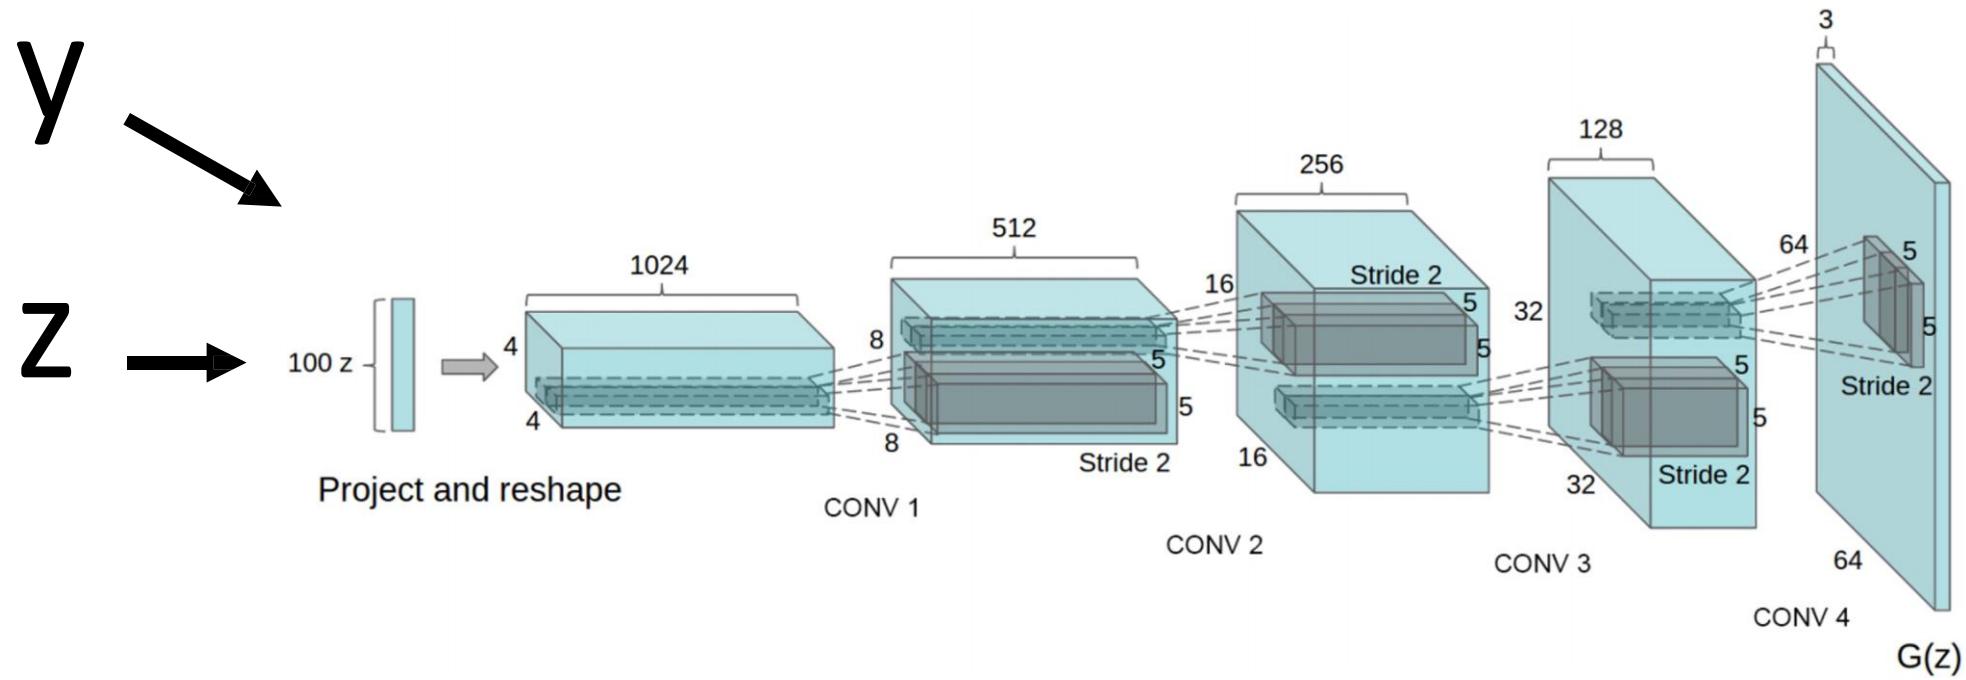


Karras et al, "Analyzing and Improving the Image Quality of StyleGAN", CVPR 2020

# Conditional GANs

**Recall:** Conditional Generative Models learn  $p(x|y)$  instead of  $p(x)$

Make generator and discriminator both take label  $y$  as an additional input!



# Conditional GANs: Conditional Batch Normalization

## Batch Normalization

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \gamma_j \hat{x}_{i,j} + \beta_j$$

Learn a separate scale and shift for each different label  $y$

## Conditional Batch Normalization

$$\mu_j = \frac{1}{N} \sum_{i=1}^N x_{i,j}$$

$$\sigma_j^2 = \frac{1}{N} \sum_{i=1}^N (x_{i,j} - \mu_j)^2$$

$$\hat{x}_{i,j} = \frac{x_{i,j} - \mu_j}{\sqrt{\sigma_j^2 + \epsilon}}$$

$$y_{i,j} = \boldsymbol{\gamma}_j^y \hat{x}_{i,j} + \boldsymbol{\beta}_j^y$$

# Conditional GANs: Spectral Normalization

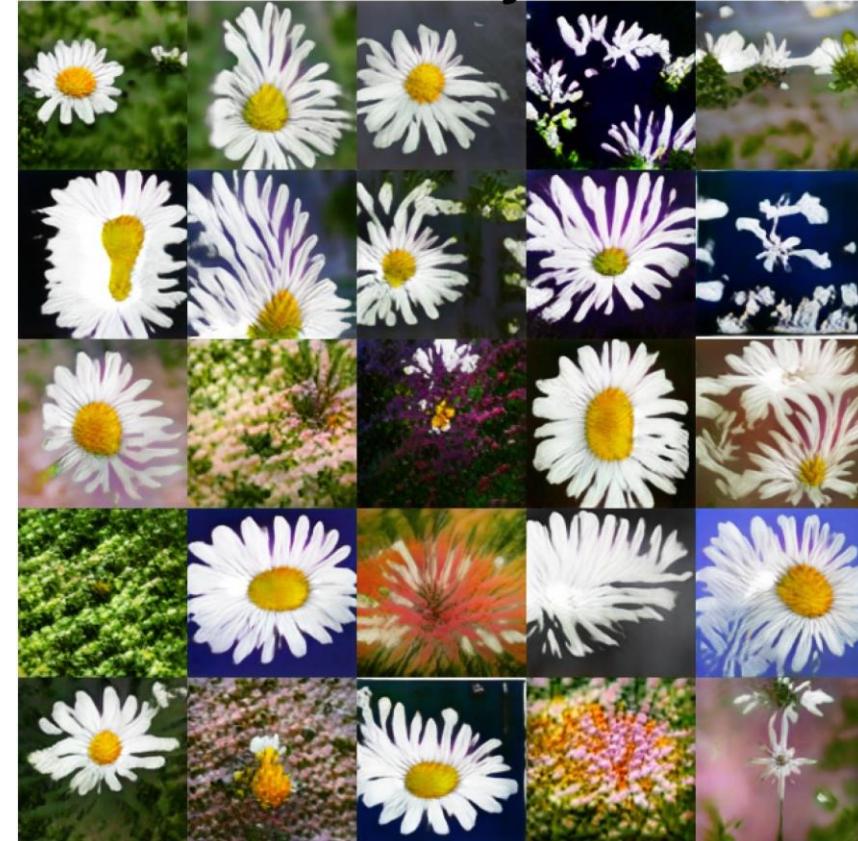
Welsh springer spaniel



Fire truck



Daisy



# Conditional GANs: Self-Attention

Goldfish



Indigo bunting



Redshank



Saint Bernard



128x128 images on ImageNet

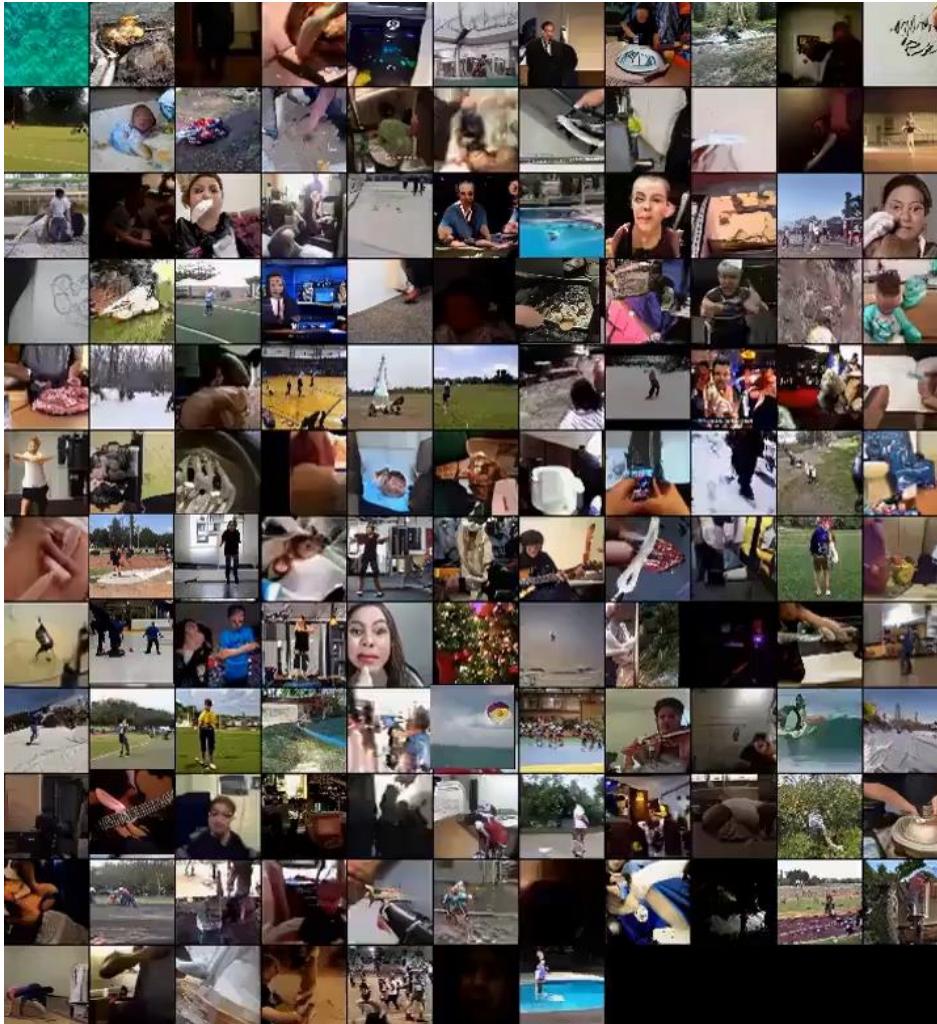
# Conditional GANs: BigGAN



512x512 images on ImageNet

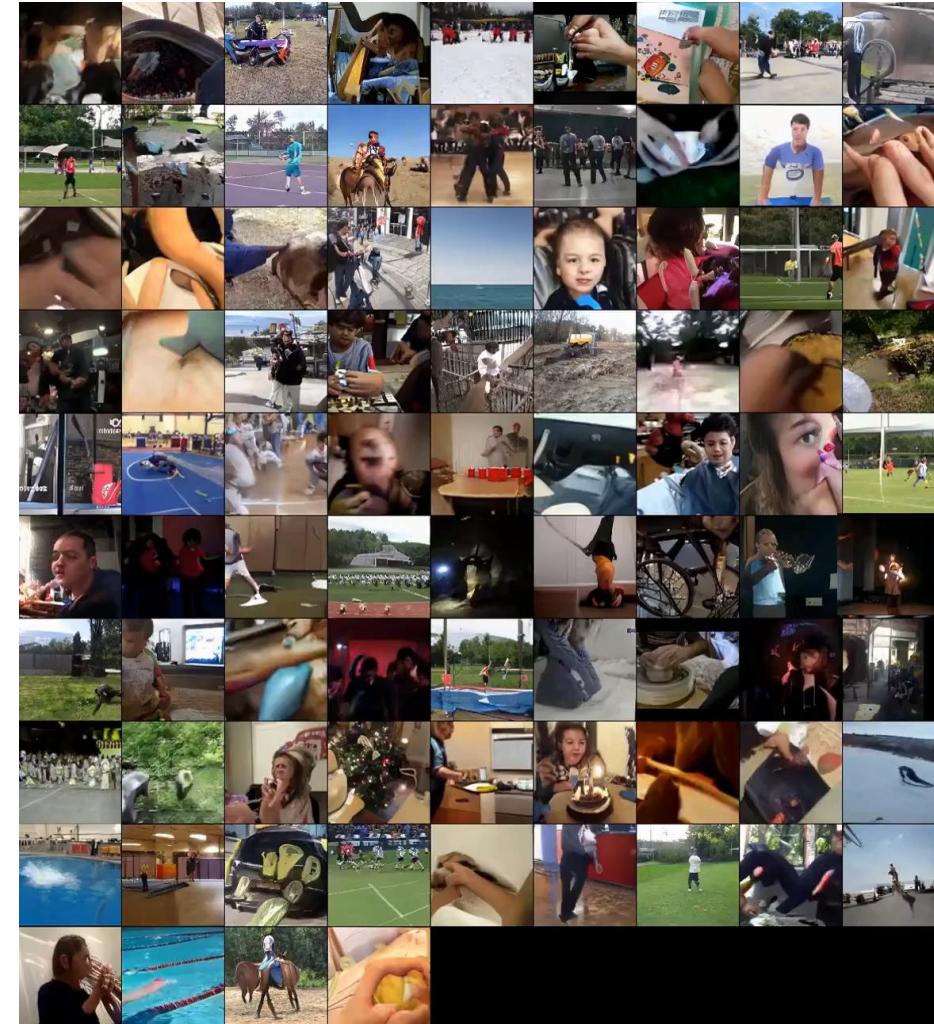
# Generating Videos with GANs

Clark et al, "Adversarial Video Generation on Complex Datasets", arXiv 2019



64x64 images, 48 frames

<https://drive.google.com/file/d/1FjOQYdUuxPXvS8yeOhXdPQMapUQaklLi/view>



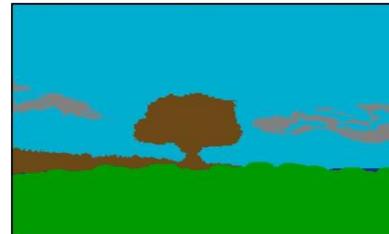
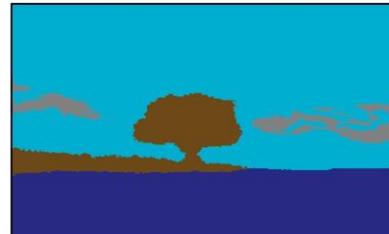
128x128 images, 12 frames

[https://drive.google.com/file/d/165Yxuvvu3viOy-39LhhSDGtczbWphj\\_i/view](https://drive.google.com/file/d/165Yxuvvu3viOy-39LhhSDGtczbWphj_i/view)

# Label Map to Image

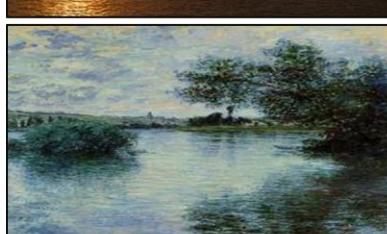
Input: Label Map

|       |          |
|-------|----------|
| cloud | sky      |
| tree  | mountain |
| sea   | grass    |



Semantic Manipulation Using Segmentation Map

Input:  
Style  
Image



Stylization using Guide Images



# Conditioning on more than labels! Text to Image

This bird is red and brown in color, with a stubby beak



The bird is short and stubby with yellow on its body



A bird with a medium orange bill white body gray wings and webbed feet



This small black bird has a short, slightly curved bill and long legs



A picture of a very clean living room



A group of people on skis stand in the snow



Eggs fruit candy nuts and meat served on white dish



A street sign on a stoplight pole in the middle of a day



Zhang et al, "StackGAN++: Realistic Image Synthesis with Stacked Generative Adversarial Networks.", TPAMI 2018

Zhang et al, "StackGAN: Text to Photo-realistic Image Synthesis with Stacked Generative Adversarial Networks.", ICCV 2017

Reed et al, "Generative Adversarial Text-to-Image Synthesis", ICML 2016

# Text to Image: DALL-E

**Step 1:** Train VQ-VAE (discrete grid of latent codes)

**Step 2:** Train autoregressive Transformer model to predict sequence of latent codes  
(Giant model on 250M image/text pairs)

**Step 3:** Given text prompt,  
sample new image codes; pass  
through VQ-VAE decoder to  
generate images



*an illustration of a baby hedgehog in a christmas sweater walking a dog*

# Text to Image: DALL-E

**Step 1:** Train VQ-VAE (discrete grid of latent codes)

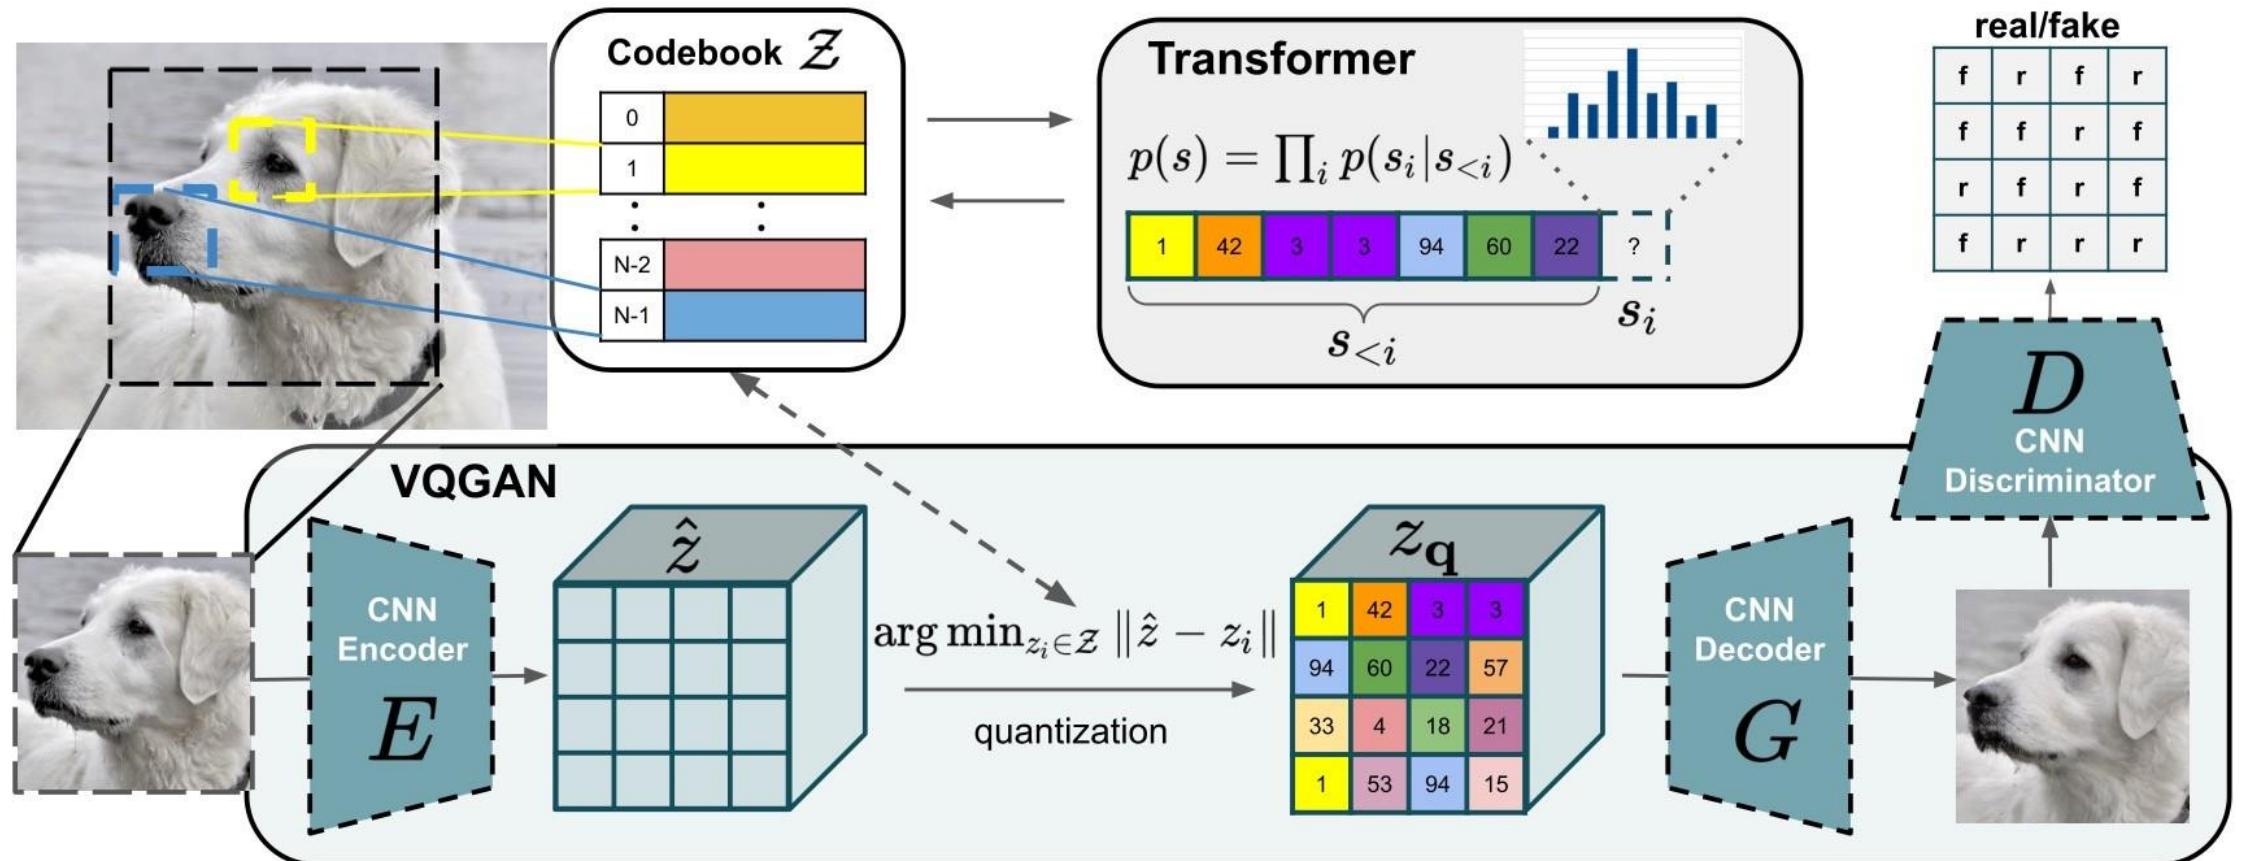
**Step 2:** Train autoregressive Transformer model to predict sequence of latent codes  
(Giant model on 250M image/text pairs)

**Step 3:** Given text prompt, sample new image codes; pass through VQ-VAE decoder to generate images



*a neon sign that reads “backprop”. a neon sign that reads “backprop”. backprop neon sign*

# VQ-GAN



# VQ-GAN (Semantic Segmentation to Image)



Esser et al, "Taming Transformers for High-Resolution Image Synthesis", CVPR 2021

# Image Super-Resolution: Low-Res to High-Res

bicubic  
(21.59dB/0.6423)



SRResNet  
(23.53dB/0.7832)



SRGAN  
(21.15dB/0.6868)



original



# Image-to-Image Translation: Pix2Pix

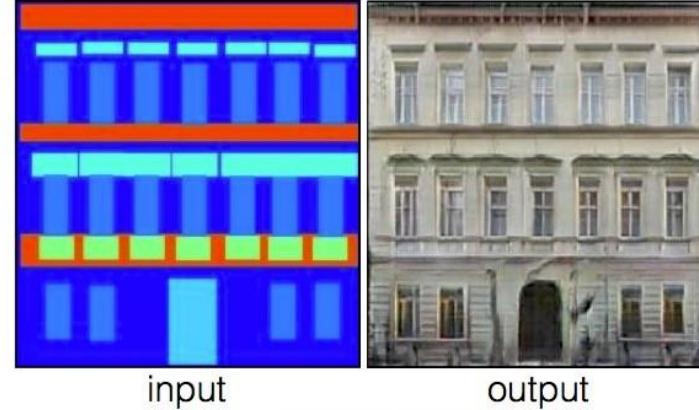
Labels to Street Scene



input

output

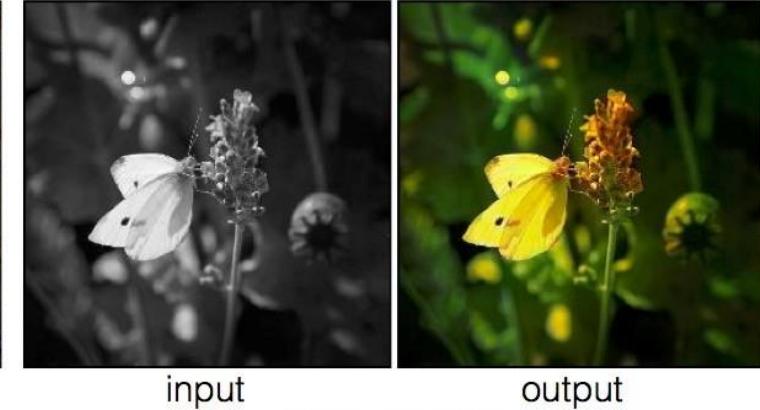
Labels to Facade



input

output

BW to Color



input

output

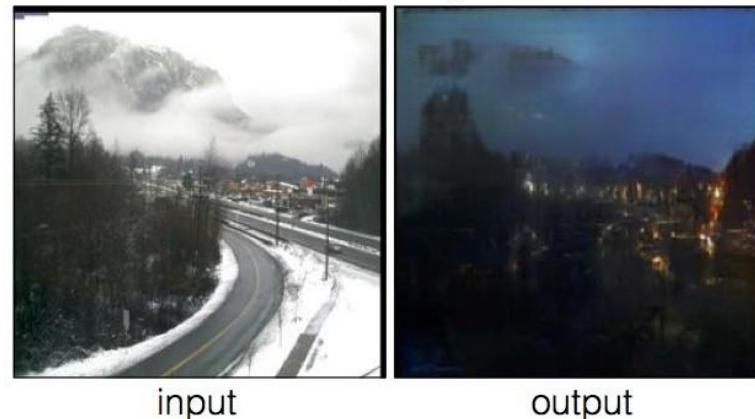
Aerial to Map



input

output

Day to Night



input

output

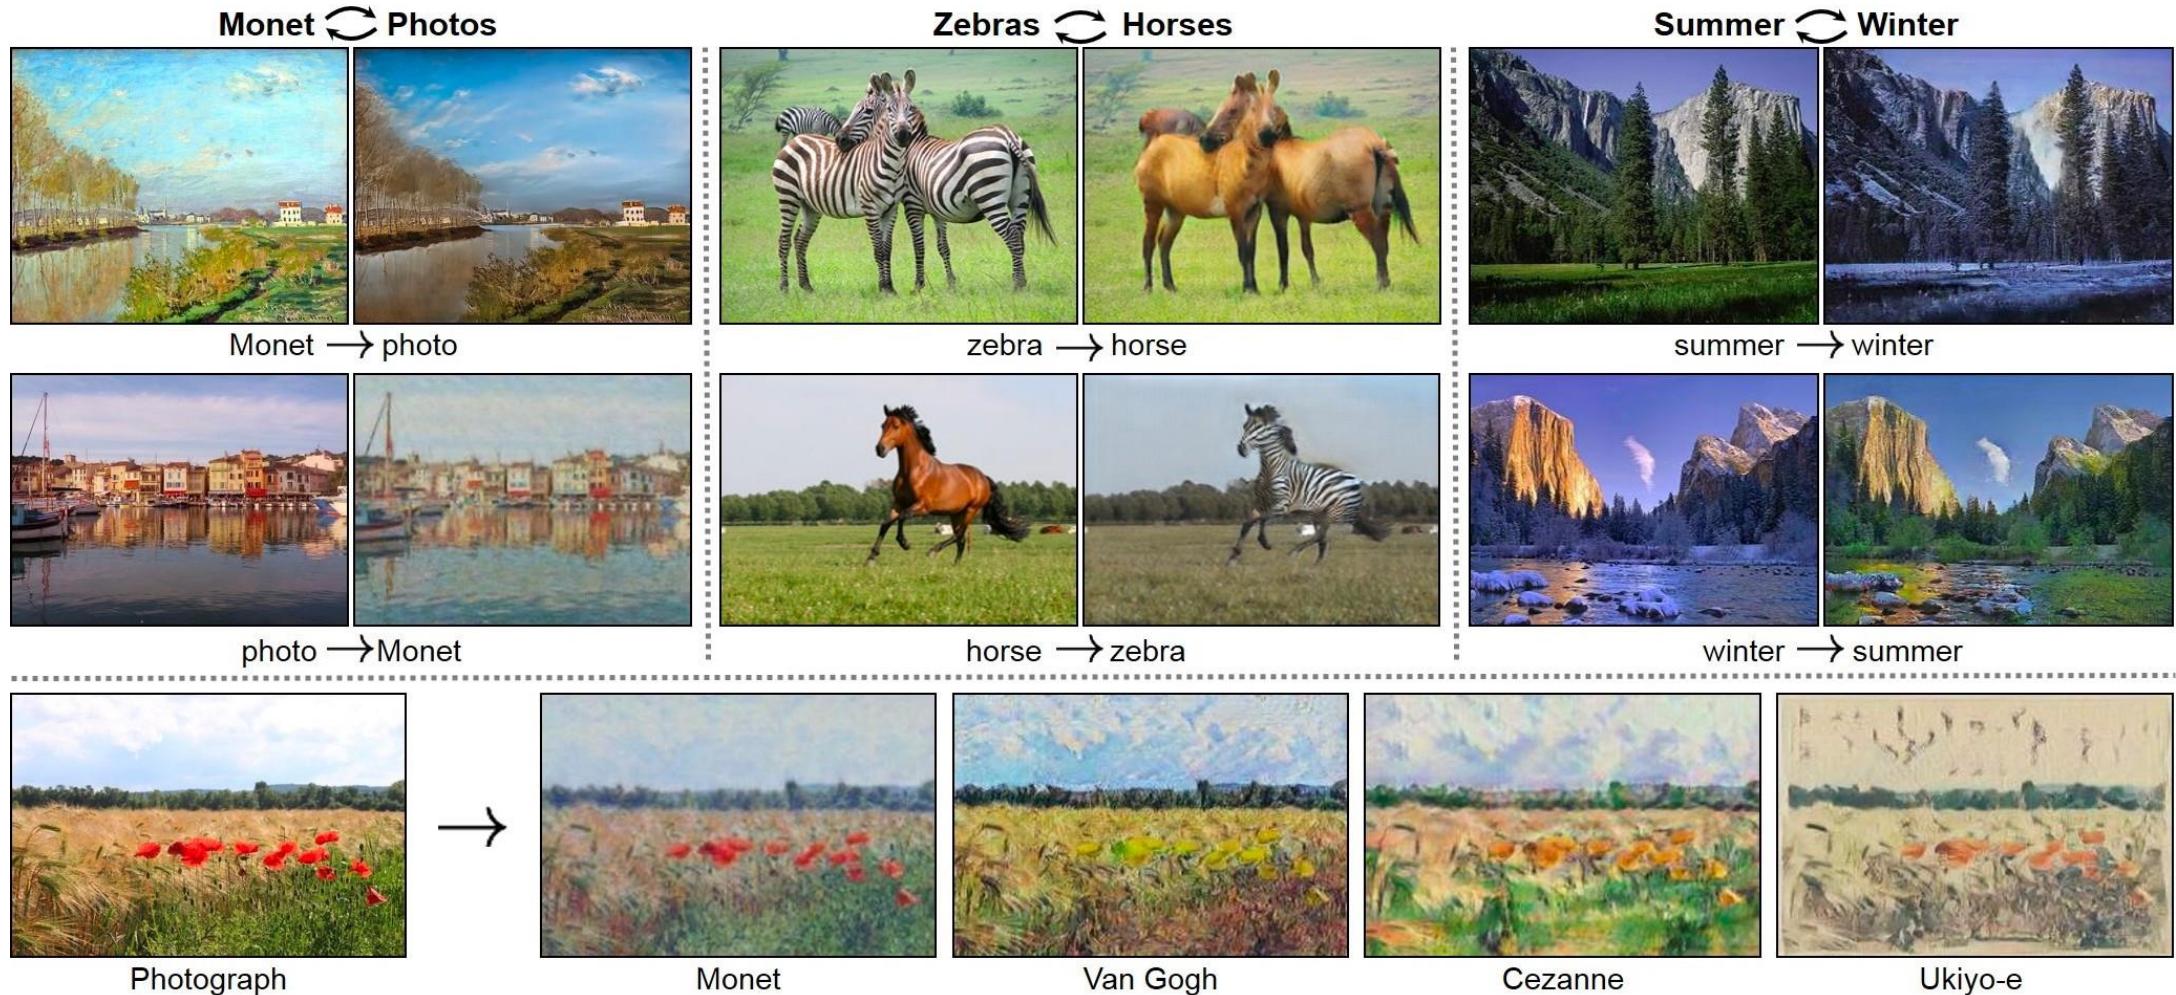
Edges to Photo



input

output

# Unpaired Image-to-Image Translation: CycleGAN



# Unpaired Image-to-Image Translation: CycleGAN

Input Video: Horse

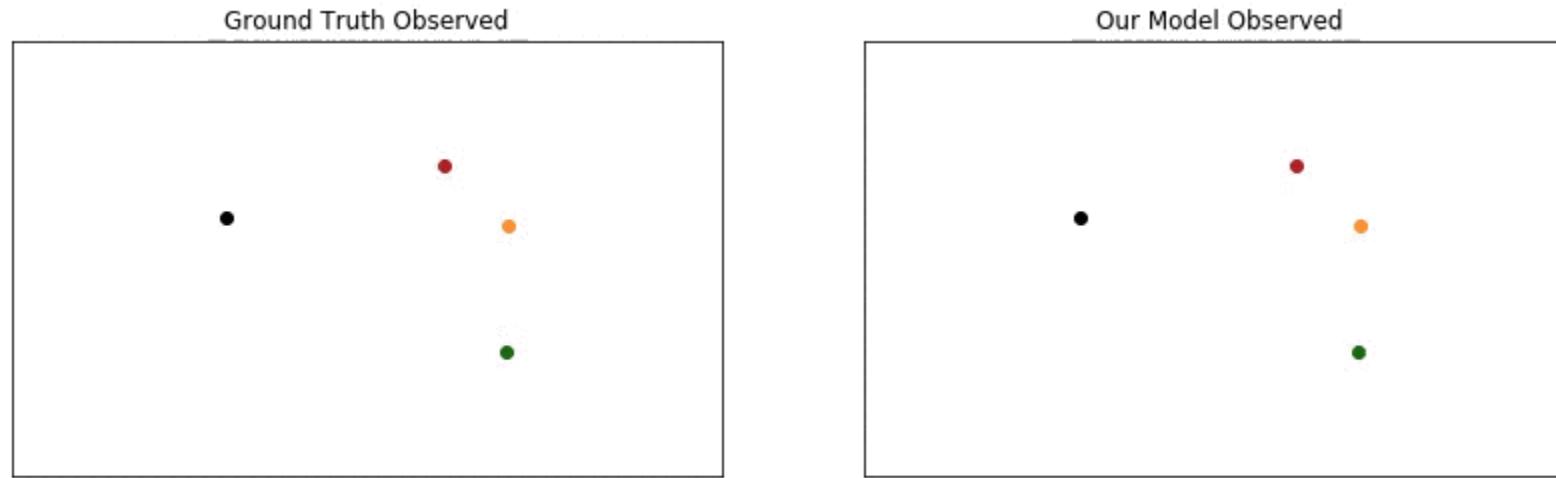


Output Video: Zebra

Zhu et al, "Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks", ICCV 2017

<https://www.youtube.com/watch?v=9reHvktowLY>

# GANs: Not just for images! Trajectory Prediction



Gupta, Johnson, Li, Savarese, Alahi, "Social GAN: Socially Acceptable Trajectories with Generative Adversarial Networks", CVPR 2018

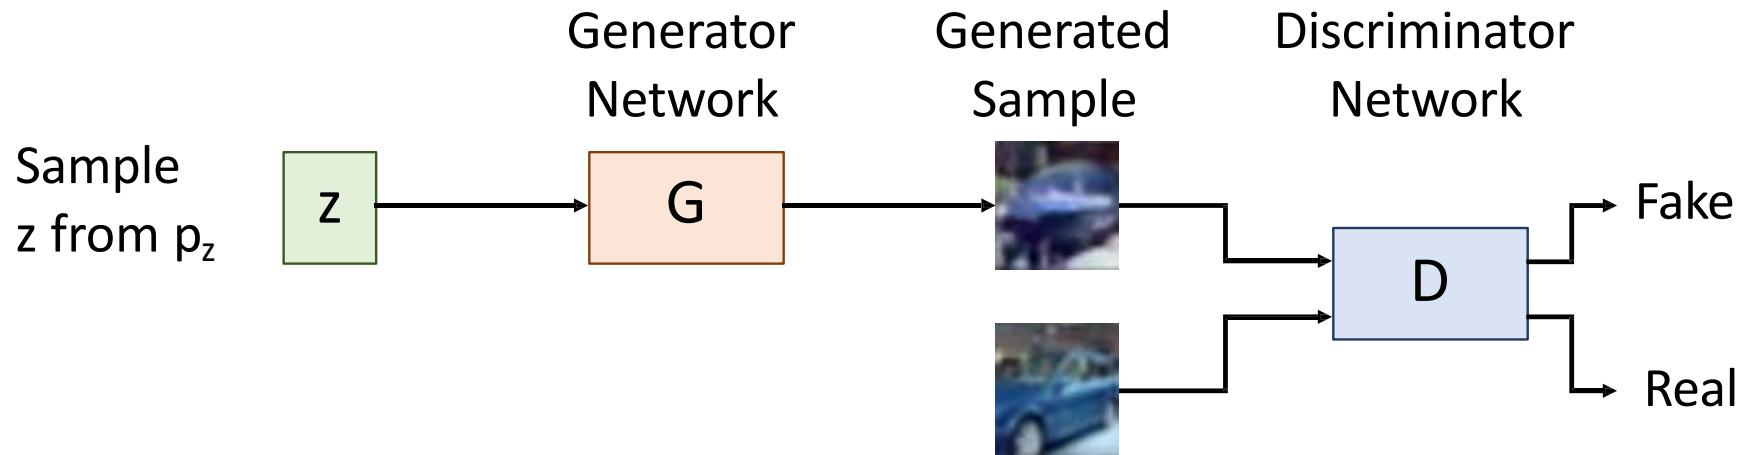
<https://github.com/agrimgupta92/sgan>

# GAN Summary

Jointly train two networks:

**Discriminator:** Classify data as real or fake

**Generator:** Generate data that fools the discriminator



Under some assumptions, generator converges to true data distribution

Many applications

Very active area of research